
Scrapy Documentation

Release 0.16.4

Scrapinghub



扫二维码获取更多资料

February 14, 2013

Contents

| | | |
|----------|--------------------------------------|-----------|
| 1 | Getting help | 3 |
| 2 | First steps | 5 |
| 2.1 | Scrapy at a glance | 5 |
| 2.2 | Installation guide | 8 |
| 2.3 | Scrapy Tutorial | 10 |
| 2.4 | Examples | 16 |
| 3 | Basic concepts | 17 |
| 3.1 | Command line tool | 17 |
| 3.2 | Items | 24 |
| 3.3 | Spiders | 28 |
| 3.4 | Link Extractors | 36 |
| 3.5 | Selectors | 38 |
| 3.6 | Item Loaders | 44 |
| 3.7 | Scrapy shell | 51 |
| 3.8 | Item Pipeline | 54 |
| 3.9 | Feed exports | 57 |
| 3.10 | Link Extractors | 61 |
| 4 | Built-in services | 65 |
| 4.1 | Logging | 65 |
| 4.2 | Stats Collection | 67 |
| 4.3 | Sending e-mail | 68 |
| 4.4 | Telnet Console | 70 |
| 4.5 | Web Service | 72 |
| 5 | Solving specific problems | 79 |
| 5.1 | Frequently Asked Questions | 79 |
| 5.2 | Debugging Spiders | 83 |
| 5.3 | Spiders Contracts | 85 |
| 5.4 | Common Practices | 87 |
| 5.5 | Broad Crawls | 89 |
| 5.6 | Using Firefox for scraping | 90 |

| | | |
|----------|---------------------------------------------|------------|
| 5.7 | Using Firebug for scraping | 91 |
| 5.8 | Debugging memory leaks | 95 |
| 5.9 | Downloading Item Images | 99 |
| 5.10 | Ubuntu packages | 104 |
| 5.11 | Scrapy Service (scrapyd) | 104 |
| 5.12 | AutoThrottle extension | 114 |
| 5.13 | Jobs: pausing and resuming crawls | 116 |
| 5.14 | DjangoItem | 117 |
| 6 | Extending Scrapy | 121 |
| 6.1 | Architecture overview | 121 |
| 6.2 | Downloader Middleware | 124 |
| 6.3 | Spider Middleware | 131 |
| 6.4 | Extensions | 135 |
| 6.5 | Core API | 140 |
| 7 | Reference | 145 |
| 7.1 | Requests and Responses | 145 |
| 7.2 | Settings | 152 |
| 7.3 | Signals | 163 |
| 7.4 | Exceptions | 166 |
| 7.5 | Item Exporters | 167 |
| 8 | All the rest | 175 |
| 8.1 | Release notes | 175 |
| 8.2 | Contributing to Scrapy | 186 |
| 8.3 | Versioning and API Stability | 188 |
| 8.4 | Experimental features | 189 |
| | Python Module Index | 191 |

This documentation contains everything you need to know about Scrapy.

Getting help

Having trouble? We'd like to help!

- Try the [FAQ](#) – it's got answers to some common questions.
- Looking for specific information? Try the *genindex* or *modindex*.
- Search for information in the [archives of the scrapy-users mailing list](#), or post a question.
- Ask a question in the #scrapy IRC channel.
- Report bugs with Scrapy in our [issue tracker](#).

First steps

2.1 Scrapy at a glance

Scrapy is an application framework for crawling web sites and extracting structured data which can be used for a wide range of useful applications, like data mining, information processing or historical archival.

Even though Scrapy was originally designed for [screen scraping](#) (more precisely, [web scraping](#)), it can also be used to extract data using APIs (such as [Amazon Associates Web Services](#)) or as a general purpose web crawler.

The purpose of this document is to introduce you to the concepts behind Scrapy so you can get an idea of how it works and decide if Scrapy is what you need.

When you're ready to start a project, you can [start with the tutorial](#).

2.1.1 Pick a website

So you need to extract some information from a website, but the website doesn't provide any API or mechanism to access that info programmatically. Scrapy can help you extract that information.

Let's say we want to extract the URL, name, description and size of all torrent files added today in the [Mininova](#) site.

The list of all torrents added today can be found on this page:

<http://www.mininova.org/today>

2.1.2 Define the data you want to scrape

The first thing is to define the data we want to scrape. In Scrapy, this is done through *Scrapy Items* (Torrent files, in this case).

This would be our Item:

```
from scrapy.item import Item, Field

class Torrent(Item):
    url = Field()
    name = Field()
    description = Field()
    size = Field()
```

2.1.3 Write a Spider to extract the data

The next thing is to write a Spider which defines the start URL (<http://www.mininova.org/today>), the rules for following links and the rules for extracting the data from pages.

If we take a look at that page content we'll see that all torrent URLs are like <http://www.mininova.org/tor/NUMBER> where NUMBER is an integer. We'll use that to construct the regular expression for the links to follow: `/tor/\d+`.

We'll use [XPath](#) for selecting the data to extract from the web page HTML source. Let's take one of those torrent pages:

<http://www.mininova.org/tor/2657665>

And look at the page HTML source to construct the XPath to select the data we want which is: torrent name, description and size.

By looking at the page HTML source we can see that the file name is contained inside a `<h1>` tag:

```
<h1>Home [2009] [Eng] XviD-ovd</h1>
```

An XPath expression to extract the name could be:

```
//h1/text()
```

And the description is contained inside a `<div>` tag with `id="description"`:

```
<h2>Description:</h2>
```

```
<div id="description">
```

```
"HOME" - a documentary film by Yann Arthus-Bertrand
```

```
<br/>
```

```
<br/>
```

```
***
```

```
<br/>
```

```
<br/>
```

```
"We are living in exceptional times. Scientists tell us that we have 10 years to change the way we live."
```

```
...
```

An XPath expression to select the description could be:

```
//div[@id='description']
```

Finally, the file size is contained in the second `<p>` tag inside the `<div>` tag with `id=specifications`:

```
<div id="specifications">
```

```
<p>
```

```
<strong>Category:</strong>
```

```
<a href="/cat/4">Movies</a> &gt; <a href="/sub/35">Documentary</a>
```

```
</p>
```

```
<p>
```

```
<strong>Total size:</strong>
```

```
699.79&nbsp;megabyte</p>
```

An XPath expression to select the description could be:

```
//div[@id='specifications']/p[2]/text()[2]
```

For more information about XPath see the [XPath reference](#).

Finally, here's the spider code:

```
class MininovaSpider(CrawlSpider):

    name = 'mininova.org'
    allowed_domains = ['mininova.org']
    start_urls = ['http://www.mininova.org/today']
    rules = [Rule(SgmlLinkExtractor(allow=['/tor/\d+']), 'parse_torrent')]

    def parse_torrent(self, response):
        x = HtmlXPathSelector(response)

        torrent = TorrentItem()
        torrent['url'] = response.url
        torrent['name'] = x.select("//h1/text()").extract()
        torrent['description'] = x.select("//div[@id='description']").extract()
        torrent['size'] = x.select("//div[@id='info-left']/p[2]/text()[2]").extract()
        return torrent
```

For brevity's sake, we intentionally left out the import statements. The `Torrent` item is *defined above*.

2.1.4 Run the spider to extract the data

Finally, we'll run the spider to crawl the site and output file `scraped_data.json` with the scraped data in JSON format:

```
scrapy crawl mininova.org -o scraped_data.json -t json
```

This uses *feed exports* to generate the JSON file. You can easily change the export format (XML or CSV, for example) or the storage backend (FTP or [Amazon S3](#), for example).

You can also write an *item pipeline* to store the items in a database very easily.

2.1.5 Review scraped data

If you check the `scraped_data.json` file after the process finishes, you'll see the scraped items there:

```
[{"url": "http://www.mininova.org/tor/2657665", "name": ["Home[2009][Eng]XviD-ovd"], "description":
# ... other items ...
}]
```

You'll notice that all field values (except for the `url` which was assigned directly) are actually lists. This is because the *selectors* return lists. You may want to store single values, or perform some additional parsing/cleansing to the values. That's what *Item Loaders* are for.

2.1.6 What else?

You've seen how to extract and store items from a website using Scrapy, but this is just the surface. Scrapy provides a lot of powerful features for making scraping easy and efficient, such as:

- Built-in support for *selecting and extracting* data from HTML and XML sources
- Built-in support for cleaning and sanitizing the scraped data using a collection of reusable filters (called *Item Loaders*) shared between all the spiders.
- Built-in support for *generating feed exports* in multiple formats (JSON, CSV, XML) and storing them in multiple backends (FTP, S3, local filesystem)

- A media pipeline for *automatically downloading images* (or any other media) associated with the scraped items
- Support for *extending Scrapy* by plugging your own functionality using *signals* and a well-defined API (middlewares, *extensions*, and *pipelines*).
- Wide range of built-in middlewares and extensions for:
 - cookies and session handling
 - HTTP compression
 - HTTP authentication
 - HTTP cache
 - user-agent spoofing
 - robots.txt
 - crawl depth restriction
 - and more
- Robust encoding support and auto-detection, for dealing with foreign, non-standard and broken encoding declarations.
- Support for creating spiders based on pre-defined templates, to speed up spider creation and make their code more consistent on large projects. See *genspider* command for more details.
- Extensible *stats collection* for multiple spider metrics, useful for monitoring the performance of your spiders and detecting when they get broken
- An *Interactive shell console* for trying XPaths, very useful for writing and debugging your spiders
- A *System service* designed to ease the deployment and run of your spiders in production.
- A built-in *Web service* for monitoring and controlling your bot
- A *Telnet console* for hooking into a Python console running inside your Scrapy process, to introspect and debug your crawler
- *Logging* facility that you can hook on to for catching errors during the scraping process.
- Support for crawling based on URLs discovered through *Sitemaps*
- A caching DNS resolver

2.1.7 What's next?

The next obvious steps are for you to [download Scrapy](#), read [the tutorial](#) and join [the community](#). Thanks for your interest!

2.2 Installation guide

2.2.1 Pre-requisites

The installation steps assume that you have the following things installed:

- Python 2.6 or 2.7
- OpenSSL. This comes preinstalled in all operating systems except Windows (see [Platform specific installation notes](#))

- `pip` or `easy_install` Python package managers

2.2.2 Installing Scrapy

You can install Scrapy using `easy_install` or `pip` (which is the canonical way to distribute and install Python packages).

Note: Check *Platform specific installation notes* first.

To install using `pip`:

```
pip install Scrapy
```

To install using `easy_install`:

```
easy_install Scrapy
```

2.2.3 Platform specific installation notes

Windows

After installing Python, follow these steps before installing Scrapy:

- add the `C:\python27\Scripts` and `C:\python27` folders to the system path by adding those directories to the `PATH` environment variable from the [Control Panel](#).
- install OpenSSL by following these steps:
 1. go to [Win32 OpenSSL page](#)
 2. download Visual C++ 2008 redistributables for your Windows and architecture
 3. download OpenSSL for your Windows and architecture (the regular version, not the light one)
 4. add the `c:\openssl-win32\bin` (or similar) directory to your `PATH`, the same way you added `python27` in the first step“ in the first step
- some binary packages that Scrapy depends on (like Twisted, lxml and pyOpenSSL) require a compiler available to install, and fail if you don't have Visual Studio installed. You can find Windows installers for those in the following links. Make sure you respect your Python version and Windows architecture.
 - pywin32: <http://sourceforge.net/projects/pywin32/files/>
 - Twisted: <http://twistedmatrix.com/trac/wiki/Downloads>
 - zope.interface: download the egg from [zope.interface pypi page](#) and install it by running `easy_install file.egg`
 - lxml: <http://pypi.python.org/pypi/lxml/>
 - pyOpenSSL: <https://launchpad.net/pyopenssl>

Ubuntu 9.10 or above

Don't use the `python-scrapy` package provided by Ubuntu, they are typically too old and slow to catch up with latest Scrapy.

Instead, use the official *Ubuntu Packages*, which already solve all dependencies for you and are continuously updated with the latest bug fixes.

2.3 Scrapy Tutorial

In this tutorial, we'll assume that Scrapy is already installed on your system. If that's not the case, see [Installation guide](#).

We are going to use [Open directory project \(dmoz\)](#) as our example domain to scrape.

This tutorial will walk you through these tasks:

1. Creating a new Scrapy project
2. Defining the Items you will extract
3. Writing a *spider* to crawl a site and extract *Items*
4. Writing an *Item Pipeline* to store the extracted Items

Scrapy is written in [Python](#). If you're new to the language you might want to start by getting an idea of what the language is like, to get the most out of Scrapy. If you're already familiar with other languages, and want to learn Python quickly, we recommend [Learn Python The Hard Way](#). If you're new to programming and want to start with Python, take a look at [this list of Python resources for non-programmers](#).

2.3.1 Creating a project

Before you start scraping, you will have set up a new Scrapy project. Enter a directory where you'd like to store your code and then run:

```
scrapy startproject tutorial
```

This will create a `tutorial` directory with the following contents:

```
tutorial/  
  scrapy.cfg  
  tutorial/  
    __init__.py  
    items.py  
    pipelines.py  
    settings.py  
    spiders/  
      __init__.py  
      ...
```

These are basically:

- `scrapy.cfg`: the project configuration file
- `tutorial/`: the project's python module, you'll later import your code from here.
- `tutorial/items.py`: the project's items file.
- `tutorial/pipelines.py`: the project's pipelines file.
- `tutorial/settings.py`: the project's settings file.
- `tutorial/spiders/`: a directory where you'll later put your spiders.

2.3.2 Defining our Item

Items are containers that will be loaded with the scraped data; they work like simple python dicts but provide additional protecting against populating undeclared fields, to prevent typos.

They are declared by creating an `scrapy.item.Item` class and defining its attributes as `scrapy.item.Field` objects, like you will in an ORM (don't worry if you're not familiar with ORMs, you will see that this is an easy task).

We begin by modeling the item that we will use to hold the sites data obtained from dmoz.org, as we want to capture the name, url and description of the sites, we define fields for each of these three attributes. To do that, we edit `items.py`, found in the `tutorial` directory. Our `Item` class looks like this:

```
from scrapy.item import Item, Field

class DmozItem(Item):
    title = Field()
    link = Field()
    desc = Field()
```

This may seem complicated at first, but defining the item allows you to use other handy components of Scrapy that need to know how your item looks like.

2.3.3 Our first Spider

Spiders are user-written classes used to scrape information from a domain (or group of domains).

They define an initial list of URLs to download, how to follow links, and how to parse the contents of those pages to extract *items*.

To create a Spider, you must subclass `scrapy.spider.BaseSpider`, and define the three main, mandatory, attributes:

- `name`: identifies the Spider. It must be unique, that is, you can't set the same name for different Spiders.
- `start_urls`: is a list of URLs where the Spider will begin to crawl from. So, the first pages downloaded will be those listed here. The subsequent URLs will be generated successively from data contained in the start URLs.
- `parse()` is a method of the spider, which will be called with the downloaded `Response` object of each start URL. The response is passed to the method as the first and only argument.

This method is responsible for parsing the response data and extracting scraped data (as scraped items) and more URLs to follow.

The `parse()` method is in charge of processing the response and returning scraped data (as `Item` objects) and more URLs to follow (as `Request` objects).

This is the code for our first Spider; save it in a file named `dmoz_spider.py` under the `dmoz/spiders` directory:

```
from scrapy.spider import BaseSpider

class DmozSpider(BaseSpider):
    name = "dmoz"
    allowed_domains = ["dmoz.org"]
    start_urls = [
        "http://www.dmoz.org/Computers/Programming/Languages/Python/Books/",
        "http://www.dmoz.org/Computers/Programming/Languages/Python/Resources/"
    ]

    def parse(self, response):
        filename = response.url.split("/")[-2]
        open(filename, 'wb').write(response.body)
```

Crawling

To put our spider to work, go to the project's top level directory and run:

```
scrapy crawl dmoz
```

The `crawl dmoz` command runs the spider for the `dmoz.org` domain. You will get an output similar to this:

```
2008-08-20 03:51:13-0300 [scrapy] INFO: Started project: dmoz
2008-08-20 03:51:13-0300 [tutorial] INFO: Enabled extensions: ...
2008-08-20 03:51:13-0300 [tutorial] INFO: Enabled downloader middlewares: ...
2008-08-20 03:51:13-0300 [tutorial] INFO: Enabled spider middlewares: ...
2008-08-20 03:51:13-0300 [tutorial] INFO: Enabled item pipelines: ...
2008-08-20 03:51:14-0300 [dmoz] INFO: Spider opened
2008-08-20 03:51:14-0300 [dmoz] DEBUG: Crawled <http://www.dmoz.org/Computers/Programming/Languages/Python/
2008-08-20 03:51:14-0300 [dmoz] DEBUG: Crawled <http://www.dmoz.org/Computers/Programming/Languages/Python/
2008-08-20 03:51:14-0300 [dmoz] INFO: Spider closed (finished)
```

Pay attention to the lines containing `[dmoz]`, which corresponds to our spider. You can see a log line for each URL defined in `start_urls`. Because these URLs are the starting ones, they have no referrers, which is shown at the end of the log line, where it says `(referer: <None>)`.

But more interesting, as our parse method instructs, two files have been created: *Books* and *Resources*, with the content of both URLs.

What just happened under the hood?

Scrapy creates `scrapy.http.Request` objects for each URL in the `start_urls` attribute of the Spider, and assigns them the `parse` method of the spider as their callback function.

These Requests are scheduled, then executed, and `scrapy.http.Response` objects are returned and then fed back to the spider, through the `parse()` method.

Extracting Items

Introduction to Selectors

There are several ways to extract data from web pages. Scrapy uses a mechanism based on [XPath](#) expressions called *XPath selectors*. For more information about selectors and other extraction mechanisms see the *XPath selectors documentation*.

Here are some examples of XPath expressions and their meanings:

- `/html/head/title`: selects the `<title>` element, inside the `<head>` element of a HTML document
- `/html/head/title/text()`: selects the text inside the aforementioned `<title>` element.
- `//td`: selects all the `<td>` elements
- `//div[@class="mine"]`: selects all `div` elements which contain an attribute `class="mine"`

These are just a couple of simple examples of what you can do with XPath, but XPath expressions are indeed much more powerful. To learn more about XPath we recommend [this XPath tutorial](#).

For working with XPaths, Scrapy provides a `XPathSelector` class, which comes in two flavours, `HtmlXPathSelector` (for HTML data) and `XmlXPathSelector` (for XML data). In order to use them you must instantiate the desired class with a `Response` object.

You can see selectors as objects that represent nodes in the document structure. So, the first instantiated selectors are associated to the root node, or the entire document.

Selectors have three methods (click on the method to see the complete API documentation).

- `select()`: returns a list of selectors, each of them representing the nodes selected by the xpath expression given as argument.
- `extract()`: returns a unicode string with the data selected by the XPath selector.
- `re()`: returns a list of unicode strings extracted by applying the regular expression given as argument.

Trying Selectors in the Shell

To illustrate the use of Selectors we're going to use the built-in *Scrapy shell*, which also requires IPython (an extended Python console) installed on your system.

To start a shell, you must go to the project's top level directory and run:

```
scrapy shell http://www.dmoz.org/Computers/Programming/Languages/Python/Books/
```

This is what the shell looks like:

```
[ ... Scrapy log here ... ]
```

```
[s] Available Scrapy objects:
[s] 2010-08-19 21:45:59-0300 [default] INFO: Spider closed (finished)
[s] hxs      <HtmlXPathSelector (http://www.dmoz.org/Computers/Programming/Languages/Python/Books/)>
[s] item     Item()
[s] request  <GET http://www.dmoz.org/Computers/Programming/Languages/Python/Books/>
[s] response <200 http://www.dmoz.org/Computers/Programming/Languages/Python/Books/>
[s] spider   <BaseSpider 'default' at 0x1b6c2d0>
[s] xxs      <XmlXPathSelector (http://www.dmoz.org/Computers/Programming/Languages/Python/Books/)>
[s] Useful shortcuts:
[s] shelp()      Print this help
[s] fetch(req_or_url) Fetch a new request or URL and update shell objects
[s] view(response) View response in a browser
```

```
In [1]:
```

After the shell loads, you will have the response fetched in a local `response` variable, so if you type `response.body` you will see the body of the response, or you can type `response.headers` to see its headers.

The shell also instantiates two selectors, one for HTML (in the `hxs` variable) and one for XML (in the `xxs` variable) with this response. So let's try them:

```
In [1]: hxs.select('//title')
Out[1]: [<HtmlXPathSelector (title) xpath=//title>]

In [2]: hxs.select('//title').extract()
Out[2]: [u'<title>Open Directory - Computers: Programming: Languages: Python: Books</title>']

In [3]: hxs.select('//title/text()')
Out[3]: [<HtmlXPathSelector (text) xpath=//title/text()>]

In [4]: hxs.select('//title/text()').extract()
Out[4]: [u'Open Directory - Computers: Programming: Languages: Python: Books']
```

```
In [5]: hxs.select('//title/text()').re('(\w+):')
Out[5]: [u'Computers', u'Programming', u'Languages', u'Python']
```

Extracting the data

Now, let's try to extract some real information from those pages.

You could type `response.body` in the console, and inspect the source code to figure out the XPath expressions you need to use. However, inspecting the raw HTML code there could become a very tedious task. To make this an easier task, you can use some Firefox extensions like Firebug. For more information see [Using Firebug for scraping](#) and [Using Firefox for scraping](#).

After inspecting the page source, you'll find that the web sites information is inside a `` element, in fact the *second* `` element.

So we can select each `` element belonging to the sites list with this code:

```
hxs.select('//ul/li')
```

And from them, the sites descriptions:

```
hxs.select('//ul/li/text()').extract()
```

The sites titles:

```
hxs.select('//ul/li/a/text()').extract()
```

And the sites links:

```
hxs.select('//ul/li/a/@href').extract()
```

As we said before, each `select()` call returns a list of selectors, so we can concatenate further `select()` calls to dig deeper into a node. We are going to use that property here, so:

```
sites = hxs.select('//ul/li')
for site in sites:
    title = site.select('a/text()').extract()
    link = site.select('a/@href').extract()
    desc = site.select('text()').extract()
    print title, link, desc
```

Note: For a more detailed description of using nested selectors, see [Nesting selectors](#) and [Working with relative XPath expressions](#) in the [Selectors](#) documentation

Let's add this code to our spider:

```
from scrapy.spider import BaseSpider
from scrapy.selector import HtmlXPathSelector

class DmozSpider(BaseSpider):
    name = "dmoz"
    allowed_domains = ["dmoz.org"]
    start_urls = [
        "http://www.dmoz.org/Computers/Programming/Languages/Python/Books/",
        "http://www.dmoz.org/Computers/Programming/Languages/Python/Resources/"
    ]
```

```
def parse(self, response):
    hxs = HtmlXPathSelector(response)
    sites = hxs.select('//ul/li')
    for site in sites:
        title = site.select('a/text()').extract()
        link = site.select('a/@href').extract()
        desc = site.select('text()').extract()
        print title, link, desc
```

Now try crawling the dmoz.org domain again and you'll see sites being printed in your output, run:

```
scrapy crawl dmoz
```

Using our item

Item objects are custom python dicts; you can access the values of their fields (attributes of the class we defined earlier) using the standard dict syntax like:

```
>>> item = DmozItem()
>>> item['title'] = 'Example title'
>>> item['title']
'Example title'
```

Spiders are expected to return their scraped data inside Item objects. So, in order to return the data we've scraped so far, the final code for our Spider would be like this:

```
from scrapy.spider import BaseSpider
from scrapy.selector import HtmlXPathSelector

from tutorial.items import DmozItem

class DmozSpider(BaseSpider):
    name = "dmoz"
    allowed_domains = ["dmoz.org"]
    start_urls = [
        "http://www.dmoz.org/Computers/Programming/Languages/Python/Books/",
        "http://www.dmoz.org/Computers/Programming/Languages/Python/Resources/"
    ]

    def parse(self, response):
        hxs = HtmlXPathSelector(response)
        sites = hxs.select('//ul/li')
        items = []
        for site in sites:
            item = DmozItem()
            item['title'] = site.select('a/text()').extract()
            item['link'] = site.select('a/@href').extract()
            item['desc'] = site.select('text()').extract()
            items.append(item)
        return items
```

Note: You can find a fully-functional variant of this spider in the [dirbot](https://github.com/scrapy/dirbot) project available at <https://github.com/scrapy/dirbot>

Now doing a crawl on the dmoz.org domain yields DmozItem's:

```
[dmoz] DEBUG: Scraped from <200 http://www.dmoz.org/Computers/Programming/Languages/Python/Books/>
{'desc': [u' - By David Mertz; Addison Wesley. Book in progress, full text, ASCII format. Asks i
'link': [u'http://gnosis.cx/TPiP/'],
'title': [u'Text Processing in Python']}
[dmoz] DEBUG: Scraped from <200 http://www.dmoz.org/Computers/Programming/Languages/Python/Books/>
{'desc': [u' - By Sean McGrath; Prentice Hall PTR, 2000, ISBN 0130211192, has CD-ROM. Methods to
'link': [u'http://www.informit.com/store/product.aspx?isbn=0130211192'],
'title': [u'XML Processing with Python']}
```

2.3.4 Storing the scraped data

The simplest way to store the scraped data is by using the *Feed exports*, with the following command:

```
scrapy crawl dmoz -o items.json -t json
```

That will generate a `items.json` file containing all scraped items, serialized in **JSON**.

In small projects (like the one in this tutorial), that should be enough. However, if you want to perform more complex things with the scraped items, you can write an *Item Pipeline*. As with Items, a placeholder file for Item Pipelines has been set up for you when the project is created, in `tutorial/pipelines.py`. Though you don't need to implement any item pipeline if you just want to store the scraped items.

2.3.5 Next steps

This tutorial covers only the basics of Scrapy, but there's a lot of other features not mentioned here. Check the *What else?* section in *Scrapy at a glance* chapter for a quick overview of the most important ones.

Then, we recommend you continue by playing with an example project (see *Examples*), and then continue with the section *Basic concepts*.

2.4 Examples

The best way to learn is with examples, and Scrapy is no exception. For this reason, there is an example Scrapy project named **dirbot**, that you can use to play and learn more about Scrapy. It contains the dmoz spider described in the tutorial.

This **dirbot** project is available at: <https://github.com/scrapy/dirbot>

It contains a README file with a detailed description of the project contents.

If you're familiar with git, you can checkout the code. Otherwise you can download a tarball or zip file of the project by clicking on [Downloads](#).

The [scrapy tag on Snipplr](#) is used for sharing code snippets such as spiders, middlewares, extensions, or scripts. Feel free (and encouraged!) to share any code there.

Scrapy at a glance Understand what Scrapy is and how it can help you.

Installation guide Get Scrapy installed on your computer.

Scrapy Tutorial Write your first Scrapy project.

Examples Learn more by playing with a pre-made Scrapy project.

Basic concepts

3.1 Command line tool

New in version 0.10. Scrapy is controlled through the `scrapy` command-line tool, to be referred here as the “Scrapy tool” to differentiate it from their sub-commands which we just call “commands”, or “Scrapy commands”.

The Scrapy tool provides several commands, for multiple purposes, and each one accepts a different set of arguments and options.

3.1.1 Default structure of Scrapy projects

Before delving into the command-line tool and its sub-commands, let’s first understand the directory structure of a Scrapy project.

Even though it can be modified, all Scrapy projects have the same file structure by default, similar to this:

```
scrapy.cfg
myproject/
  __init__.py
  items.py
  pipelines.py
  settings.py
  spiders/
    __init__.py
    spider1.py
    spider2.py
    ...
```

The directory where the `scrapy.cfg` file resides is known as the *project root directory*. That file contains the name of the python module that defines the project settings. Here is an example:

```
[settings]
default = myproject.settings
```

3.1.2 Using the `scrapy` tool

You can start by running the Scrapy tool with no arguments and it will print some usage help and the available commands:

```
Scrapy X.Y - no active project
```

```
Usage:
```

```
  scrapy <command> [options] [args]
```

```
Available commands:
```

```
  crawl      Start crawling a spider or URL
  fetch      Fetch a URL using the Scrapy downloader
  [...]     [...]
```

The first line will print the currently active project, if you're inside a Scrapy project. In this, it was run from outside a project. If run from inside a project it would have printed something like this:

```
Scrapy X.Y - project: myproject
```

```
Usage:
```

```
  scrapy <command> [options] [args]
```

```
[...]
```

Creating projects

The first thing you typically do with the `scrapy` tool is create your Scrapy project:

```
scrapy startproject myproject
```

That will create a Scrapy project under the `myproject` directory.

Next, you go inside the new project directory:

```
cd myproject
```

And you're ready to use the `scrapy` command to manage and control your project from there.

Controlling projects

You use the `scrapy` tool from inside your projects to control and manage them.

For example, to create a new spider:

```
scrapy genspider mydomain mydomain.com
```

Some Scrapy commands (like `crawl`) must be run from inside a Scrapy project. See the [commands reference](#) below for more information on which commands must be run from inside projects, and which not.

Also keep in mind that some commands may have slightly different behaviours when running them from inside projects. For example, the `fetch` command will use spider-overridden behaviours (such as the `user_agent` attribute to override the user-agent) if the url being fetched is associated with some specific spider. This is intentional, as the `fetch` command is meant to be used to check how spiders are downloading pages.

3.1.3 Available tool commands

This section contains a list of the available built-in commands with a description and some usage examples. Remember you can always get more info about each command by running:

```
scrapy <command> -h
```

And you can see all available commands with:

```
scrapy -h
```

There are two kinds of commands, those that only work from inside a Scrapy project (Project-specific commands) and those that also work without an active Scrapy project (Global commands), though they may behave slightly different when running from inside a project (as they would use the project overridden settings).

Global commands:

- `startproject`
- `settings`
- `runspider`
- `shell`
- `fetch`
- `view`
- `version`

Project-only commands:

- `crawl`
- `check`
- `list`
- `edit`
- `parse`
- `genspider`
- `server`
- `deploy`

startproject

- Syntax: `scrapy startproject <project_name>`
- Requires project: *no*

Creates a new Scrapy project named `project_name`, under the `project_name` directory.

Usage example:

```
$ scrapy startproject myproject
```

genspider

- Syntax: `scrapy genspider [-t template] <name> <domain>`
- Requires project: *yes*

Create a new spider in the current project.

This is just a convenient shortcut command for creating spiders based on pre-defined templates, but certainly not the only way to create spiders. You can just create the spider source code files yourself, instead of using this command.

Usage example:

```
$ scrapy genspider -l
Available templates:
  basic
  crawl
  csvfeed
  xmlfeed

$ scrapy genspider -d basic
from scrapy.spider import BaseSpider

class $classname(BaseSpider):
    name = "$name"
    allowed_domains = ["$domain"]
    start_urls = (
        'http://www.$domain/',
    )

    def parse(self, response):
        pass

$ scrapy genspider -t basic example example.com
Created spider 'example' using template 'basic' in module:
mybot.spiders.example
```

crawl

- Syntax: `scrapy crawl <spider>`
- Requires project: *yes*

Start crawling a spider.

Usage examples:

```
$ scrapy crawl myspider
[ ... myspider starts crawling ... ]
```

check

- Syntax: `scrapy check [-l] <spider>`
- Requires project: *yes*

Run contract checks.

Usage examples:

```
$ scrapy check -l
first_spider
  * parse
  * parse_item
second_spider
```



```
* parse
* parse_item

$ scrapy check
[FAILED] first_spider:parse_item
>>> 'RetailPricex' field is missing

[FAILED] first_spider:parse
>>> Returned 92 requests, expected 0..4
```

server

- Syntax: scrapy server
- Requires project: *yes*

Start Scrapy server for this project, which can be referred from the JSON API with the project name `default`. For more info see: [Scrapy Service \(scrapyd\)](#).

Usage example:

```
$ scrapy server
[ ... scrapyd starts and stays idle waiting for spiders to get scheduled ... ]
```

To schedule spiders, use the Scrapy JSON API.

list

- Syntax: scrapy list
- Requires project: *yes*

List all available spiders in the current project. The output is one spider per line.

Usage example:

```
$ scrapy list
spider1
spider2
```

edit

- Syntax: scrapy edit <spider>
- Requires project: *yes*

Edit the given spider using the editor defined in the `EDITOR` setting.

This command is provided only as a convenient shortcut for the most common case, the developer is of course free to choose any tool or IDE to write and debug his spiders.

Usage example:

```
$ scrapy edit spider1
```

fetch

- Syntax: `scrapy fetch <url>`
- Requires project: *no*

Downloads the given URL using the Scrapy downloader and writes the contents to standard output.

The interesting thing about this command is that it fetches the page how the spider would download it. For example, if the spider has an `USER_AGENT` attribute which overrides the User Agent, it will use that one.

So this command can be used to “see” how your spider would fetch certain page.

If used outside a project, no particular per-spider behaviour would be applied and it will just use the default Scrapy downloader settings.

Usage examples:

```
$ scrapy fetch --nolog http://www.example.com/some/page.html
[ ... html content here ... ]
```

```
$ scrapy fetch --nolog --headers http://www.example.com/
{'Accept-Ranges': ['bytes'],
 'Age': ['1263'],
 'Connection': ['close'],
 'Content-Length': ['596'],
 'Content-Type': ['text/html; charset=UTF-8'],
 'Date': ['Wed, 18 Aug 2010 23:59:46 GMT'],
 'Etag': ['"573c1-254-48c9c87349680"'],
 'Last-Modified': ['Fri, 30 Jul 2010 15:30:18 GMT'],
 'Server': ['Apache/2.2.3 (CentOS)']}
```

view

- Syntax: `scrapy view <url>`
- Requires project: *no*

Opens the given URL in a browser, as your Scrapy spider would “see” it. Sometimes spiders see pages differently from regular users, so this can be used to check what the spider “sees” and confirm it’s what you expect.

Usage example:

```
$ scrapy view http://www.example.com/some/page.html
[ ... browser starts ... ]
```

shell

- Syntax: `scrapy shell [url]`
- Requires project: *no*

Starts the Scrapy shell for the given URL (if given) or empty if not URL is given. See [Scrapy shell](#) for more info.

Usage example:

```
$ scrapy shell http://www.example.com/some/page.html
[ ... scrapy shell starts ... ]
```

parse

- Syntax: `scrapy parse <url> [options]`
- Requires project: *yes*

Fetches the given URL and parses with the spider that handles it, using the method passed with the `--callback` option, or `parse` if not given.

Supported options:

- `--callback` or `-c`: spider method to use as callback for parsing the response
- `--rules` or `-r`: use `CrawlSpider` rules to discover the callback (ie. spider method) to use for parsing the response
- `--noitems`: don't show scraped items
- `--nolinks`: don't show extracted links
- `--depth` or `-d`: depth level for which the requests should be followed recursively (default: 1)
- `--verbose` or `-v`: display information for each depth level

Usage example:

```
$ scrapy parse http://www.example.com/ -c parse_item
[ ... scrapy log lines crawling example.com spider ... ]

>>> STATUS DEPTH LEVEL 1 <<<
# Scraped Items -----
[{'name': u'Example item',
  'category': u'Furniture',
  'length': u'12 cm'}]

# Requests -----
[]
```

settings

- Syntax: `scrapy settings [options]`
- Requires project: *no*

Get the value of a Scrapy setting.

If used inside a project it'll show the project setting value, otherwise it'll show the default Scrapy value for that setting.

Example usage:

```
$ scrapy settings --get BOT_NAME
scrapybot
$ scrapy settings --get DOWNLOAD_DELAY
0
```

runspider

- Syntax: `scrapy runspider <spider_file.py>`
- Requires project: *no*

Run a spider self-contained in a Python file, without having to create a project.

Example usage:

```
$ scrapy runspider myspider.py
[ ... spider starts crawling ... ]
```

version

- Syntax: `scrapy version [-v]`
- Requires project: *no*

Prints the Scrapy version. If used with `-v` it also prints Python, Twisted and Platform info, which is useful for bug reports.

deploy

New in version 0.11.

- Syntax: `scrapy deploy [<target:project> | -l <target> | -L]`
- Requires project: *yes*

Deploy the project into a Scrapy server. See *Deploying your project*.

3.1.4 Custom project commands

You can also add your custom project commands by using the `COMMANDS_MODULE` setting. See the Scrapy commands in [scrapy/commands](#) for examples on how to implement your commands.

COMMANDS_MODULE

Default: "" (empty string)

A module to use for looking custom Scrapy commands. This is used to add custom commands for your Scrapy project.

Example:

```
COMMANDS_MODULE = 'mybot.commands'
```

3.2 Items

The main goal in scraping is to extract structured data from unstructured sources, typically, web pages. Scrapy provides the `Item` class for this purpose.

`Item` objects are simple containers used to collect the scraped data. They provide a *dictionary-like* API with a convenient syntax for declaring their available fields.

3.2.1 Declaring Items

Items are declared using a simple class definition syntax and `Field` objects. Here is an example:

```
from scrapy.item import Item, Field

class Product(Item):
    name = Field()
    price = Field()
    stock = Field()
    last_updated = Field(serializer=str)
```

Note: Those familiar with `Django` will notice that Scrapy Items are declared similar to `Django Models`, except that Scrapy Items are much simpler as there is no concept of different field types.

3.2.2 Item Fields

`Field` objects are used to specify metadata for each field. For example, the serializer function for the `last_updated` field illustrated in the example above.

You can specify any kind of metadata for each field. There is no restriction on the values accepted by `Field` objects. For this same reason, there isn't a reference list of all available metadata keys. Each key defined in `Field` objects could be used by a different components, and only those components know about it. You can also define and use any other `Field` key in your project too, for your own needs. The main goal of `Field` objects is to provide a way to define all field metadata in one place. Typically, those components whose behaviour depends on each field use certain field keys to configure that behaviour. You must refer to their documentation to see which metadata keys are used by each component.

It's important to note that the `Field` objects used to declare the item do not stay assigned as class attributes. Instead, they can be accessed through the `Item.fields` attribute.

And that's all you need to know about declaring items.

3.2.3 Working with Items

Here are some examples of common tasks performed with items, using the `Product` item *declared above*. You will notice the API is very similar to the `dict` API.

Creating items

```
>>> product = Product(name='Desktop PC', price=1000)
>>> print product
Product(name='Desktop PC', price=1000)
```

Getting field values

```
>>> product['name']
Desktop PC
>>> product.get('name')
Desktop PC

>>> product['price']
```

```
1000

>>> product['last_updated']
Traceback (most recent call last):
...
KeyError: 'last_updated'

>>> product.get('last_updated', 'not set')
not set

>>> product['lala'] # getting unknown field
Traceback (most recent call last):
...
KeyError: 'lala'

>>> product.get('lala', 'unknown field')
'unknown field'

>>> 'name' in product # is name field populated?
True

>>> 'last_updated' in product # is last_updated populated?
False

>>> 'last_updated' in product.fields # is last_updated a declared field?
True

>>> 'lala' in product.fields # is lala a declared field?
False
```

Setting field values

```
>>> product['last_updated'] = 'today'
>>> product['last_updated']
today

>>> product['lala'] = 'test' # setting unknown field
Traceback (most recent call last):
...
KeyError: 'Product does not support field: lala'
```

Accessing all populated values

To access all populated values, just use the typical `dict` API:

```
>>> product.keys()
['price', 'name']

>>> product.items()
[('price', 1000), ('name', 'Desktop PC')]
```

Other common tasks

Copying items:

```
>>> product2 = Product(product)
>>> print product2
Product(name='Desktop PC', price=1000)
```

Creating dicts from items:

```
>>> dict(product) # create a dict from all populated values
{'price': 1000, 'name': 'Desktop PC'}
```

Creating items from dicts:

```
>>> Product({'name': 'Laptop PC', 'price': 1500})
Product(price=1500, name='Laptop PC')

>>> Product({'name': 'Laptop PC', 'lala': 1500}) # warning: unknown field in dict
Traceback (most recent call last):
...
KeyError: 'Product does not support field: lala'
```

3.2.4 Extending Items

You can extend Items (to add more fields or to change some metadata for some fields) by declaring a subclass of your original Item.

For example:

```
class DiscountedProduct(Product):
    discount_percent = Field(serializer=str)
    discount_expiration_date = Field()
```

You can also extend field metadata by using the previous field metadata and appending more values, or changing existing values, like this:

```
class SpecificProduct(Product):
    name = Field(Product.fields['name'], serializer=my_serializer)
```

That adds (or replaces) the `serializer` metadata key for the `name` field, keeping all the previously existing metadata values.

3.2.5 Item objects

```
class scrapy.item.Item([arg])
```

Return a new Item optionally initialized from the given argument.

Items replicate the standard `dict` API, including its constructor. The only additional attribute provided by Items is:

fields

A dictionary containing *all declared fields* for this Item, not only those populated. The keys are the field names and the values are the `Field` objects used in the *Item declaration*.

3.2.6 Field objects

```
class scrapy.item.Field([arg])
```

The `Field` class is just an alias to the built-in `dict` class and doesn't provide any extra functionality or at-

tributes. In other words, `Field` objects are plain-old Python dicts. A separate class is used to support the *item declaration syntax* based on class attributes.

3.3 Spiders

Spiders are classes which define how a certain site (or group of sites) will be scraped, including how to perform the crawl (ie. follow links) and how to extract structured data from their pages (ie. scraping items). In other words, Spiders are the place where you define the custom behaviour for crawling and parsing pages for a particular site (or, in some cases, group of sites).

For spiders, the scraping cycle goes through something like this:

1. You start by generating the initial Requests to crawl the first URLs, and specify a callback function to be called with the response downloaded from those requests.

The first requests to perform are obtained by calling the `start_requests()` method which (by default) generates `Request` for the URLs specified in the `start_urls` and the `parse` method as callback function for the Requests.

2. In the callback function, you parse the response (web page) and return either `Item` objects, `Request` objects, or an iterable of both. Those Requests will also contain a callback (maybe the same) and will then be downloaded by Scrapy and then their response handled by the specified callback.
3. In callback functions, you parse the page contents, typically using *Selectors* (but you can also use BeautifulSoup, lxml or whatever mechanism you prefer) and generate items with the parsed data.
4. Finally, the items returned from the spider will be typically persisted to a database (in some *Item Pipeline*) or written to a file using *Feed exports*.

Even though this cycle applies (more or less) to any kind of spider, there are different kinds of default spiders bundled into Scrapy for different purposes. We will talk about those types here.

3.3.1 Spider arguments

Spiders can receive arguments that modify their behaviour. Some common uses for spider arguments are to define the start URLs or to restrict the crawl to certain sections of the site, but they can be used to configure any functionality of the spider.

Spider arguments are passed through the `crawl` command using the `-a` option. For example:

```
scrapy crawl myspider -a category=electronics
```

Spiders receive arguments in their constructors:

```
class MySpider(BaseSpider):
    name = 'myspider'

    def __init__(self, category=None):
        self.start_urls = ['http://www.example.com/categories/%s' % category]
        # ...
```

Spider arguments can also be passed through the *schedule.json* API.

3.3.2 Built-in spiders reference

Scrapy comes with some useful generic spiders that you can use, to subclass your spiders from. Their aim is to provide convenient functionality for a few common scraping cases, like following all links on a site based on certain rules, crawling from [Sitemaps](#), or parsing a XML/CSV feed.

For the examples used in the following spiders, we'll assume you have a project with a `TestItem` declared in a `myproject.items` module:

```
from scrapy.item import Item

class TestItem(Item):
    id = Field()
    name = Field()
    description = Field()
```

BaseSpider

class scrapy.spider.BaseSpider

This is the simplest spider, and the one from which every other spider must inherit from (either the ones that come bundled with Scrapy, or the ones that you write yourself). It doesn't provide any special functionality. It just requests the given `start_urls/start_requests`, and calls the spider's method `parse` for each of the resulting responses.

name

A string which defines the name for this spider. The spider name is how the spider is located (and instantiated) by Scrapy, so it must be unique. However, nothing prevents you from instantiating more than one instance of the same spider. This is the most important spider attribute and it's required.

If the spider scrapes a single domain, a common practice is to name the spider after the domain, or without the TLD. So, for example, a spider that crawls `mywebsite.com` would often be called `mywebsite`.

allowed_domains

An optional list of strings containing domains that this spider is allowed to crawl. Requests for URLs not belonging to the domain names specified in this list won't be followed if `OffsiteMiddleware` is enabled.

start_urls

A list of URLs where the spider will begin to crawl from, when no particular URLs are specified. So, the first pages downloaded will be those listed here. The subsequent URLs will be generated successively from data contained in the start URLs.

start_requests()

This method must return an iterable with the first Requests to crawl for this spider.

This is the method called by Scrapy when the spider is opened for scraping when no particular URLs are specified. If particular URLs are specified, the `make_requests_from_url()` is used instead to create the Requests. This method is also called only once from Scrapy, so it's safe to implement it as a generator.

The default implementation uses `make_requests_from_url()` to generate Requests for each url in `start_urls`.

If you want to change the Requests used to start scraping a domain, this is the method to override. For example, if you need to start by logging in using a POST request, you could do:

```
def start_requests(self):
    return [FormRequest("http://www.example.com/login",
                        formdata={'user': 'john', 'pass': 'secret'})]
```

```
        callback=self.logged_in)]

def logged_in(self, response):
    # here you would extract links to follow and return Requests for
    # each of them, with another callback
    pass

make_requests_from_url(url)
    A method that receives a URL and returns a Request object (or a list of Request objects) to scrape. This method is used to construct the initial requests in the start_requests() method, and is typically used to convert urls to requests.

    Unless overridden, this method returns Requests with the parse() method as their callback function, and with dont_filter parameter enabled (see Request class for more info).

parse(response)
    This is the default callback used by Scrapy to process downloaded responses, when their requests don't specify a callback.

    The parse method is in charge of processing the response and returning scraped data and/or more URLs to follow. Other Requests callbacks have the same requirements as the BaseSpider class.

    This method, as well as any other Request callback, must return an iterable of Request and/or Item objects.

    Parameters response (:class:`~scrapy.http.Response`) – the response to parse

log(message[, level, component])
    Log a message using the scrapy.log.msg() function, automatically populating the spider argument with the name of this spider. For more information see Logging.
```

BaseSpider example

Let's see an example:

```
from scrapy import log # This module is useful for printing out debug information
from scrapy.spider import BaseSpider

class MySpider(BaseSpider):
    name = 'example.com'
    allowed_domains = ['example.com']
    start_urls = [
        'http://www.example.com/1.html',
        'http://www.example.com/2.html',
        'http://www.example.com/3.html',
    ]

    def parse(self, response):
        self.log('A response from %s just arrived!' % response.url)
```

Another example returning multiples Requests and Items from a single callback:

```
from scrapy.selector import HtmlXPathSelector
from scrapy.spider import BaseSpider
from scrapy.http import Request
from myproject.items import MyItem

class MySpider(BaseSpider):
    name = 'example.com'
```

```

allowed_domains = ['example.com']
start_urls = [
    'http://www.example.com/1.html',
    'http://www.example.com/2.html',
    'http://www.example.com/3.html',
]

def parse(self, response):
    hxs = HtmlXPathSelector(response)
    for h3 in hxs.select('//h3').extract():
        yield MyItem(title=h3)

    for url in hxs.select('//a/@href').extract():
        yield Request(url, callback=self.parse)

```

CrawlSpider

class scrapy.contrib.spiders.CrawlSpider

This is the most commonly used spider for crawling regular websites, as it provides a convenient mechanism for following links by defining a set of rules. It may not be the best suited for your particular web sites or project, but it's generic enough for several cases, so you can start from it and override it as needed for more custom functionality, or just implement your own spider.

Apart from the attributes inherited from BaseSpider (that you must specify), this class supports a new attribute:

rules

Which is a list of one (or more) [Rule](#) objects. Each [Rule](#) defines a certain behaviour for crawling the site. Rules objects are described below. If multiple rules match the same link, the first one will be used, according to the order they're defined in this attribute.

Crawling rules

class scrapy.contrib.spiders.Rule(link_extractor, callback=None, cb_kwargs=None, follow=None, process_links=None, process_request=None)

`link_extractor` is a [Link Extractor](#) object which defines how links will be extracted from each crawled page.

`callback` is a callable or a string (in which case a method from the spider object with that name will be used) to be called for each link extracted with the specified `link_extractor`. This callback receives a response as its first argument and must return a list containing [Item](#) and/or [Request](#) objects (or any subclass of them).

Warning: When writing crawl spider rules, avoid using `parse` as callback, since the [CrawlSpider](#) uses the `parse` method itself to implement its logic. So if you override the `parse` method, the crawl spider will no longer work.

`cb_kwargs` is a dict containing the keyword arguments to be passed to the callback function

`follow` is a boolean which specifies if links should be followed from each response extracted with this rule. If `callback` is `None` `follow` defaults to `True`, otherwise it default to `False`.

`process_links` is a callable, or a string (in which case a method from the spider object with that name will be used) which will be called for each list of links extracted from each response using the specified `link_extractor`. This is mainly used for filtering purposes.

`process_request` is a callable, or a string (in which case a method from the spider object with that name will be used) which will be called with every request extracted by this rule, and must return a request or None (to filter out the request).

CrawlSpider example

Let's now take a look at an example CrawlSpider with rules:

```
from scrapy.contrib.spiders import CrawlSpider, Rule
from scrapy.contrib.linkextractors.sgml import SgmlLinkExtractor
from scrapy.selector import HtmlXPathSelector
from scrapy.item import Item

class MySpider(CrawlSpider):
    name = 'example.com'
    allowed_domains = ['example.com']
    start_urls = ['http://www.example.com']

    rules = (
        # Extract links matching 'category.php' (but not matching 'subsection.php')
        # and follow links from them (since no callback means follow=True by default).
        Rule(SgmlLinkExtractor(allow=('category\.php', ), deny=('subsection\.php', ))),

        # Extract links matching 'item.php' and parse them with the spider's method parse_item
        Rule(SgmlLinkExtractor(allow=('item\.php', ), callback='parse_item'),
    )

    def parse_item(self, response):
        self.log('Hi, this is an item page! %s' % response.url)

        hxs = HtmlXPathSelector(response)
        item = Item()
        item['id'] = hxs.select('//td[@id="item_id"]/text()').re(r'ID: (\d+)')
        item['name'] = hxs.select('//td[@id="item_name"]/text()').extract()
        item['description'] = hxs.select('//td[@id="item_description"]/text()').extract()
        return item
```

This spider would start crawling `example.com`'s home page, collecting category links, and item links, parsing the latter with the `parse_item` method. For each item response, some data will be extracted from the HTML using XPath, and a `Item` will be filled with it.

XMLFeedSpider

class scrapy.contrib.spiders.XMLFeedSpider

XMLFeedSpider is designed for parsing XML feeds by iterating through them by a certain node name. The iterator can be chosen from: `iternodes`, `xml`, and `html`. It's recommended to use the `iternodes` iterator for performance reasons, since the `xml` and `html` iterators generate the whole DOM at once in order to parse it. However, using `html` as the iterator may be useful when parsing XML with bad markup.

To set the iterator and the tag name, you must define the following class attributes:

iterator

A string which defines the iterator to use. It can be either:

- `'iternodes'` - a fast iterator based on regular expressions

- `'html'` - an iterator which uses `HtmlXPathSelector`. Keep in mind this uses DOM parsing and must load all DOM in memory which could be a problem for big feeds
- `'xml'` - an iterator which uses `XmlXPathSelector`. Keep in mind this uses DOM parsing and must load all DOM in memory which could be a problem for big feeds

It defaults to: `'iternodes'`.

itertag

A string with the name of the node (or element) to iterate in. Example:

```
itertag = 'product'
```

namespaces

A list of (prefix, uri) tuples which define the namespaces available in that document that will be processed with this spider. The prefix and uri will be used to automatically register namespaces using the `register_namespace()` method.

You can then specify nodes with namespaces in the `itertag` attribute.

Example:

```
class YourSpider(XMLFeedSpider):

    namespaces = [('n', 'http://www.sitemaps.org/schemas/sitemap/0.9')]
    itertag = 'n:url'
    # ...
```

Apart from these new attributes, this spider has the following overrideable methods too:

adapt_response (*response*)

A method that receives the response as soon as it arrives from the spider middleware, before the spider starts parsing it. It can be used to modify the response body before parsing it. This method receives a response and also returns a response (it could be the same or another one).

parse_node (*response*, *selector*)

This method is called for the nodes matching the provided tag name (`itertag`). Receives the response and an `XPathSelector` for each node. Overriding this method is mandatory. Otherwise, your spider won't work. This method must return either a `Item` object, a `Request` object, or an iterable containing any of them.

process_results (*response*, *results*)

This method is called for each result (item or request) returned by the spider, and it's intended to perform any last time processing required before returning the results to the framework core, for example setting the item IDs. It receives a list of results and the response which originated those results. It must return a list of results (Items or Requests).

XMLFeedSpider example

These spiders are pretty easy to use, let's have a look at one example:

```
from scrapy import log
from scrapy.contrib.spiders import XMLFeedSpider
from myproject.items import TestItem

class MySpider(XMLFeedSpider):
    name = 'example.com'
    allowed_domains = ['example.com']
    start_urls = ['http://www.example.com/feed.xml']
```

```
iterator = 'iternodes' # This is actually unnecessary, since it's the default value
itertag = 'item'

def parse_node(self, response, node):
    log.msg('Hi, this is a <%s> node!: %s' % (self.itertag, ''.join(node.extract())))

    item = Item()
    item['id'] = node.select('@id').extract()
    item['name'] = node.select('name').extract()
    item['description'] = node.select('description').extract()
    return item
```

Basically what we did up there was to create a spider that downloads a feed from the given `start_urls`, and then iterates through each of its `item` tags, prints them out, and stores some random data in an `Item`.

CSVFeedSpider

class scrapy.contrib.spiders.CSVFeedSpider

This spider is very similar to the `XMLFeedSpider`, except that it iterates over rows, instead of nodes. The method that gets called in each iteration is `parse_row()`.

delimiter

A string with the separator character for each field in the CSV file Defaults to `,` (comma).

headers

A list of the rows contained in the file CSV feed which will be used to extract fields from it.

parse_row(response, row)

Receives a response and a dict (representing each row) with a key for each provided (or detected) header of the CSV file. This spider also gives the opportunity to override `adapt_response` and `process_results` methods for pre- and post-processing purposes.

CSVFeedSpider example

Let's see an example similar to the previous one, but using a `CSVFeedSpider`:

```
from scrapy import log
from scrapy.contrib.spiders import CSVFeedSpider
from myproject.items import TestItem

class MySpider(CSVFeedSpider):
    name = 'example.com'
    allowed_domains = ['example.com']
    start_urls = ['http://www.example.com/feed.csv']
    delimiter = ';'
    headers = ['id', 'name', 'description']

    def parse_row(self, response, row):
        log.msg('Hi, this is a row!: %r' % row)

        item = TestItem()
        item['id'] = row['id']
        item['name'] = row['name']
        item['description'] = row['description']
        return item
```

SitemapSpider

class scrapy.contrib.spiders.SitemapSpider

SitemapSpider allows you to crawl a site by discovering the URLs using [Sitemaps](#).

It supports nested sitemaps and discovering sitemap urls from [robots.txt](#).

sitemap_urls

A list of urls pointing to the sitemaps whose urls you want to crawl.

You can also point to a [robots.txt](#) and it will be parsed to extract sitemap urls from it.

sitemap_rules

A list of tuples (regex, callback) where:

- regex is a regular expression to match urls extracted from sitemaps. regex can be either a str or a compiled regex object.
- callback is the callback to use for processing the urls that match the regular expression. callback can be a string (indicating the name of a spider method) or a callable.

For example:

```
sitemap_rules = [('/product/', 'parse_product')]
```

Rules are applied in order, and only the first one that matches will be used.

If you omit this attribute, all urls found in sitemaps will be processed with the parse callback.

sitemap_follow

A list of regexes of sitemap that should be followed. This is only for sites that use [Sitemap index files](#) that point to other sitemap files.

By default, all sitemaps are followed.

SitemapSpider examples

Simplest example: process all urls discovered through sitemaps using the parse callback:

```
from scrapy.contrib.spiders import SitemapSpider

class MySpider(SitemapSpider):
    sitemap_urls = ['http://www.example.com/sitemap.xml']

    def parse(self, response):
        pass # ... scrape item here ...
```

Process some urls with certain callback and other urls with a different callback:

```
from scrapy.contrib.spiders import SitemapSpider

class MySpider(SitemapSpider):
    sitemap_urls = ['http://www.example.com/sitemap.xml']
    sitemap_rules = [
        ('/product/', 'parse_product'),
        ('/category/', 'parse_category'),
    ]

    def parse_product(self, response):
        pass # ... scrape product ...
```

```
def parse_category(self, response):  
    pass # ... scrape category ...
```

Follow sitemaps defined in the `robots.txt` file and only follow sitemaps whose url contains `/sitemap_shop`:

```
from scrapy.contrib.spiders import SitemapSpider  
  
class MySpider(SitemapSpider):  
    sitemap_urls = ['http://www.example.com/robots.txt']  
    sitemap_rules = [  
        ('/shop/', 'parse_shop'),  
    ]  
    sitemap_follow = ['/sitemap_shops']  
  
    def parse_shop(self, response):  
        pass # ... scrape shop here ...
```

Combine SitemapSpider with other sources of urls:

```
from scrapy.contrib.spiders import SitemapSpider  
  
class MySpider(SitemapSpider):  
    sitemap_urls = ['http://www.example.com/robots.txt']  
    sitemap_rules = [  
        ('/shop/', 'parse_shop'),  
    ]  
  
    other_urls = ['http://www.example.com/about']  
  
    def start_requests(self):  
        requests = list(super(MySpider, self).start_requests())  
        requests += [Request(x, callback=self.parse_other) for x in self.other_urls]  
        return requests  
  
    def parse_shop(self, response):  
        pass # ... scrape shop here ...  
  
    def parse_other(self, response):  
        pass # ... scrape other here ...
```

3.4 Link Extractors

LinkExtractors are objects whose only purpose is to extract links from web pages (`scrapy.http.Response` objects) which will be eventually followed.

There are two Link Extractors available in Scrapy by default, but you create your own custom Link Extractors to suit your needs by implementing a simple interface.

The only public method that every LinkExtractor has is `extract_links`, which receives a `Response` object and returns a list of links. Link Extractors are meant to be instantiated once and their `extract_links` method called several times with different responses, to extract links to follow.

Link extractors are used in the `CrawlSpider` class (available in Scrapy), through a set of rules, but you can also use it in your spiders, even if you don't subclass from `CrawlSpider`, as its purpose is very simple: to extract links.

3.4.1 Built-in link extractors reference

All available link extractors classes bundled with Scrapy are provided in the `scrapy.contrib.linkextractors` module.

SgmlLinkExtractor

```
class scrapy.contrib.linkextractors.sgml.SgmlLinkExtractor(allow=(), deny=(),
                                                         allow_domains=(),
                                                         deny_domains=(),
                                                         deny_extensions=None,
                                                         restrict_xpaths=(),
                                                         tags=('a', 'area'),
                                                         attrs=('href'), canonicalize=True,
                                                         unique=True, process_value=None)
```

The `SgmlLinkExtractor` extends the base `BaseSgmlLinkExtractor` by providing additional filters that you can specify to extract links, including regular expressions patterns that the links must match to be extracted. All those filters are configured through these constructor parameters:

Parameters

- **allow** (*a regular expression (or list of)*) – a single regular expression (or list of regular expressions) that the (absolute) urls must match in order to be extracted. If not given (or empty), it will match all links.
- **deny** (*a regular expression (or list of)*) – a single regular expression (or list of regular expressions) that the (absolute) urls must match in order to be excluded (ie. not extracted). It has precedence over the `allow` parameter. If not given (or empty) it won't exclude any links.
- **allow_domains** (*str or list*) – a single value or a list of string containing domains which will be considered for extracting the links
- **deny_domains** (*str or list*) – a single value or a list of strings containing domains which won't be considered for extracting the links
- **deny_extensions** (*list*) – a list of extensions that should be ignored when extracting links. If not given, it will default to the `IGNORED_EXTENSIONS` list defined in the `scrapy.linkextractor` module.
- **restrict_xpaths** (*str or list*) – is a XPath (or list of XPath's) which defines regions inside the response where links should be extracted from. If given, only the text selected by those XPath will be scanned for links. See examples below.
- **tags** (*str or list*) – a tag or a list of tags to consider when extracting links. Defaults to `('a', 'area')`.
- **attrs** (*list*) – list of attributes which should be considered when looking for links to extract (only for those tags specified in the `tags` parameter). Defaults to `('href',)`
- **canonicalize** (*boolean*) – canonicalize each extracted url (using `scrapy.utils.url.canonicalize_url`). Defaults to `True`.
- **unique** (*boolean*) – whether duplicate filtering should be applied to extracted links.
- **process_value** (*callable*) – see `process_value` argument of `BaseSgmlLinkExtractor` class constructor

BaseSgmlLinkExtractor

```
class scrapy.contrib.linkextractors.sgml.BaseSgmlLinkExtractor(tag="a",
                                                                attr="href",
                                                                unique=False, process_value=None)
```

The purpose of this Link Extractor is only to serve as a base class for the `SgmlLinkExtractor`. You should use that one instead.

The constructor arguments are:

Parameters

- **tag** (*str or callable*) – either a string (with the name of a tag) or a function that receives a tag name and returns `True` if links should be extracted from that tag, or `False` if they shouldn't. Defaults to `'a'`. request (once it's downloaded) as its first parameter. For more information, see *Passing additional data to callback functions*.
- **attr** (*str or callable*) – either string (with the name of a tag attribute), or a function that receives an attribute name and returns `True` if links should be extracted from it, or `False` if they shouldn't. Defaults to `href`.
- **unique** (*boolean*) – is a boolean that specifies if a duplicate filtering should be applied to links extracted.
- **process_value** (*callable*) – a function which receives each value extracted from the tag and attributes scanned and can modify the value and return a new one, or return `None` to ignore the link altogether. If not given, `process_value` defaults to `lambda x: x`.

For example, to extract links from this code:

```
<a href="javascript:goToPage('../other/page.html'); return false">Link text</a>
```

You can use the following function in `process_value`:

```
def process_value(value):
    m = re.search("javascript:goToPage\('(.*?)'", value)
    if m:
        return m.group(1)
```

3.5 Selectors

When you're scraping web pages, the most common task you need to perform is to extract data from the HTML source. There are several libraries available to achieve this:

- `BeautifulSoup` is a very popular screen scraping library among Python programmers which constructs a Python object based on the structure of the HTML code and also deals with bad markup reasonably well, but it has one drawback: it's slow.
- `lxml` is a XML parsing library (which also parses HTML) with a pythonic API based on `ElementTree` (which is not part of the Python standard library).

Scrapy comes with its own mechanism for extracting data. They're called XPath selectors (or just "selectors", for short) because they "select" certain parts of the HTML document specified by `XPath` expressions.

`XPath` is a language for selecting nodes in XML documents, which can also be used with HTML.

Both `lxml` and Scrapy Selectors are built over the `libxml2` library, which means they're very similar in speed and parsing accuracy.

This page explains how selectors work and describes their API which is very small and simple, unlike the `lxml` API which is much bigger because the `lxml` library can be used for many other tasks, besides selecting markup documents.

For a complete reference of the selectors API see the *XPath selector reference*.

3.5.1 Using selectors

Constructing selectors

There are two types of selectors bundled with Scrapy. Those are:

- `HtmlXPathSelector` - for working with HTML documents
- `XmlXPathSelector` - for working with XML documents

Both share the same selector API, and are constructed with a `Response` object as their first parameter. This is the `Response` they're going to be "selecting".

Example:

```
hxs = HtmlXPathSelector(response) # a HTML selector
xbs = XmlXPathSelector(response) # a XML selector
```

Using selectors with XPaths

To explain how to use the selectors we'll use the *Scrapy shell* (which provides interactive testing) and an example page located in the Scrapy documentation server:

http://doc.scrapy.org/en/latest/_static/selectors-sample1.html

Here's its HTML code:

```
<html>
<head>
  <base href='http://example.com/' />
  <title>Example website</title>
</head>
<body>
  <div id='images'>
    <a href='image1.html'>Name: My image 1 <br /><img src='image1_thumb.jpg' /></a>
    <a href='image2.html'>Name: My image 2 <br /><img src='image2_thumb.jpg' /></a>
    <a href='image3.html'>Name: My image 3 <br /><img src='image3_thumb.jpg' /></a>
    <a href='image4.html'>Name: My image 4 <br /><img src='image4_thumb.jpg' /></a>
    <a href='image5.html'>Name: My image 5 <br /><img src='image5_thumb.jpg' /></a>
  </div>
</body>
</html>
```

First, let's open the shell:

```
scrapy shell http://doc.scrapy.org/en/latest/_static/selectors-sample1.html
```

Then, after the shell loads, you'll have some selectors already instantiated and ready to use.

Since we're dealing with HTML, we'll be using the `HtmlXPathSelector` object which is found, by default, in the `hxs` shell variable.

So, by looking at the *HTML code* of that page, let's construct an XPath (using an HTML selector) for selecting the text inside the title tag:

```
>>> hxs.select('//title/text()')
[<HtmlXPathSelector (text) xpath=//title/text(>]
```

As you can see, the `select()` method returns an `XPathSelectorList`, which is a list of new selectors. This API can be used quickly for extracting nested data.

To actually extract the textual data, you must call the selector `extract()` method, as follows:

```
>>> hxs.select('//title/text()').extract()
[u'Example website']
```

Now we're going to get the base URL and some image links:

```
>>> hxs.select('//base/@href').extract()
[u'http://example.com/']

>>> hxs.select('//a[contains(@href, "image")]/@href').extract()
[u'image1.html',
 u'image2.html',
 u'image3.html',
 u'image4.html',
 u'image5.html']

>>> hxs.select('//a[contains(@href, "image")]/img/@src').extract()
[u'image1_thumb.jpg',
 u'image2_thumb.jpg',
 u'image3_thumb.jpg',
 u'image4_thumb.jpg',
 u'image5_thumb.jpg']
```

Using selectors with regular expressions

Selectors also have a `re()` method for extracting data using regular expressions. However, unlike using the `select()` method, the `re()` method does not return a list of `XPathSelector` objects, so you can't construct nested `.re()` calls.

Here's an example used to extract images names from the *HTML code* above:

```
>>> hxs.select('//a[contains(@href, "image")]/text()').re(r'Name:\s*(.*)')
[u'My image 1',
 u'My image 2',
 u'My image 3',
 u'My image 4',
 u'My image 5']
```

Nesting selectors

The `select()` selector method returns a list of selectors, so you can call the `select()` for those selectors too. Here's an example:

```
>>> links = hxs.select('//a[contains(@href, "image")]')
>>> links.extract()
[u'<a href="image1.html">Name: My image 1 <br></a>',
 u'<a href="image2.html">Name: My image 2 <br></a>',
```

```
u'<a href="image3.html">Name: My image 3 <br></a>',
u'<a href="image4.html">Name: My image 4 <br></a>',
u'<a href="image5.html">Name: My image 5 <br></a>']
```

```
>>> for index, link in enumerate(links):
    args = (index, link.select('@href').extract(), link.select('img/@src').extract())
    print 'Link number %d points to url %s and image %s' % args
```

```
Link number 0 points to url [u'image1.html'] and image [u'image1_thumb.jpg']
Link number 1 points to url [u'image2.html'] and image [u'image2_thumb.jpg']
Link number 2 points to url [u'image3.html'] and image [u'image3_thumb.jpg']
Link number 3 points to url [u'image4.html'] and image [u'image4_thumb.jpg']
Link number 4 points to url [u'image5.html'] and image [u'image5_thumb.jpg']
```

Working with relative XPaths

Keep in mind that if you are nesting XPathSelectors and use an XPath that starts with /, that XPath will be absolute to the document and not relative to the XPathSelector you're calling it from.

For example, suppose you want to extract all <p> elements inside <div> elements. First, you would get all <div> elements:

```
>>> divs = hxs.select('//div')
```

At first, you may be tempted to use the following approach, which is wrong, as it actually extracts all <p> elements from the document, not only those inside <div> elements:

```
>>> for p in divs.select('/p') # this is wrong - gets all <p> from the whole document
>>>     print p.extract()
```

This is the proper way to do it (note the dot prefixing the ./p XPath):

```
>>> for p in divs.select('./p') # extracts all <p> inside
>>>     print p.extract()
```

Another common case would be to extract all direct <p> children:

```
>>> for p in divs.select('p')
>>>     print p.extract()
```

For more details about relative XPaths see the [Location Paths](#) section in the XPath specification.

3.5.2 Built-in XPath Selectors reference

There are two types of selectors bundled with Scrapy: `HtmlXPathSelector` and `XmlXPathSelector`. Both of them implement the same `XPathSelector` interface. The only different is that one is used to process HTML data and the other XML data.

XPathSelector objects

class scrapy.selector.XPathSelector(response)

A `XPathSelector` object is a wrapper over response to select certain parts of its content.

response is a `Response` object that will be used for selecting and extracting data

select (*xpath*)

Apply the given XPath relative to this `XPathSelector` and return a list of `XPathSelector` objects (ie. a `XPathSelectorList`) with the result.

`xpath` is a string containing the XPath to apply

re (*regex*)

Apply the given regex and return a list of unicode strings with the matches.

`regex` can be either a compiled regular expression or a string which will be compiled to a regular expression using `re.compile(regex)`

extract ()

Return a unicode string with the content of this `XPathSelector` object.

register_namespace (*prefix, uri*)

Register the given namespace to be used in this `XPathSelector`. Without registering namespaces you can't select or extract data from non-standard namespaces. See examples below.

__nonzero__ ()

Returns `True` if there is any real content selected by this `XPathSelector` or `False` otherwise. In other words, the boolean value of an `XPathSelector` is given by the contents it selects.

XPathSelectorList objects

class scrapy.selector.**XPathSelectorList**

The `XPathSelectorList` class is subclass of the builtin `list` class, which provides a few additional methods.

select (*xpath*)

Call the `XPathSelector.select()` method for all `XPathSelector` objects in this list and return their results flattened, as a new `XPathSelectorList`.

`xpath` is the same argument as the one in `XPathSelector.select()`

re (*regex*)

Call the `XPathSelector.re()` method for all `XPathSelector` objects in this list and return their results flattened, as a list of unicode strings.

`regex` is the same argument as the one in `XPathSelector.re()`

extract ()

Call the `XPathSelector.extract()` method for all `XPathSelector` objects in this list and return their results flattened, as a list of unicode strings.

extract_unquoted ()

Call the `XPathSelector.extract_unquoted()` method for all `XPathSelector` objects in this list and return their results flattened, as a list of unicode strings. This method should not be applied to all kinds of `XPathSelectors`. For more info see `XPathSelector.extract_unquoted()`.

HtmlXPathSelector objects

class scrapy.selector.**HtmlXPathSelector** (*response*)

A subclass of `XPathSelector` for working with HTML content. It uses the `libxml2` HTML parser. See the `XPathSelector` API for more info.

HtmlXPathSelector examples

Here's a couple of `HtmlXPathSelector` examples to illustrate several concepts. In all cases, we assume there is already an `HtmlPathSelector` instantiated with a `Response` object like this:

```
x = HtmlXPathSelector(html_response)
```

1. Select all `<h1>` elements from a HTML response body, returning a list of `XPathSelector` objects (ie. a `XPathSelectorList` object):

```
x.select("//h1")
```

2. Extract the text of all `<h1>` elements from a HTML response body, returning a list of unicode strings:

```
x.select("//h1").extract()           # this includes the h1 tag
x.select("//h1/text()").extract()    # this excludes the h1 tag
```

3. Iterate over all `<p>` tags and print their class attribute:

```
for node in x.select("//p"):
    ...     print node.select("@href")
```

4. Extract textual data from all `<p>` tags without entities, as a list of unicode strings:

```
x.select("//p/text()").extract_unquoted()

# the following line is wrong. extract_unquoted() should only be used
# with textual XPathSelectors
x.select("//p").extract_unquoted()  # it may work but output is unpredictable
```

XmlXPathSelector objects

class scrapy.selector.**XmlXPathSelector**(response)

A subclass of `XPathSelector` for working with XML content. It uses the `libxml2` XML parser. See the `XPathSelector` API for more info.

XmlXPathSelector examples

Here's a couple of `XmlXPathSelector` examples to illustrate several concepts. In both cases we assume there is already an `XmlPathSelector` instantiated with a `Response` object like this:

```
x = XmlPathSelector(xml_response)
```

1. Select all `<product>` elements from a XML response body, returning a list of `XPathSelector` objects (ie. a `XPathSelectorList` object):

```
x.select("//product")
```

2. Extract all prices from a [Google Base XML feed](#) which requires registering a namespace:

```
x.register_namespace("g", "http://base.google.com/ns/1.0")
x.select("//g:price").extract()
```

3.6 Item Loaders

Item Loaders provide a convenient mechanism for populating scraped *Items*. Even though Items can be populated using their own dictionary-like API, the Item Loaders provide a much more convenient API for populating them from a scraping process, by automating some common tasks like parsing the raw extracted data before assigning it.

In other words, *Items* provide the *container* of scraped data, while Item Loaders provide the mechanism for *populating* that container.

Item Loaders are designed to provide a flexible, efficient and easy mechanism for extending and overriding different field parsing rules, either by spider, or by source format (HTML, XML, etc) without becoming a nightmare to maintain.

3.6.1 Using Item Loaders to populate items

To use an Item Loader, you must first instantiate it. You can either instantiate it with an dict-like object (e.g. Item or dict) or without one, in which case an Item is automatically instantiated in the Item Loader constructor using the Item class specified in the `ItemLoader.default_item_class` attribute.

Then, you start collecting values into the Item Loader, typically using *XPath Selectors*. You can add more than one value to the same item field; the Item Loader will know how to “join” those values later using a proper processing function.

Here is a typical Item Loader usage in a *Spider*, using the *Product item* declared in the *Items chapter*:

```
from scrapy.contrib.loader import XPathItemLoader
from myproject.items import Product

def parse(self, response):
    l = XPathItemLoader(item=Product(), response=response)
    l.add_xpath('name', '//div[@class="product_name"]')
    l.add_xpath('name', '//div[@class="product_title"]')
    l.add_xpath('price', '//p[@id="price"]')
    l.add_xpath('stock', '//p[@id="stock"]')
    l.add_value('last_updated', 'today') # you can also use literal values
    return l.load_item()
```

By quickly looking at that code, we can see the name field is being extracted from two different XPath locations in the page:

1. `//div[@class="product_name"]`
2. `//div[@class="product_title"]`

In other words, data is being collected by extracting it from two XPath locations, using the `add_xpath()` method. This is the data that will be assigned to the name field later.

Afterwards, similar calls are used for price and stock fields, and finally the last_update field is populated directly with a literal value (today) using a different method: `add_value()`.

Finally, when all data is collected, the `ItemLoader.load_item()` method is called which actually populates and returns the item populated with the data previously extracted and collected with the `add_xpath()` and `add_value()` calls.

3.6.2 Input and Output processors

An Item Loader contains one input processor and one output processor for each (item) field. The input processor processes the extracted data as soon as it's received (through the `add_xpath()` or `add_value()` methods) and the result of the input processor is collected and kept inside the ItemLoader. After collecting all data, the

`ItemLoader.load_item()` method is called to populate and get the populated `Item` object. That's when the output processor is called with the data previously collected (and processed using the input processor). The result of the output processor is the final value that gets assigned to the item.

Let's see an example to illustrate how the input and output processors are called for a particular field (the same applies for any other field):

```
l = XPathItemLoader(Product(), some_xpath_selector)
l.add_xpath('name', xpath1) # (1)
l.add_xpath('name', xpath2) # (2)
l.add_value('name', 'test') # (3)
return l.load_item() # (4)
```

So what happens is:

1. Data from `xpath1` is extracted, and passed through the *input processor* of the `name` field. The result of the input processor is collected and kept in the Item Loader (but not yet assigned to the item).
2. Data from `xpath2` is extracted, and passed through the same *input processor* used in (1). The result of the input processor is appended to the data collected in (1) (if any).
3. This case is similar to the previous ones, except that the value to be collected is assigned directly, instead of being extracted from a XPath. However, the value is still passed through the input processors. In this case, since the value is not iterable it is converted to an iterable of a single element before passing it to the input processor, because input processor always receive iterables.
4. The data collected in (1) and (2) is passed through the *output processor* of the `name` field. The result of the output processor is the value assigned to the `name` field in the item.

It's worth noticing that processors are just callable objects, which are called with the data to be parsed, and return a parsed value. So you can use any function as input or output processor. The only requirement is that they must accept one (and only one) positional argument, which will be an iterator.

Note: Both input and output processors must receive an iterator as their first argument. The output of those functions can be anything. The result of input processors will be appended to an internal list (in the Loader) containing the collected values (for that field). The result of the output processors is the value that will be finally assigned to the item.

The other thing you need to keep in mind is that the values returned by input processors are collected internally (in lists) and then passed to output processors to populate the fields.

Last, but not least, Scrapy comes with some *commonly used processors* built-in for convenience.

3.6.3 Declaring Item Loaders

Item Loaders are declared like Items, by using a class definition syntax. Here is an example:

```
from scrapy.contrib.loader import ItemLoader
from scrapy.contrib.loader.processor import TakeFirst, MapCompose, Join

class ProductLoader(ItemLoader):

    default_output_processor = TakeFirst()

    name_in = MapCompose(unicode.title)
    name_out = Join()

    price_in = MapCompose(unicode.strip)
```

```
# ...
```

As you can see, input processors are declared using the `_in` suffix while output processors are declared using the `_out` suffix. And you can also declare a default input/output processors using the `ItemLoader.default_input_processor` and `ItemLoader.default_output_processor` attributes.

3.6.4 Declaring Input and Output Processors

As seen in the previous section, input and output processors can be declared in the Item Loader definition, and it's very common to declare input processors this way. However, there is one more place where you can specify the input and output processors to use: in the *Item Field* metadata. Here is an example:

```
from scrapy.item import Item, Field
from scrapy.contrib.loader.processor import MapCompose, Join, TakeFirst

from scrapy.utils.markup import remove_entities
from myproject.utils import filter_prices

class Product(Item):
    name = Field(
        input_processor=MapCompose(remove_entities),
        output_processor=Join(),
    )
    price = Field(
        default=0,
        input_processor=MapCompose(remove_entities, filter_prices),
        output_processor=TakeFirst(),
    )
```

The precedence order, for both input and output processors, is as follows:

1. Item Loader field-specific attributes: `field_in` and `field_out` (most precedence)
2. Field metadata (`input_processor` and `output_processor` key)
3. Item Loader defaults: `ItemLoader.default_input_processor()` and `ItemLoader.default_output_processor()` (least precedence)

See also: *Reusing and extending Item Loaders*.

3.6.5 Item Loader Context

The Item Loader Context is a dict of arbitrary key/values which is shared among all input and output processors in the Item Loader. It can be passed when declaring, instantiating or using Item Loader. They are used to modify the behaviour of the input/output processors.

For example, suppose you have a function `parse_length` which receives a text value and extracts a length from it:

```
def parse_length(text, loader_context):
    unit = loader_context.get('unit', 'm')
    # ... length parsing code goes here ...
    return parsed_length
```

By accepting a `loader_context` argument the function is explicitly telling the Item Loader that is able to receive an Item Loader context, so the Item Loader passes the currently active context when calling it, and the processor function (`parse_length` in this case) can thus use them.

There are several ways to modify Item Loader context values:

1. By modifying the currently active Item Loader context (`context` attribute):

```
loader = ItemLoader(product)
loader.context['unit'] = 'cm'
```

2. On Item Loader instantiation (the keyword arguments of Item Loader constructor are stored in the Item Loader context):

```
loader = ItemLoader(product, unit='cm')
```

3. On Item Loader declaration, for those input/output processors that support instantiating them with a Item Loader context. `MapCompose` is one of them:

```
class ProductLoader(ItemLoader):
    length_out = MapCompose(parse_length, unit='cm')
```

3.6.6 ItemLoader objects

class scrapy.contrib.loader.ItemLoader([item], **kwargs)

Return a new Item Loader for populating the given Item. If no item is given, one is instantiated automatically using the class in `default_item_class`.

The item and the remaining keyword arguments are assigned to the Loader context (accessible through the `context` attribute).

get_value (value, *processors, **kwargs)

Process the given value by the given processors and keyword arguments.

Available keyword arguments:

Parameters `re` (*str or compiled regex*) – a regular expression to use for extracting data from the given value using `extract_regex()` method, applied before processors

Examples:

```
>>> from scrapy.contrib.loader.processor import TakeFirst
>>> loader.get_value(u'name: foo', TakeFirst(), unicode.upper, re='name: (.+)')
'FOO'
```

add_value (field_name, value, *processors, **kwargs)

Process and then add the given value for the given field.

The value is first passed through `get_value()` by giving the processors and kwargs, and then passed through the *field input processor* and its result appended to the data collected for that field. If the field already contains collected data, the new data is added.

The given `field_name` can be None, in which case values for multiple fields may be added. And the processed value should be a dict with `field_name` mapped to values.

Examples:

```
loader.add_value('name', u'Color TV')
loader.add_value('colours', [u'white', u'blue'])
loader.add_value('length', u'100')
loader.add_value('name', u'name: foo', TakeFirst(), re='name: (.+)')
loader.add_value(None, {'name': u'foo', 'sex': u'male'})
```

replace_value (field_name, value)

Similar to `add_value()` but replaces the collected data with the new value instead of adding it.

load_item()

Populate the item with the data collected so far, and return it. The data collected is first passed through the *output processors* to get the final value to assign to each item field.

get_collected_values (*field_name*)

Return the collected values for the given field.

get_output_value (*field_name*)

Return the collected values parsed using the output processor, for the given field. This method doesn't populate or modify the item at all.

get_input_processor (*field_name*)

Return the input processor for the given field.

get_output_processor (*field_name*)

Return the output processor for the given field.

item

The *Item* object being parsed by this Item Loader.

context

The currently active *Context* of this Item Loader.

default_item_class

An Item class (or factory), used to instantiate items when not given in the constructor.

default_input_processor

The default input processor to use for those fields which don't specify one.

default_output_processor

The default output processor to use for those fields which don't specify one.

class scrapy.contrib.loader.XPathItemLoader ([*item*, *selector*, *response*], ****kwargs**)

The *XPathItemLoader* class extends the *ItemLoader* class providing more convenient mechanisms for extracting data from web pages using *XPath selectors*.

XPathItemLoader objects accept two more additional parameters in their constructors:

Parameters

- **selector** (*XPathSelector* object) – The selector to extract data from, when using the *add_xpath()* or *replace_xpath()* method.
- **response** (*Response* object) – The response used to construct the selector using the *default_selector_class*, unless the selector argument is given, in which case this argument is ignored.

get_xpath (*xpath*, **processors*, ****kwargs**)

Similar to *ItemLoader.get_value()* but receives an XPath instead of a value, which is used to extract a list of unicode strings from the selector associated with this *XPathItemLoader*.

Parameters

- **xpath** (*str*) – the XPath to extract data from
- **re** (*str* or *compiled regex*) – a regular expression to use for extracting data from the selected XPath region

Examples:

```
# HTML snippet: <p class="product-name">Color TV</p>
loader.get_xpath('//p[@class="product-name"]')
# HTML snippet: <p id="price">the price is $1200</p>
loader.get_xpath('//p[@id="price"]', TakeFirst(), re='the price is (.*)')
```

add_xpath (*field_name*, *xpath*, **processors*, ***kwargs*)

Similar to `ItemLoader.add_value()` but receives an XPath instead of a value, which is used to extract a list of unicode strings from the selector associated with this `XPathItemLoader`.

See `get_xpath()` for kwargs.

Parameters *xpath* (*str*) – the XPath to extract data from

Examples:

```
# HTML snippet: <p class="product-name">Color TV</p>
loader.add_xpath('name', '//p[@class="product-name"]')
# HTML snippet: <p id="price">the price is $1200</p>
loader.add_xpath('price', '//p[@id="price"]', re='the price is (.*)')
```

replace_xpath (*field_name*, *xpath*, **processors*, ***kwargs*)

Similar to `add_xpath()` but replaces collected data instead of adding it.

default_selector_class

The class used to construct the `selector` of this `XPathItemLoader`, if only a response is given in the constructor. If a selector is given in the constructor this attribute is ignored. This attribute is sometimes overridden in subclasses.

selector

The `XPathSelector` object to extract data from. It's either the selector given in the constructor or one created from the response given in the constructor using the `default_selector_class`. This attribute is meant to be read-only.

3.6.7 Reusing and extending Item Loaders

As your project grows bigger and acquires more and more spiders, maintenance becomes a fundamental problem, specially when you have to deal with many different parsing rules for each spider, having a lot of exceptions, but also wanting to reuse the common processors.

Item Loaders are designed to ease the maintenance burden of parsing rules, without losing flexibility and, at the same time, providing a convenient mechanism for extending and overriding them. For this reason Item Loaders support traditional Python class inheritance for dealing with differences of specific spiders (or groups of spiders).

Suppose, for example, that some particular site encloses their product names in three dashes (ie. ---Plasma TV---) and you don't want to end up scraping those dashes in the final product names.

Here's how you can remove those dashes by reusing and extending the default Product Item Loader (`ProductLoader`):

```
from scrapy.contrib.loader.processor import MapCompose
from myproject.ItemLoaders import ProductLoader

def strip_dashes(x):
    return x.strip('-')

class SiteSpecificLoader(ProductLoader):
    name_in = MapCompose(strip_dashes, ProductLoader.name_in)
```

Another case where extending Item Loaders can be very helpful is when you have multiple source formats, for example XML and HTML. In the XML version you may want to remove CDATA occurrences. Here's an example of how to do it:

```
from scrapy.contrib.loader.processor import MapCompose
from myproject.ItemLoaders import ProductLoader
```

```
from myproject.utils.xml import remove_cdata

class XmlProductLoader(ProductLoader):
    name_in = MapCompose(remove_cdata, ProductLoader.name_in)
```

And that's how you typically extend input processors.

As for output processors, it is more common to declare them in the field metadata, as they usually depend only on the field and not on each specific site parsing rule (as input processors do). See also: [Declaring Input and Output Processors](#).

There are many other possible ways to extend, inherit and override your Item Loaders, and different Item Loaders hierarchies may fit better for different projects. Scrapy only provides the mechanism; it doesn't impose any specific organization of your Loaders collection - that's up to you and your project's needs.

3.6.8 Available built-in processors

Even though you can use any callable function as input and output processors, Scrapy provides some commonly used processors, which are described below. Some of them, like the `MapCompose` (which is typically used as input processor) compose the output of several functions executed in order, to produce the final parsed value.

Here is a list of all built-in processors:

class scrapy.contrib.loader.processor.**Identity**

The simplest processor, which doesn't do anything. It returns the original values unchanged. It doesn't receive any constructor arguments nor accepts Loader contexts.

Example:

```
>>> from scrapy.contrib.loader.processor import Identity
>>> proc = Identity()
>>> proc(['one', 'two', 'three'])
['one', 'two', 'three']
```

class scrapy.contrib.loader.processor.**TakeFirst**

Return the first non-null/non-empty value from the values received, so it's typically used as an output processor to single-valued fields. It doesn't receive any constructor arguments, nor accept Loader contexts.

Example:

```
>>> from scrapy.contrib.loader.processor import TakeFirst
>>> proc = TakeFirst()
>>> proc(['', 'one', 'two', 'three'])
'one'
```

class scrapy.contrib.loader.processor.**Join**(separator=u' ')

Returns the values joined with the separator given in the constructor, which defaults to u' '. It doesn't accept Loader contexts.

When using the default separator, this processor is equivalent to the function: u' '.join

Examples:

```
>>> from scrapy.contrib.loader.processor import Join
>>> proc = Join()
>>> proc(['one', 'two', 'three'])
u'one two three'
>>> proc = Join('<br>')
>>> proc(['one', 'two', 'three'])
u'one<br>two<br>three'
```

class scrapy.contrib.loader.processor.**Compose**(*functions, **default_loader_context)

A processor which is constructed from the composition of the given functions. This means that each input value of this processor is passed to the first function, and the result of that function is passed to the second function, and so on, until the last function returns the output value of this processor.

By default, stop process on None value. This behaviour can be changed by passing keyword argument `stop_on_none=False`.

Example:

```
>>> from scrapy.contrib.loader.processor import Compose
>>> proc = Compose(lambda v: v[0], str.upper)
>>> proc(['hello', 'world'])
'HELLO'
```

Each function can optionally receive a `loader_context` parameter. For those which do, this processor will pass the currently active *Loader context* through that parameter.

The keyword arguments passed in the constructor are used as the default Loader context values passed to each function call. However, the final Loader context values passed to functions are overridden with the currently active Loader context accessible through the `ItemLoader.context()` attribute.

class scrapy.contrib.loader.processor.**MapCompose**(*functions, **default_loader_context)

A processor which is constructed from the composition of the given functions, similar to the *Compose* processor. The difference with this processor is the way internal results are passed among functions, which is as follows:

The input value of this processor is *iterated* and each element is passed to the first function, and the result of that function (for each element) is concatenated to construct a new iterable, which is then passed to the second function, and so on, until the last function is applied for each value of the list of values collected so far. The output values of the last function are concatenated together to produce the output of this processor.

Each particular function can return a value or a list of values, which is flattened with the list of values returned by the same function applied to the other input values. The functions can also return `None` in which case the output of that function is ignored for further processing over the chain.

This processor provides a convenient way to compose functions that only work with single values (instead of iterables). For this reason the *MapCompose* processor is typically used as input processor, since data is often extracted using the `extract()` method of *selectors*, which returns a list of unicode strings.

The example below should clarify how it works:

```
>>> def filter_world(x):
...     return None if x == 'world' else x
...
>>> from scrapy.contrib.loader.processor import MapCompose
>>> proc = MapCompose(filter_world, unicode.upper)
>>> proc([u'hello', u'world', u'this', u'is', u'scrappy'])
[u'HELLO', u'THIS', u'IS', u'SCRAPY']
```

As with the *Compose* processor, functions can receive Loader contexts, and constructor keyword arguments are used as default context values. See *Compose* processor for more info.

3.7 Scrapy shell

The Scrapy shell is an interactive shell where you can try and debug your scraping code very quickly, without having to run the spider. It's meant to be used for testing data extraction code, but you can actually use it for testing any kind of code as it is also a regular Python shell.

The shell is used for testing XPath expressions and see how they work and what data they extract from the web pages you're trying to scrape. It allows you to interactively test your XPaths while you're writing your spider, without having to run the spider to test every change.

Once you get familiarized with the Scrapy shell, you'll see that it's an invaluable tool for developing and debugging your spiders.

If you have [IPython](#) installed, the Scrapy shell will use it (instead of the standard Python console). The [IPython](#) console is much more powerful and provides smart auto-completion and colorized output, among other things.

We highly recommend you install [IPython](#), specially if you're working on Unix systems (where [IPython](#) excels). See the [IPython installation guide](#) for more info.

3.7.1 Launch the shell

To launch the Scrapy shell you can use the `shell` command like this:

```
scrapy shell <url>
```

Where the `<url>` is the URL you want to scrape.

3.7.2 Using the shell

The Scrapy shell is just a regular Python console (or *IPython* console if you have it available) which provides some additional shortcut functions for convenience.

Available Shortcuts

- `shelp()` - print a help with the list of available objects and shortcuts
- `fetch(request_or_url)` - fetch a new response from the given request or URL and update all related objects accordingly.
- `view(response)` - open the given response in your local web browser, for inspection. This will add a `<base>` tag to the response body in order for external links (such as images and style sheets) to display properly. Note, however, that this will create a temporary file in your computer, which won't be removed automatically.

Available Scrapy objects

The Scrapy shell automatically creates some convenient objects from the downloaded page, like the [Response](#) object and the [XPathSelector](#) objects (for both HTML and XML content).

Those objects are:

- `spider` - the Spider which is known to handle the URL, or a [BaseSpider](#) object if there is no spider found for the current URL
- `request` - a [Request](#) object of the last fetched page. You can modify this request using `replace()` or fetch a new request (without leaving the shell) using the `fetch` shortcut.
- `response` - a [Response](#) object containing the last fetched page
- `hxs` - a [HtmlXPathSelector](#) object constructed with the last response fetched
- `xxs` - a [XmlXPathSelector](#) object constructed with the last response fetched
- `settings` - the current *Scrapy settings*

3.7.3 Example of shell session

Here's an example of a typical shell session where we start by scraping the <http://scrapy.org> page, and then proceed to scrape the <http://slashdot.org> page. Finally, we modify the (Slashdot) request method to POST and re-fetch it getting a HTTP 405 (method not allowed) error. We end the session by typing Ctrl-D (in Unix systems) or Ctrl-Z in Windows.

Keep in mind that the data extracted here may not be the same when you try it, as those pages are not static and could have changed by the time you test this. The only purpose of this example is to get you familiarized with how the Scrapy shell works.

First, we launch the shell:

```
scrapy shell http://scrapy.org --nolog
```

Then, the shell fetches the URL (using the Scrapy downloader) and prints the list of available objects and useful shortcuts (you'll notice that these lines all start with the `[s]` prefix):

```
[s] Available objects
[s]   hxs      <HtmlXPathSelector (http://scrapy.org) xpath=None>
[s]   item     Item()
[s]   request  <http://scrapy.org>
[s]   response <http://scrapy.org>
[s]   settings <Settings 'mybot.settings'>
[s]   spider   <scrapy.spider.models.BaseSpider object at 0x2bed9d0>
[s]   xxs      <XmlXPathSelector (http://scrapy.org) xpath=None>
[s] Useful shortcuts:
[s]   shelp()      Prints this help.
[s]   fetch(req_or_url) Fetch a new request or URL and update objects
[s]   view(response) View response in a browser

>>>
```

After that, we can start playing with the objects:

```
>>> hxs.select("//h2/text()").extract()[0]
u'Welcome to Scrapy'

>>> fetch("http://slashdot.org")
[s] Available Scrapy objects:
[s]   hxs      <HtmlXPathSelector (http://slashdot.org) xpath=None>
[s]   item     JobItem()
[s]   request  <GET http://slashdot.org>
[s]   response <200 http://slashdot.org>
[s]   settings <Settings 'jobshot.settings'>
[s]   spider   <BaseSpider 'default' at 0x3c44a10>
[s]   xxs      <XmlXPathSelector (http://slashdot.org) xpath=None>
[s] Useful shortcuts:
[s]   shelp()      Shell help (print this help)
[s]   fetch(req_or_url) Fetch request (or URL) and update local objects
[s]   view(response) View response in a browser

>>> hxs.select("//h2/text()").extract()
[u'News for nerds, stuff that matters']

>>> request = request.replace(method="POST")

>>> fetch(request)
2009-04-03 00:57:39-0300 [default] ERROR: Downloading <http://slashdot.org> from <None>: 405 Method Not Allowed

>>>
```

3.7.4 Invoking the shell from spiders to inspect responses

Sometimes you want to inspect the responses that are being processed in a certain point of your spider, if only to check that response you expect is getting there.

This can be achieved by using the `scrapy.shell.inspect_response` function.

Here's an example of how you would call it from your spider:

```
class MySpider(BaseSpider):
    ...

    def parse(self, response):
        if response.url == 'http://www.example.com/products.php':
            from scrapy.shell import inspect_response
            inspect_response(response)

        # ... your parsing code ..
```

When you run the spider, you will get something similar to this:

```
2009-08-27 19:15:25-0300 [example.com] DEBUG: Crawled <http://www.example.com/> (referer: <None>)
2009-08-27 19:15:26-0300 [example.com] DEBUG: Crawled <http://www.example.com/products.php> (referer: <None>)
[s] Available objects
[s]   hxs          <HtmlXPathSelector (http://www.example.com/products.php) xpath=None>
...

>>> response.url
'http://www.example.com/products.php'
```

Then, you can check if the extraction code is working:

```
>>> hxs.select('//h1')
[]
```

Nope, it doesn't. So you can open the response in your web browser and see if it's the response you were expecting:

```
>>> view(response)
>>>
```

Finally you hit Ctrl-D (or Ctrl-Z in Windows) to exit the shell and resume the crawling:

```
>>> ^D
2009-08-27 19:15:25-0300 [example.com] DEBUG: Crawled <http://www.example.com/product.php?id=1> (referer: <None>)
2009-08-27 19:15:25-0300 [example.com] DEBUG: Crawled <http://www.example.com/product.php?id=2> (referer: <None>)
# ...
```

Note that you can't use the `fetch` shortcut here since the Scrapy engine is blocked by the shell. However, after you leave the shell, the spider will continue crawling where it stopped, as shown above.

3.8 Item Pipeline

After an item has been scraped by a spider, it is sent to the Item Pipeline which process it through several components that are executed sequentially.

Each item pipeline component (sometimes referred as just "Item Pipeline") is a Python class that implements a simple method. They receive an Item and perform an action over it, also deciding if the Item should continue through the pipeline or be dropped and no longer processed.

Typical use for item pipelines are:

- cleansing HTML data
- validating scraped data (checking that the items contain certain fields)
- checking for duplicates (and dropping them)
- storing the scraped item in a database

3.8.1 Writing your own item pipeline

Writing your own item pipeline is easy. Each item pipeline component is a single Python class that must implement the following method:

process_item (*item*, *spider*)

This method is called for every item pipeline component and must either return a `Item` (or any descendant class) object or raise a `DropItem` exception. Dropped items are no longer processed by further pipeline components.

Parameters

- **item** (`Item` object) – the item scraped
- **spider** (`BaseSpider` object) – the spider which scraped the item

Additionally, they may also implement the following methods:

open_spider (*spider*)

This method is called when the spider is opened.

Parameters **spider** (`BaseSpider` object) – the spider which was opened

close_spider (*spider*)

This method is called when the spider is closed.

Parameters **spider** (`BaseSpider` object) – the spider which was closed

3.8.2 Item pipeline example

Price validation and dropping items with no prices

Let's take a look at the following hypothetical pipeline that adjusts the `price` attribute for those items that do not include VAT (`price_excludes_vat` attribute), and drops those items which don't contain a price:

```
from scrapy.exceptions import DropItem

class PricePipeline(object):

    vat_factor = 1.15

    def process_item(self, item, spider):
        if item['price']:
            if item['price_excludes_vat']:
                item['price'] = item['price'] * self.vat_factor
            return item
        else:
            raise DropItem("Missing price in %s" % item)
```

Write items to a JSON file

The following pipeline stores all scraped items (from all spiders) into a single `items.jsonl` file, containing one item per line serialized in JSON format:

```
import json

class JsonWriterPipeline(object):

    def __init__(self):
        self.file = open('items.jsonl', 'wb')

    def process_item(self, item, spider):
        line = json.dumps(dict(item)) + "\n"
        self.file.write(line)
        return item
```

Note: The purpose of `JsonWriterPipeline` is just to introduce how to write item pipelines. If you really want to store all scraped items into a JSON file you should use the *Feed exports*.

Duplicates filter

A filter that looks for duplicate items, and drops those items that were already processed. Let say that our items have an unique id, but our spider returns multiples items with the same id:

```
from scrapy import signals
from scrapy.exceptions import DropItem

class DuplicatesPipeline(object):

    def __init__(self):
        self.ids_seen = set()

    def process_item(self, item, spider):
        if item['id'] in self.ids_seen:
            raise DropItem("Duplicate item found: %s" % item)
        else:
            self.ids_seen.add(item['id'])
            return item
```

3.8.3 Activating an Item Pipeline component

To activate an Item Pipeline component you must add its class to the `ITEM_PIPELINES` list, like in the following example:

```
ITEM_PIPELINES = [
    'myproject.pipeline.PricePipeline',
    'myproject.pipeline.JsonWriterPipeline',
]
```

3.9 Feed exports

New in version 0.10. One of the most frequently required features when implementing scrapers is being able to store the scraped data properly and, quite often, that means generating a “export file” with the scraped data (commonly called “export feed”) to be consumed by other systems.

Scrapy provides this functionality out of the box with the Feed Exports, which allows you to generate a feed with the scraped items, using multiple serialization formats and storage backends.

3.9.1 Serialization formats

For serializing the scraped data, the feed exports use the *Item exporters* and these formats are supported out of the box:

- *JSON*
- *JSON lines*
- *CSV*
- *XML*

But you can also extend the supported format through the `FEED_EXPORTERS` setting.

JSON

- `FEED_FORMAT: json`
- Exporter used: `JsonItemExporter`
- See *this warning* if you’re using JSON with large feeds

JSON lines

- `FEED_FORMAT: jsonlines`
- Exporter used: `JsonLinesItemExporter`

CSV

- `FEED_FORMAT: csv`
- Exporter used: `CsvItemExporter`

XML

- `FEED_FORMAT: xml`
- Exporter used: `XmlItemExporter`

Pickle

- `FEED_FORMAT: pickle`
- Exporter used: `PickleItemExporter`

Marshal

- `FEED_FORMAT`: `marshal`
- Exporter used: `MarshalItemExporter`

3.9.2 Storages

When using the feed exports you define where to store the feed using a [URI](#) (through the `FEED_URI` setting). The feed exports supports multiple storage backend types which are defined by the URI scheme.

The storages backends supported out of the box are:

- *Local filesystem*
- *FTP*
- *S3* (requires `boto`)
- *Standard output*

Some storage backends may be unavailable if the required external libraries are not available. For example, the S3 backend is only available if the `boto` library is installed.

3.9.3 Storage URI parameters

The storage URI can also contain parameters that get replaced when the feed is being created. These parameters are:

- `%(time)s` - gets replaced by a timestamp when the feed is being created
- `%(name)s` - gets replaced by the spider name

Any other named parameter gets replaced by the spider attribute of the same name. For example, `%(site_id)s` would get replaced by the `spider.site_id` attribute the moment the feed is being created.

Here are some examples to illustrate:

- Store in FTP using one directory per spider:
 - `ftp://user:password@ftp.example.com/scraping/feeds/%(name)s/%(time)s.json`
- Store in S3 using one directory per spider:
 - `s3://mybucket/scraping/feeds/%(name)s/%(time)s.json`

3.9.4 Storage backends

Local filesystem

The feeds are stored in the local filesystem.

- URI scheme: `file`
- Example URI: `file:///tmp/export.csv`
- Required external libraries: `none`

Note that for the local filesystem storage (only) you can omit the scheme if you specify an absolute path like `/tmp/export.csv`. This only works on Unix systems though.

FTP

The feeds are stored in a FTP server.

- URI scheme: `ftp`
- Example URI: `ftp://user:pass@ftp.example.com/path/to/export.csv`
- Required external libraries: `none`

S3

The feeds are stored on [Amazon S3](#).

- URI scheme: `s3`
- Example URIs:
 - `s3://mybucket/path/to/export.csv`
 - `s3://aws_key:aws_secret@mybucket/path/to/export.csv`
- Required external libraries: `boto`

The AWS credentials can be passed as user/password in the URI, or they can be passed through the following settings:

- `AWS_ACCESS_KEY_ID`
- `AWS_SECRET_ACCESS_KEY`

Standard output

The feeds are written to the standard output of the Scrapy process.

- URI scheme: `stdout`
- Example URI: `stdout:`
- Required external libraries: `none`

3.9.5 Settings

These are the settings used for configuring the feed exports:

- `FEED_URI` (mandatory)
- `FEED_FORMAT`
- `FEED_STORAGES`
- `FEED_EXPORTERS`
- `FEED_STORE_EMPTY`

FEED_URI

Default: `None`

The URI of the export feed. See [Storage backends](#) for supported URI schemes.

This setting is required for enabling the feed exports.

FEED_FORMAT

The serialization format to be used for the feed. See *Serialization formats* for possible values.

FEED_STORE_EMPTY

Default: `False`

Whether to export empty feeds (ie. feeds with no items).

FEED_STORAGES

Default:: `{}`

A dict containing additional feed storage backends supported by your project. The keys are URI schemes and the values are paths to storage classes.

FEED_STORAGES_BASE

Default:

```
{
    '': 'scrapy.contrib.feedexport.FileFeedStorage',
    'file': 'scrapy.contrib.feedexport.FileFeedStorage',
    'stdout': 'scrapy.contrib.feedexport.StdoutFeedStorage',
    's3': 'scrapy.contrib.feedexport.S3FeedStorage',
    'ftp': 'scrapy.contrib.feedexport.FTPFeedStorage',
}
```

A dict containing the built-in feed storage backends supported by Scrapy.

FEED_EXPORTERS

Default:: `{}`

A dict containing additional exporters supported by your project. The keys are URI schemes and the values are paths to *Item exporter* classes.

FEED_EXPORTERS_BASE

Default:

```
FEED_EXPORTERS_BASE = {
    'json': 'scrapy.contrib.exporter.JsonItemExporter',
    'jsonlines': 'scrapy.contrib.exporter.JsonLinesItemExporter',
    'csv': 'scrapy.contrib.exporter.CsvItemExporter',
    'xml': 'scrapy.contrib.exporter.XmlItemExporter',
    'marshal': 'scrapy.contrib.exporter.MarshalItemExporter',
}
```

A dict containing the built-in feed exporters supported by Scrapy.

3.10 Link Extractors

LinkExtractors are objects whose only purpose is to extract links from web pages (`scrapy.http.Response` objects) which will be eventually followed.

There are two Link Extractors available in Scrapy by default, but you create your own custom Link Extractors to suit your needs by implementing a simple interface.

The only public method that every LinkExtractor has is `extract_links`, which receives a `Response` object and returns a list of links. Link Extractors are meant to be instantiated once and their `extract_links` method called several times with different responses, to extract links to follow.

Link extractors are used in the `CrawlSpider` class (available in Scrapy), through a set of rules, but you can also use it in your spiders, even if you don't subclass from `CrawlSpider`, as its purpose is very simple: to extract links.

3.10.1 Built-in link extractors reference

All available link extractors classes bundled with Scrapy are provided in the `scrapy.contrib.linkextractors` module.

SgmlLinkExtractor

```
class scrapy.contrib.linkextractors.sgml.SgmlLinkExtractor (allow=(), deny=(),
                                                           allow_domains=(),
                                                           deny_domains=(),
                                                           deny_extensions=None,
                                                           restrict_xpaths=(),
                                                           tags=('a', 'area'),
                                                           attrs=('href'), canonical-
                                                           ize=True, unique=True,
                                                           process_value=None)
```

The `SgmlLinkExtractor` extends the base `BaseSgmlLinkExtractor` by providing additional filters that you can specify to extract links, including regular expressions patterns that the links must match to be extracted. All those filters are configured through these constructor parameters:

Parameters

- **allow** (*a regular expression (or list of)*) – a single regular expression (or list of regular expressions) that the (absolute) urls must match in order to be extracted. If not given (or empty), it will match all links.
- **deny** (*a regular expression (or list of)*) – a single regular expression (or list of regular expressions) that the (absolute) urls must match in order to be excluded (ie. not extracted). It has precedence over the `allow` parameter. If not given (or empty) it won't exclude any links.
- **allow_domains** (*str or list*) – a single value or a list of string containing domains which will be considered for extracting the links
- **deny_domains** (*str or list*) – a single value or a list of strings containing domains which won't be considered for extracting the links
- **deny_extensions** (*list*) – a list of extensions that should be ignored when extracting links. If not given, it will default to the `IGNORED_EXTENSIONS` list defined in the `scrapy.linkextractor` module.

- **restrict_xpaths** (*str or list*) – is a XPath (or list of XPath's) which defines regions inside the response where links should be extracted from. If given, only the text selected by those XPath will be scanned for links. See examples below.
- **tags** (*str or list*) – a tag or a list of tags to consider when extracting links. Defaults to ('a', 'area').
- **attrs** (*list*) – list of attributes which should be considered when looking for links to extract (only for those tags specified in the `tags` parameter). Defaults to ('href',)
- **canonicalize** (*boolean*) – canonicalize each extracted url (using `scrapy.utils.url.canonicalize_url`). Defaults to `True`.
- **unique** (*boolean*) – whether duplicate filtering should be applied to extracted links.
- **process_value** (*callable*) – see `process_value` argument of `BaseSgmlLinkExtractor` class constructor

BaseSgmlLinkExtractor

```
class scrapy.contrib.linkextractors.sgml.BaseSgmlLinkExtractor(tag="a",
                                                                attr="href",
                                                                unique=False, process_value=None)
```

The purpose of this Link Extractor is only to serve as a base class for the `SgmlLinkExtractor`. You should use that one instead.

The constructor arguments are:

Parameters

- **tag** (*str or callable*) – either a string (with the name of a tag) or a function that receives a tag name and returns `True` if links should be extracted from that tag, or `False` if they shouldn't. Defaults to 'a'. request (once it's downloaded) as its first parameter. For more information, see [Passing additional data to callback functions](#).
- **attr** (*str or callable*) – either string (with the name of a tag attribute), or a function that receives an attribute name and returns `True` if links should be extracted from it, or `False` if they shouldn't. Defaults to `href`.
- **unique** (*boolean*) – is a boolean that specifies if a duplicate filtering should be applied to links extracted.
- **process_value** (*callable*) – a function which receives each value extracted from the tag and attributes scanned and can modify the value and return a new one, or return `None` to ignore the link altogether. If not given, `process_value` defaults to `lambda x: x`.

For example, to extract links from this code:

```
<a href="javascript:goToPage('../other/page.html'); return false">Link text</a>
```

You can use the following function in `process_value`:

```
def process_value(value):
    m = re.search("javascript:goToPage\('(.*?)'", value)
    if m:
        return m.group(1)
```

Command line tool Learn about the command-line tool used to manage your Scrapy project.

Items Define the data you want to scrape.

Spiders Write the rules to crawl your websites.

Selectors Extract the data from web pages using XPath.

Scrapy shell Test your extraction code in an interactive environment.

Item Loaders Populate your items with the extracted data.

Item Pipeline Post-process and store your scraped data.

Feed exports Output your scraped data using different formats and storages.

Link Extractors Convenient classes to extract links to follow from pages.

Built-in services

4.1 Logging

Scrapy provides a logging facility which can be used through the `scrapy.log` module. The current underlying implementation uses [Twisted logging](#) but this may change in the future.

The logging service must be explicitly started through the `scrapy.log.start()` function.

4.1.1 Log levels

Scrapy provides 5 logging levels:

1. `CRITICAL` - for critical errors
2. `ERROR` - for regular errors
3. `WARNING` - for warning messages
4. `INFO` - for informational messages
5. `DEBUG` - for debugging messages

4.1.2 How to set the log level

You can set the log level using the `-loglevel/-L` command line option, or using the `LOG_LEVEL` setting.

4.1.3 How to log messages

Here's a quick example of how to log a message using the `WARNING` level:

```
from scrapy import log
log.msg("This is a warning", level=log.WARNING)
```

4.1.4 Logging from Spiders

The recommended way to log from spiders is by using the Spider `log()` method, which already populates the `spider` argument of the `scrapy.log.msg()` function. The other arguments are passed directly to the `msg()` function.

4.1.5 scrapy.log module

`scrapy.log.start` (*logfile=None, loglevel=None, logstdout=None*)

Start the logging facility. This must be called before actually logging any messages. Otherwise, messages logged before this call will get lost.

Parameters

- **logfile** (*str*) – the file path to use for logging output. If omitted, the `LOG_FILE` setting will be used. If both are `None`, the log will be sent to standard error.
- **loglevel** – the minimum logging level to log. Available values are: `CRITICAL`, `ERROR`, `WARNING`, `INFO` and `DEBUG`.
- **logstdout** (*boolean*) – if `True`, all standard output (and error) of your application will be logged instead. For example if you “print ‘hello’” it will appear in the Scrapy log. If omitted, the `LOG_STDOUT` setting will be used.

`scrapy.log.msg` (*message, level=INFO, spider=None*)

Log a message

Parameters

- **message** (*str*) – the message to log
- **level** – the log level for this message. See [Log levels](#).
- **spider** (`BaseSpider` object) – the spider to use for logging this message. This parameter should always be used when logging things related to a particular spider.

`scrapy.log.CRITICAL`

Log level for critical errors

`scrapy.log.ERROR`

Log level for errors

`scrapy.log.WARNING`

Log level for warnings

`scrapy.log.INFO`

Log level for informational messages (recommended level for production deployments)

`scrapy.log.DEBUG`

Log level for debugging messages (recommended level for development)

4.1.6 Logging settings

These settings can be used to configure the logging:

- `LOG_ENABLED`
- `LOG_ENCODING`
- `LOG_FILE`

- `LOG_LEVEL`
- `LOG_STDOUT`

4.2 Stats Collection

Scrapy provides a convenient facility for collecting stats in the form of key/values, where values are often counters. The facility is called the Stats Collector, and can be accessed through the `stats` attribute of the *Crawler API*, as illustrated by the examples in the *Common Stats Collector uses* section below.

However, the Stats Collector is always available, so you can always import it in your module and use its API (to increment or set new stat keys), regardless of whether the stats collection is enabled or not. If it's disabled, the API will still work but it won't collect anything. This is aimed at simplifying the stats collector usage: you should spend no more than one line of code for collecting stats in your spider, Scrapy extension, or whatever code you're using the Stats Collector from.

Another feature of the Stats Collector is that it's very efficient (when enabled) and extremely efficient (almost unnoticeable) when disabled.

The Stats Collector keeps a stats table per open spider which is automatically opened when the spider is opened, and closed when the spider is closed.

4.2.1 Common Stats Collector uses

Access the stats collector through the `stats` attribute. Here is an example of an extension that access stats:

```
class ExtensionThatAccessStats(object):

    def __init__(self, stats):
        self.stats = stats

    @classmethod
    def from_crawler(cls, crawler):
        return cls(crawler.stats)
```

Set stat value:

```
stats.set_value('hostname', socket.gethostname())
```

Increment stat value:

```
stats.inc_value('pages_crawled')
```

Set stat value only if greater than previous:

```
stats.max_value('max_items_scraped', value)
```

Set stat value only if lower than previous:

```
stats.min_value('min_free_memory_percent', value)
```

Get stat value:

```
>>> stats.get_value('pages_crawled')
8
```

Get all stats:

```
>>> stats.get_stats()
{'pages_crawled': 1238, 'start_time': datetime.datetime(2009, 7, 14, 21, 47, 28, 977139)}
```

4.2.2 Available Stats Collectors

Besides the basic `StatsCollector` there are other Stats Collectors available in Scrapy which extend the basic Stats Collector. You can select which Stats Collector to use through the `STATS_CLASS` setting. The default Stats Collector used is the `MemoryStatsCollector`.

MemoryStatsCollector

class scrapy.statscol.MemoryStatsCollector

A simple stats collector that keeps the stats of the last scraping run (for each spider) in memory, after they're closed. The stats can be accessed through the `spider_stats` attribute, which is a dict keyed by spider domain name.

This is the default Stats Collector used in Scrapy.

spider_stats

A dict of dicts (keyed by spider name) containing the stats of the last scraping run for each spider.

DummyStatsCollector

class scrapy.statscol.DummyStatsCollector

A Stats collector which does nothing but is very efficient (because it does nothing). This stats collector can be set via the `STATS_CLASS` setting, to disable stats collect in order to improve performance. However, the performance penalty of stats collection is usually marginal compared to other Scrapy workload like parsing pages.

4.3 Sending e-mail

Although Python makes sending e-mails relatively easy via the `smtplib` library, Scrapy provides its own facility for sending e-mails which is very easy to use and it's implemented using `Twisted non-blocking IO`, to avoid interfering with the non-blocking IO of the crawler. It also provides a simple API for sending attachments and it's very easy to configure, with a few *settings*.

4.3.1 Quick example

There are two ways to instantiate the mail sender. You can instantiate it using the standard constructor:

```
from scrapy.mail import MailSender
mailer = MailSender()
```

Or you can instantiate it passing a Scrapy settings object, which will respect the *settings*:

```
mailer = MailSender.from_settings(settings)
```

And here is how to use it to send an e-mail (without attachments):

```
mailer.send(to=["someone@example.com"], subject="Some subject", body="Some body", cc=["another@examp
```


4.3.2 MailSender class reference

MailSender is the preferred class to use for sending emails from Scrapy, as it uses [Twisted non-blocking IO](#), like the rest of the framework.

```
class scrapy.mail.MailSender (smtphost=None, mailfrom=None, smtpuser=None, smtppass=None,
                             smtpport=None)
```

Parameters

- **smtphost** (*str*) – the SMTP host to use for sending the emails. If omitted, the [MAIL_HOST](#) setting will be used.
- **mailfrom** (*str*) – the address used to send emails (in the From: header). If omitted, the [MAIL_FROM](#) setting will be used.
- **smtpuser** – the SMTP user. If omitted, the [MAIL_USER](#) setting will be used. If not given, no SMTP authentication will be performed.
- **smtppass** (*str*) – the SMTP pass for authentication.
- **smtpport** (*int*) – the SMTP port to connect to

```
classmethod from_settings (settings)
```

Instantiate using a Scrapy settings object, which will respect [these Scrapy settings](#).

Parameters **settings** ([scrapy.settings.Settings](#) object) – the e-mail recipients

```
send (to, subject, body, cc=None, attachs=())
```

Send email to the given recipients.

Parameters

- **to** (*list*) – the e-mail recipients
- **subject** (*str*) – the subject of the e-mail
- **cc** (*list*) – the e-mails to CC
- **body** (*str*) – the e-mail body
- **attachs** (*iterable*) – an iterable of tuples (attach_name, mimetype, file_object) where attach_name is a string with the name that will appear on the e-mail's attachment, mimetype is the mimetype of the attachment and file_object is a readable file object with the contents of the attachment

4.3.3 Mail settings

These settings define the default constructor values of the [MailSender](#) class, and can be used to configure e-mail notifications in your project without writing any code (for those extensions and code that uses [MailSender](#)).

MAIL_FROM

Default: ' scrapy@localhost '

Sender email to use (From: header) for sending emails.

MAIL_HOST

Default: 'localhost'

SMTP host to use for sending emails.

MAIL_PORT

Default: 25

SMTP port to use for sending emails.

MAIL_USER

Default: None

User to use for SMTP authentication. If disabled no SMTP authentication will be performed.

MAIL_PASS

Default: None

Password to use for SMTP authentication, along with `MAIL_USER`.

4.4 Telnet Console

Scrapy comes with a built-in telnet console for inspecting and controlling a Scrapy running process. The telnet console is just a regular python shell running inside the Scrapy process, so you can do literally anything from it.

The telnet console is a *built-in Scrapy extension* which comes enabled by default, but you can also disable it if you want. For more information about the extension itself see *Telnet console extension*.

4.4.1 How to access the telnet console

The telnet console listens in the TCP port defined in the `TELNETCONSOLE_PORT` setting, which defaults to 6023. To access the console you need to type:

```
telnet localhost 6023
>>>
```

You need the telnet program which comes installed by default in Windows, and most Linux distros.

4.4.2 Available variables in the telnet console

The telnet console is like a regular Python shell running inside the Scrapy process, so you can do anything from it including importing new modules, etc.

However, the telnet console comes with some default variables defined for convenience:

| Shortcut | Description |
|-------------------------|-----------------------------------------------------------------------|
| <code>crawler</code> | the Scrapy Crawler (<code>scrapy.crawler.Crawler</code> object) |
| <code>engine</code> | <code>Crawler.engine</code> attribute |
| <code>spider</code> | the active spider |
| <code>slot</code> | the engine slot |
| <code>extensions</code> | the Extension Manager (<code>Crawler.extensions</code> attribute) |
| <code>stats</code> | the Stats Collector (<code>Crawler.stats</code> attribute) |
| <code>settings</code> | the Scrapy settings object (<code>Crawler.settings</code> attribute) |
| <code>est</code> | print a report of the engine status |
| <code>prefs</code> | for memory debugging (see Debugging memory leaks) |
| <code>p</code> | a shortcut to the <code>pprint.pprint</code> function |
| <code>hpy</code> | for memory debugging (see Debugging memory leaks) |

4.4.3 Telnet console usage examples

Here are some example tasks you can do with the telnet console:

View engine status

You can use the `est()` method of the Scrapy engine to quickly show its state using the telnet console:

```
telnet localhost 6023
>>> est()
Execution engine status

time()-engine.start_time           : 21.3188259602
engine.is_idle()                   : False
engine.has_capacity()               : True
engine.scheduler.is_idle()         : False
len(engine.scheduler.pending_requests) : 1
engine.downloader.is_idle()        : False
len(engine.downloader.slots)       : 1
engine.scrapers.is_idle()          : False
len(engine.scrapers.slots)         : 1

Spider: <GayotSpider 'gayotcom' at 0x2dc2b10>
  engine.spider_is_idle(spider)      : False
  engine.slots[spider].closing       : False
  len(engine.scheduler.pending_requests[spider]) : 11504
  len(engine.downloader.slots[spider].queue) : 9
  len(engine.downloader.slots[spider].active) : 17
  len(engine.downloader.slots[spider].transferring) : 8
  engine.downloader.slots[spider].lastseen : 1311311093.61
  len(engine.scrapers.slots[spider].queue) : 0
  len(engine.scrapers.slots[spider].active) : 0
  engine.scrapers.slots[spider].active_size : 0
  engine.scrapers.slots[spider].itemproc_size : 0
  engine.scrapers.slots[spider].needs_backout() : False
```

Pause, resume and stop the Scrapy engine

To pause:

```
telnet localhost 6023
>>> engine.pause()
>>>
```

To resume:

```
telnet localhost 6023
>>> engine.unpause()
>>>
```

To stop:

```
telnet localhost 6023
>>> engine.stop()
Connection closed by foreign host.
```

4.4.4 Telnet Console signals

`scrapy.telnet.update_telnet_vars` (*telnet_vars*)

Sent just before the telnet console is opened. You can hook up to this signal to add, remove or update the variables that will be available in the telnet local namespace. In order to do that, you need to update the `telnet_vars` dict in your handler.

Parameters `telnet_vars` (*dict*) – the dict of telnet variables

4.4.5 Telnet settings

These are the settings that control the telnet console's behaviour:

TELNETCONSOLE_PORT

Default: [6023, 6073]

The port range to use for the telnet console. If set to `None` or `0`, a dynamically assigned port is used.

TELNETCONSOLE_HOST

Default: '0.0.0.0'

The interface the telnet console should listen on

4.5 Web Service

Scrapy comes with a built-in web service for monitoring and controlling a running crawler. The service exposes most resources using the [JSON-RPC 2.0](#) protocol, but there are also other (read-only) resources which just output JSON data.

Provides an extensible web service for managing a Scrapy process. It's enabled by the [WEBSERVICE_ENABLED](#) setting. The web server will listen in the port specified in [WEBSERVICE_PORT](#), and will log to the file specified in [WEBSERVICE_LOGFILE](#).

The web service is a *built-in Scrapy extension* which comes enabled by default, but you can also disable it if you're running tight on memory.

4.5.1 Web service resources

The web service contains several resources, defined in the `WEBSERVICE_RESOURCES` setting. Each resource provides a different functionality. See *Available JSON-RPC resources* for a list of resources available by default.

Although you can implement your own resources using any protocol, there are two kinds of resources bundled with Scrapy:

- Simple JSON resources - which are read-only and just output JSON data
- JSON-RPC resources - which provide direct access to certain Scrapy objects using the [JSON-RPC 2.0](#) protocol

Available JSON-RPC resources

These are the JSON-RPC resources available by default in Scrapy:

Crawler JSON-RPC resource

class `scrapy.contrib.webservice.crawler.CrawlerResource`
Provides access to the main Crawler object that controls the Scrapy process.
Available by default at: <http://localhost:6080/crawler>

Stats Collector JSON-RPC resource

class `scrapy.contrib.webservice.stats.StatsResource`
Provides access to the Stats Collector used by the crawler.
Available by default at: <http://localhost:6080/stats>

Spider Manager JSON-RPC resource

You can access the spider manager JSON-RPC resource through the *Crawler JSON-RPC resource* at: <http://localhost:6080/crawler/spiders>

Extension Manager JSON-RPC resource

You can access the extension manager JSON-RPC resource through the *Crawler JSON-RPC resource* at: <http://localhost:6080/crawler/spiders>

Available JSON resources

These are the JSON resources available by default:

Engine status JSON resource

class `scrapy.contrib.webservice.enginestatus.EngineStatusResource`
Provides access to engine status metrics.
Available by default at: <http://localhost:6080/enginestatus>

4.5.2 Web service settings

These are the settings that control the web service behaviour:

WEBSERVICE_ENABLED

Default: `True`

A boolean which specifies if the web service will be enabled (provided its extension is also enabled).

WEBSERVICE_LOGFILE

Default: `None`

A file to use for logging HTTP requests made to the web service. If unset web the log is sent to standard scrapy log.

WEBSERVICE_PORT

Default: `[6080, 7030]`

The port range to use for the web service. If set to `None` or `0`, a dynamically assigned port is used.

WEBSERVICE_HOST

Default: `'0.0.0.0'`

The interface the web service should listen on

WEBSERVICE_RESOURCES

Default: `{}`

The list of web service resources enabled for your project. See [Web service resources](#). These are added to the ones available by default in Scrapy, defined in the `WEBSERVICE_RESOURCES_BASE` setting.

WEBSERVICE_RESOURCES_BASE

Default:

```
{
    'scrapy.contrib.webservice.crawler.CrawlerResource': 1,
    'scrapy.contrib.webservice.enginestatus.EngineStatusResource': 1,
    'scrapy.contrib.webservice.stats.StatsResource': 1,
}
```

The list of web service resources available by default in Scrapy. You shouldn't change this setting in your project, change `WEBSERVICE_RESOURCES` instead. If you want to disable some resource set its value to `None` in `WEBSERVICE_RESOURCES`.

4.5.3 Writing a web service resource

Web service resources are implemented using the Twisted Web API. See this [Twisted Web guide](#) for more information on Twisted web and Twisted web resources.

To write a web service resource you should subclass the `JsonResource` or `JsonRpcResource` classes and implement the `renderGET` method.

class `scrapy.webservice.JsonResource`

A subclass of `twisted.web.resource.Resource` that implements a JSON web service resource. See

ws_name

The name by which the Scrapy web service will know this resource, and also the path where this resource will listen. For example, assuming Scrapy web service is listening on <http://localhost:6080/> and the `ws_name` is `'resource1'` the URL for that resource will be:

<http://localhost:6080/resource1/>

class `scrapy.webservice.JsonRpcResource` (*crawler*, *target=None*)

This is a subclass of `JsonResource` for implementing JSON-RPC resources. JSON-RPC resources wrap Python (Scrapy) objects around a JSON-RPC API. The resource wrapped must be returned by the `get_target()` method, which returns the target passed in the constructor by default

get_target()

Return the object wrapped by this JSON-RPC resource. By default, it returns the object passed on the constructor.

4.5.4 Examples of web service resources

StatsResource (JSON-RPC resource)

```
from scrapy.webservice import JsonRpcResource

class StatsResource(JsonRpcResource):

    ws_name = 'stats'

    def __init__(self, crawler):
        JsonRpcResource.__init__(self, crawler, crawler.stats)
```

EngineStatusResource (JSON resource)

```
from scrapy.webservice import JsonResource
from scrapy.utils.engine import get_engine_status

class EngineStatusResource(JsonResource):

    ws_name = 'enginestatus'

    def __init__(self, crawler, spider_name=None):
        JsonResource.__init__(self, crawler)
        self._spider_name = spider_name
        self.isLeaf = spider_name is not None

    def render_GET(self, txrequest):
        status = get_engine_status(self.crawler.engine)
```

```
    if self._spider_name is None:
        return status
    for sp, st in status['spiders'].items():
        if sp.name == self._spider_name:
            return st

    def getChild(self, name, txrequest):
        return EngineStatusResource(name, self.crawler)
```

4.5.5 Example of web service client

scrapy-ws.py script

```
#!/usr/bin/env python
"""
Example script to control a Scrapy server using its JSON-RPC web service.

It only provides a reduced functionality as its main purpose is to illustrate
how to write a web service client. Feel free to improve or write you own.

Also, keep in mind that the JSON-RPC API is not stable. The recommended way for
controlling a Scrapy server is through the execution queue (see the "queue"
command).

"""

import sys, optparse, urllib, json
from urlparse import urljoin

from scrapy.utils.jsonrpc import jsonrpc_client_call, JsonRpcError

def get_commands():
    return {
        'help': cmd_help,
        'stop': cmd_stop,
        'list-available': cmd_list_available,
        'list-running': cmd_list_running,
        'list-resources': cmd_list_resources,
        'get-global-stats': cmd_get_global_stats,
        'get-spider-stats': cmd_get_spider_stats,
    }

def cmd_help(args, opts):
    """help - list available commands"""
    print "Available commands:"
    for _, func in sorted(get_commands().items()):
        print "  ", func.__doc__

def cmd_stop(args, opts):
    """stop <spider> - stop a running spider"""
    jsonrpc_call(opts, 'crawler/engine', 'close_spider', args[0])

def cmd_list_running(args, opts):
    """list-running - list running spiders"""
    for x in json_get(opts, 'crawler/engine/open_spiders'):
        print x
```

```

def cmd_list_available(args, opts):
    """list-available - list name of available spiders"""
    for x in jsonrpc_call(opts, 'crawler/spiders', 'list'):
        print x

def cmd_list_resources(args, opts):
    """list-resources - list available web service resources"""
    for x in json_get(opts, '')['resources']:
        print x

def cmd_get_spider_stats(args, opts):
    """get-spider-stats <spider> - get stats of a running spider"""
    stats = jsonrpc_call(opts, 'stats', 'get_stats', args[0])
    for name, value in stats.items():
        print "%-40s %s" % (name, value)

def cmd_get_global_stats(args, opts):
    """get-global-stats - get global stats"""
    stats = jsonrpc_call(opts, 'stats', 'get_stats')
    for name, value in stats.items():
        print "%-40s %s" % (name, value)

def get_wsurl(opts, path):
    return urljoin("http://%s:%s/" % (opts.host, opts.port), path)

def jsonrpc_call(opts, path, method, *args, **kwargs):
    url = get_wsurl(opts, path)
    return jsonrpc_client_call(url, method, *args, **kwargs)

def json_get(opts, path):
    url = get_wsurl(opts, path)
    return json.loads(urllib.urlopen(url).read())

def parse_opts():
    usage = "%prog [options] <command> [arg] ..."
    description = "Scrapy web service control script. Use '%prog help' " \
        "to see the list of available commands."
    op = optparse.OptionParser(usage=usage, description=description)
    op.add_option("-H", dest="host", default="localhost", \
        help="Scrapy host to connect to")
    op.add_option("-P", dest="port", type="int", default=6080, \
        help="Scrapy port to connect to")
    opts, args = op.parse_args()
    if not args:
        op.print_help()
        sys.exit(2)
    cmdname, cmdargs, opts = args[0], args[1:], opts
    commands = get_commands()
    if cmdname not in commands:
        sys.stderr.write("Unknown command: %s\n\n" % cmdname)
        cmd_help(None, None)
        sys.exit(1)
    return commands[cmdname], cmdargs, opts

def main():
    cmd, args, opts = parse_opts()
    try:
        cmd(args, opts)

```

```
except IndexError:
    print cmd.__doc__
except JsonRpcError, e:
    print str(e)
    if e.data:
        print "Server Traceback below:"
        print e.data

if __name__ == '__main__':
    main()
```

Logging Understand the simple logging facility provided by Scrapy.

Stats Collection Collect statistics about your scraping crawler.

Sending e-mail Send email notifications when certain events occur.

Telnet Console Inspect a running crawler using a built-in Python console.

Web Service Monitor and control a crawler using a web service.

Solving specific problems

5.1 Frequently Asked Questions

5.1.1 How does Scrapy compare to BeautifulSoup or lxml?

[BeautifulSoup](#) and [lxml](#) are libraries for parsing HTML and XML. Scrapy is an application framework for writing web spiders that crawl web sites and extract data from them.

Scrapy provides a built-in mechanism for extracting data (called *selectors*) but you can easily use [BeautifulSoup](#) (or [lxml](#)) instead, if you feel more comfortable working with them. After all, they're just parsing libraries which can be imported and used from any Python code.

In other words, comparing [BeautifulSoup](#) (or [lxml](#)) to Scrapy is like comparing [jinja2](#) to [Django](#).

5.1.2 What Python versions does Scrapy support?

Scrapy runs in Python 2.6 and 2.7.

5.1.3 Does Scrapy work with Python 3.0?

No, and there are no plans to port Scrapy to Python 3.0 yet. At the moment, Scrapy works with Python 2.6 and 2.7.

See Also:

What Python versions does Scrapy support?.

5.1.4 Did Scrapy “steal” X from Django?

Probably, but we don't like that word. We think [Django](#) is a great open source project and an example to follow, so we've used it as an inspiration for Scrapy.

We believe that, if something is already done well, there's no need to reinvent it. This concept, besides being one of the foundations for open source and free software, not only applies to software but also to documentation, procedures, policies, etc. So, instead of going through each problem ourselves, we choose to copy ideas from those projects that have already solved them properly, and focus on the real problems we need to solve.

We'd be proud if Scrapy serves as an inspiration for other projects. Feel free to steal from us!

5.1.5 Does Scrapy work with HTTP proxies?

Yes. Support for HTTP proxies is provided (since Scrapy 0.8) through the HTTP Proxy downloader middleware. See [HttpProxyMiddleware](#).

5.1.6 Scrapy crashes with: ImportError: No module named win32api

You need to install [pywin32](#) because of [this Twisted bug](#).

5.1.7 How can I simulate a user login in my spider?

See *Using FormRequest.from_response() to simulate a user login*.

5.1.8 Does Scrapy crawl in breath-first or depth-first order?

By default, Scrapy uses a **LIFO** queue for storing pending requests, which basically means that it crawls in **DFO** order. This order is more convenient in most cases. If you do want to crawl in true **BFO** order, you can do it by setting the following settings:

```
DEPTH_PRIORITY = 1
SCHEDULER_DISK_QUEUE = 'scrapy.squeue.PickleFifoDiskQueue'
SCHEDULER_MEMORY_QUEUE = 'scrapy.squeue.FifoMemoryQueue'
```

5.1.9 My Scrapy crawler has memory leaks. What can I do?

See *Debugging memory leaks*.

Also, Python has a builtin memory leak issue which is described in *Leaks without leaks*.

5.1.10 How can I make Scrapy consume less memory?

See previous question.

5.1.11 Can I use Basic HTTP Authentication in my spiders?

Yes, see [HttpAuthMiddleware](#).

5.1.12 Why does Scrapy download pages in English instead of my native language?

Try changing the default [Accept-Language](#) request header by overriding the `DEFAULT_REQUEST_HEADERS` setting.

5.1.13 Where can I find some example Scrapy projects?

See *Examples*.

5.1.14 Can I run a spider without creating a project?

Yes. You can use the `runspider` command. For example, if you have a spider written in a `my_spider.py` file you can run it with:

```
scrapy runspider my_spider.py
```

See `runspider` command for more info.

5.1.15 I get “Filtered offsite request” messages. How can I fix them?

Those messages (logged with `DEBUG` level) don’t necessarily mean there is a problem, so you may not need to fix them.

Those message are thrown by the Offsite Spider Middleware, which is a spider middleware (enabled by default) whose purpose is to filter out requests to domains outside the ones covered by the spider.

For more info see: `OffsiteMiddleware`.

5.1.16 What is the recommended way to deploy a Scrapy crawler in production?

See *Scrapy Service* (`scrapyd`).

5.1.17 Can I use JSON for large exports?

It’ll depend on how large your output is. See *this warning* in `JsonItemExporter` documentation.

5.1.18 Can I return (Twisted) deferreds from signal handlers?

Some signals support returning deferreds from their handlers, others don’t. See the *Built-in signals reference* to know which ones.

5.1.19 What does the response status code 999 means?

999 is a custom reponse status code used by Yahoo sites to throttle requests. Try slowing down the crawling speed by using a download delay of 2 (or higher) in your spider:

```
class MySpider(CrawlSpider):  
    name = 'myspider'  
  
    DOWNLOAD_DELAY = 2  
  
    # [ ... rest of the spider code ... ]
```

Or by setting a global download delay in your project with the `DOWNLOAD_DELAY` setting.

5.1.20 Can I call `pdb.set_trace()` from my spiders to debug them?

Yes, but you can also use the Scrapy shell which allows you too quickly analyze (and even modify) the response being processed by your spider, which is, quite often, more useful than plain old `pdb.set_trace()`.

For more info see *Invoking the shell from spiders to inspect responses*.

5.1.21 Simplest way to dump all my scraped items into a JSON/CSV/XML file?

To dump into a JSON file:

```
scrapy crawl myspider -o items.json -t json
```

To dump into a CSV file:

```
scrapy crawl myspider -o items.csv -t csv
```

To dump into a XML file:

```
scrapy crawl myspider -o items.xml -t xml
```

For more information see *Feed exports*

5.1.22 What's this huge cryptic `__VIEWSTATE` parameter used in some forms?

The `__VIEWSTATE` parameter is used in sites built with ASP.NET/VB.NET. For more info on how it works see [this page](#). Also, here's an [example spider](#) which scrapes one of these sites.

5.1.23 What's the best way to parse big XML/CSV data feeds?

Parsing big feeds with XPath selectors can be problematic since they need to build the DOM of the entire feed in memory, and this can be quite slow and consume a lot of memory.

In order to avoid parsing all the entire feed at once in memory, you can use the functions `xmliter` and `csviter` from `scrapy.utils.iterators` module. In fact, this is what the feed spiders (see *Spiders*) use under the cover.

5.1.24 Does Scrapy manage cookies automatically?

Yes, Scrapy receives and keeps track of cookies sent by servers, and sends them back on subsequent requests, like any regular web browser does.

For more info see *Requests and Responses* and *CookiesMiddleware*.

5.1.25 How can I see the cookies being sent and received from Scrapy?

Enable the `COOKIES_DEBUG` setting.

5.1.26 How can I instruct a spider to stop itself?

Raise the `CloseSpider` exception from a callback. For more info see: `CloseSpider`.

5.1.27 How can I prevent my Scrapy bot from getting banned?

See *Avoiding getting banned*.

5.1.28 Should I use spider arguments or settings to configure my spider?

Both *spider arguments* and *settings* can be used to configure your spider. There is no strict rule that mandates to use one or the other, but settings are more suited for parameters that, once set, don't change much, while spider arguments are meant to change more often, even on each spider run and sometimes are required for the spider to run at all (for example, to set the start url of a spider).

To illustrate with an example, assuming you have a spider that needs to log into a site to scrape data, and you only want to scrape data from a certain section of the site (which varies each time). In that case, the credentials to log in would be settings, while the url of the section to scrape would be a spider argument.

5.2 Debugging Spiders

This document explains the most common techniques for debugging spiders. Consider the following scrapy spider below:

```
class MySpider(BaseSpider):
    name = 'myspider'
    start_urls = (
        'http://example.com/page1',
        'http://example.com/page2',
    )

    def parse(self, response):
        # collect 'item_urls'
        for item_url in item_urls:
            yield Request(url=item_url, callback=self.parse_item)

    def parse_item(self, response):
        item = MyItem()
        # populate 'item' fields
        yield Request(url=item_details_url, meta={'item': item},
            callback=self.parse_details)

    def parse_details(self, response):
        item = response.meta['item']
        # populate more 'item' fields
        return item
```

Basically this is a simple spider which parses two pages of items (the start_urls). Items also have a details page with additional information, so we use the meta functionality of `Request` to pass a partially populated item.

5.2.1 Parse Command

The most basic way of checking the output of your spider is to use the `parse` command. It allows to check the behaviour of different parts of the spider at the method level. It has the advantage of being flexible and simple to use, but does not allow debugging code inside a method.

In order to see the item scraped from a specific url:

```
$ scrapy parse --spider=mypspider -c parse_item -d 2 <item_url>
[ ... scrapy log lines crawling example.com spider ... ]

>>> STATUS DEPTH LEVEL 2 <<<
# Scraped Items -----
[{'url': <item_url>}]

# Requests -----
[]
```

Using the `--verbose` or `-v` option we can see the status at each depth level:

```
$ scrapy parse --spider=mypspider -c parse_item -d 2 -v <item_url>
[ ... scrapy log lines crawling example.com spider ... ]

>>> DEPTH LEVEL: 1 <<<
# Scraped Items -----
[]

# Requests -----
[<GET item_details_url>]

>>> DEPTH LEVEL: 2 <<<
# Scraped Items -----
[{'url': <item_url>}]

# Requests -----
[]
```

Checking items scraped from a single `start_url`, can also be easily achieved using:

```
$ scrapy parse --spider=mypspider -d 3 'http://example.com/page1'
```

5.2.2 Scrapy Shell

While the `parse` command is very useful for checking behaviour of a spider, it is of little help to check what happens inside a callback, besides showing the response received and the output. How to debug the situation when `parse_details` sometimes receives no item?

Fortunately, the `shell` is your bread and butter in this case (see *Invoking the shell from spiders to inspect responses*):

```
from scrapy.shell import inspect_response

def parse_details(self, response):
    item = response.meta.get('item', None)
    if item:
        # populate more 'item' fields
        return item
    else:
        inspect_response(response, self)
```

See also: *Invoking the shell from spiders to inspect responses*.

5.2.3 Open in browser

Sometimes you just want to see how a certain response looks in a browser, you can use the `open_in_browser` function for that. Here is an example of how you would use it:

```
from scrapy.utils.response import open_in_browser

def parse_details(self, response):
    if "item name" not in response.body:
        open_in_browser(response)
```

`open_in_browser` will open a browser with the response received by Scrapy at that point, adjusting the `base` tag so that images and styles are displayed properly.

5.2.4 Logging

Logging is another useful option for getting information about your spider run. Although not as convenient, it comes with the advantage that the logs will be available in all future runs should they be necessary again:

```
from scrapy import log

def parse_details(self, response):
    item = response.meta.get('item', None)
    if item:
        # populate more 'item' fields
        return item
    else:
        self.log('No item received for %s' % response.url,
                 level=log.WARNING)
```

For more information, check the [Logging](#) section.

5.3 Spiders Contracts

New in version 0.15.

Note: This is a new feature (introduced in Scrapy 0.15) and may be subject to minor functionality/API updates. Check the [release notes](#) to be notified of updates.

Testing spiders can get particularly annoying and while nothing prevents you from writing unit tests the task gets cumbersome quickly. Scrapy offers an integrated way of testing your spiders by the means of contracts.

This allows you to test each callback of your spider by hardcoding a sample url and check various constraints for how the callback processes the response. Each contract is prefixed with an `@` and included in the docstring. See the following example:

```
def parse(self, response):
    """ This function parses a sample response. Some contracts are mingled
    with this docstring.

    @url http://www.amazon.com/s?field-keywords=selfish+gene
    @returns items 1 16
    @returns requests 0 0
    @scrapes Title Author Year Price
    """
```

This callback is tested using three built-in contracts:

class scrapy.contracts.default.**UrlContract**

This contract (`@url`) sets the sample url used when checking other contract conditions for this spider. This contract is mandatory. All callbacks lacking this contract are ignored when running the checks:

```
@url url
```

class scrapy.contracts.default.**ReturnsContract**

This contract (`@returns`) sets lower and upper bounds for the items and requests returned by the spider. The upper bound is optional:

```
@returns item(s) | request(s) [min [max]]
```

class scrapy.contracts.default.**ScrapesContract**

This contract (`@scrapes`) checks that all the items returned by the callback have the specified fields:

```
@scrapes field_1 field_2 ...
```

Use the `check` command to run the contract checks.

5.3.1 Custom Contracts

If you find you need more power than the built-in scrapy contracts you can create and load your own contracts in the project by using the `SPIDER_CONTRACTS` setting:

```
SPIDER_CONTRACTS = {
    'myproject.contracts.ResponseCheck': 10,
    'myproject.contracts.ItemValidate': 10,
}
```

Each contract must inherit from `scrapy.contracts.Contract` and can override three methods:

class scrapy.contracts.**Contract** (*method*, *args)

Parameters

- **method** (*function*) – callback function to which the contract is associated
- **args** (*list*) – list of arguments passed into the docstring (whitespace separated)

adjust_request_args (*args*)

This receives a `dict` as an argument containing default arguments for `Request` object. Must return the same or a modified version of it.

pre_process (*response*)

This allows hooking in various checks on the response received from the sample request, before it's being passed to the callback.

post_process (*output*)

This allows processing the output of the callback. Iterators are converted listified before being passed to this hook.

Here is a demo contract which checks the presence of a custom header in the response received. Raise `scrapy.exceptions.ContractFail` in order to get the failures pretty printed:

```
from scrapy.contracts import Contract
from scrapy.exceptions import ContractFail

class HasHeaderContract(Contract):
    """ Demo contract which checks the presence of a custom header
```

```

    @has_header X-CustomHeader
    """

    name = 'has_header'

    def pre_process(self, response):
        for header in self.args:
            if header not in response.headers:
                raise ContractFail('X-CustomHeader not present')

```

5.4 Common Practices

This section documents common practices when using Scrapy. These are things that cover many topics and don't often fall into any other specific section.

5.4.1 Run Scrapy from a script

You can use the [API](#) to run Scrapy from a script, instead of the typical way of running Scrapy via `scrapy crawl`.

What follows is a working example of how to do that, using the [testspiders](#) project as example. Remember that Scrapy is built on top of the Twisted asynchronous networking library, so you need run it inside the Twisted reactor.

```

from twisted.internet import reactor
from scrapy.crawler import Crawler
from scrapy.settings import Settings
from scrapy import log
from testspiders.spiders.followall import FollowAllSpider

spider = FollowAllSpider(domain='scrapinghub.com')
crawler = Crawler(Settings())
crawler.configure()
crawler.crawl(spider)
crawler.start()
log.start()
reactor.run() # the script will block here

```

See Also:

[Twisted Reactor Overview](#).

5.4.2 Running multiple spiders in the same process

By default, Scrapy runs a single spider per process when you run `scrapy crawl`. However, Scrapy supports running multiple spiders per process using the [internal API](#).

Here is an example, using the [testspiders](#) project:

```

from twisted.internet import reactor
from scrapy.crawler import Crawler
from scrapy.settings import Settings
from scrapy import log
from testspiders.spiders.followall import FollowAllSpider

def setup_crawler(domain):

```

```
spider = FollowAllSpider(domain=domain)
crawler = Crawler(Settings())
crawler.configure()
crawler.crawl(spider)
crawler.start()

for domain in ['scrapinghub.com', 'insophia.com']:
    setup_crawler(domain)
log.start()
reactor.run()
```

See Also:

Run Scrapy from a script.

5.4.3 Distributed crawls

Scrapy doesn't provide any built-in facility for running crawls in a distribute (multi-server) manner. However, there are some ways to distribute crawls, which vary depending on how you plan to distribute them.

If you have many spiders, the obvious way to distribute the load is to setup many Scrapy instances and distribute spider runs among those.

If you instead want to run a single (big) spider through many machines, what you usually do is partition the urls to crawl and send them to each separate spider. Here is a concrete example:

First, you prepare the list of urls to crawl and put them into separate files/urls:

```
http://somedomain.com/urls-to-crawl/spider1/part1.list
http://somedomain.com/urls-to-crawl/spider1/part2.list
http://somedomain.com/urls-to-crawl/spider1/part3.list
```

Then you fire a spider run on 3 different Scrapy servers. The spider would receive a (spider) argument `part` with the number of the partition to crawl:

```
curl http://scrapy1.mycompany.com:6800/schedule.json -d project=myproject -d spider=spider1 -d part=1
curl http://scrapy2.mycompany.com:6800/schedule.json -d project=myproject -d spider=spider1 -d part=2
curl http://scrapy3.mycompany.com:6800/schedule.json -d project=myproject -d spider=spider1 -d part=3
```

5.4.4 Avoiding getting banned

Some websites implement certain measures to prevent bots from crawling them, with varying degrees of sophistication. Getting around those measures can be difficult and tricky, and may sometimes require special infrastructure. Please consider contacting [commercial support](#) if in doubt.

Here are some tips to keep in mind when dealing with these kind of sites:

- rotate your user agent from a pool of well-known ones from browsers (google around to get a list of them)
- disable cookies (see `COOKIES_ENABLED`) as some sites may use cookies to spot bot behaviour
- use download delays (2 or higher). See `DOWNLOAD_DELAY` setting.
- if possible, use [Google cache](#) to fetch pages, instead of hitting the sites directly
- use a pool of rotating IPs. For example, the free [Tor project](#) or paid services like [ProxyMesh](#)

If you are still unable to prevent your bot getting banned, consider contacting [commercial support](#).

5.5 Broad Crawls

Scrapy defaults are optimized for crawling specific sites. These sites are often handled by a single Scrapy spider, although this is not necessary or required (for example, there are generic spiders that handle any given site thrown at them).

In addition to this “focused crawl”, there is another common type of crawling which covers a large (potentially unlimited) number of domains, and is only limited by time or other arbitrary constraint, rather than stopping when the domain was crawled to completion or when there are no more requests to perform. These are called “broad crawls” and is the typical crawlers employed by search engines.

These are some common properties often found in broad crawls:

- they crawl many domains (often, unbounded) instead of a specific set of sites
- they don’t necessarily crawl domains to completion, because it would impractical (or impossible) to do so, and instead limit the crawl by time or number of pages crawled
- they are simpler in logic (as opposed to very complex spiders with many extraction rules) because data is often post-processed in a separate stage
- they crawl many domains concurrently, which allows them to achieve faster crawl speeds by not being limited by any particular site constraint (each site is crawled slowly to respect politeness, but many sites are crawled in parallel)

As said above, Scrapy default settings are optimized for focused crawls, not broad crawls. However, due to its asynchronous architecture, Scrapy is very well suited for performing fast broad crawls. This page summarize some things you need to keep in mind when using Scrapy for doing broad crawls, along with concrete suggestions of Scrapy settings to tune in order to achieve an efficient broad crawl.

5.5.1 Increase concurrency

Concurrency is the number of requests that are processed in parallel. There is a global limit and a per-domain limit.

The default global concurrency limit in Scrapy is not suitable for crawling many different domains in parallel, so you will want to increase it. How much to increase it will depend on how much CPU you crawler will have available. A good starting point is 100, but the best way to find out is by doing some trials and identifying at what concurrency your Scrapy process gets CPU bounded. For optimum performance, You should pick a concurrency where CPU usage is at 80-90%.

To increase the global concurrency use:

```
CONCURRENT_REQUESTS = 100
```

5.5.2 Reduce log level

When doing broad crawls you are often only interested in the crawl rates you get and any errors found. These stats are reported by Scrapy when using the `INFO` log level. In order to save CPU (and log storage requirements) you should not use `DEBUG` log level when performing large broad crawls in production. Using `DEBUG` level when developing your (broad) crawler may fine though.

To set the log level use:

```
LOG_LEVEL = 'INFO'
```

5.5.3 Disable cookies

Disable cookies unless you *really* need. Cookies are often not needed when doing broad crawls (search engine crawlers ignore them), and they improve performance by saving some CPU cycles and reducing the memory footprint of your Scrapy crawler.

To disable cookies use:

```
COOKIES_ENABLED = False
```

5.5.4 Disable retries

Retrying failed HTTP requests can slow down the crawls substantially, specially when sites causes are very slow (or fail) to respond, thus causing a timeout error which gets retried many times, unnecessarily, preventing crawler capacity to be reused for other domains.

To disable retries use:

```
RETRY_ENABLED = False
```

5.5.5 Reduce download timeout

Unless you are crawling from a very slow connection (which shouldn't be the case for broad crawls) reduce the download timeout so that stuck requests are discarded quickly and free up capacity to process the next ones.

To reduce the download timeout use:

```
DOWNLOAD_TIMEOUT = 15
```

5.5.6 Disable redirects

Consider disabling redirects, unless you are interested in following them. When doing broad crawls it's common to save redirects and resolve them when revisiting the site at a later crawl. This also help to keep the number of request constant per crawl batch, otherwise redirect loops may cause the crawler to dedicate too many resources on any specific domain.

To disable redirects use:

```
REDIRECT_ENABLED = False
```

5.6 Using Firefox for scraping

Here is a list of tips and advice on using Firefox for scraping, along with a list of useful Firefox add-ons to ease the scraping process.

5.6.1 Caveats with inspecting the live browser DOM

Since Firefox add-ons operate on a live browser DOM, what you'll actually see when inspecting the page source is not the original HTML, but a modified one after applying some browser clean up and executing Javascript code. Firefox, in particular, is known for adding `<tbody>` elements to tables. Scrapy, on the other hand, does not modify the original page HTML, so you won't be able to extract any data if you use `<tbody>` in your XPath expressions.

Therefore, you should keep in mind the following things when working with Firefox and XPath:

- Disable Firefox Javascript while inspecting the DOM looking for XPaths to be used in Scrapy
- Never use full XPath paths, use relative and clever ones based on attributes (such as `id`, `class`, `width`, etc) or any identifying features like `contains(@href, 'image')`.
- Never include `<tbody>` elements in your XPath expressions unless you really know what you're doing

5.6.2 Useful Firefox add-ons for scraping

Firebug

[Firebug](#) is a widely known tool among web developers and it's also very useful for scraping. In particular, its [Inspect Element](#) feature comes very handy when you need to construct the XPaths for extracting data because it allows you to view the HTML code of each page element while moving your mouse over it.

See [Using Firebug for scraping](#) for a detailed guide on how to use Firebug with Scrapy.

XPather

[XPather](#) allows you to test XPath expressions directly on the pages.

XPath Checker

[XPath Checker](#) is another Firefox add-on for testing XPaths on your pages.

Tamper Data

[Tamper Data](#) is a Firefox add-on which allows you to view and modify the HTTP request headers sent by Firefox. Firebug also allows to view HTTP headers, but not to modify them.

Firecookie

[Firecookie](#) makes it easier to view and manage cookies. You can use this extension to create a new cookie, delete existing cookies, see a list of cookies for the current site, manage cookies permissions and a lot more.

5.7 Using Firebug for scraping

Note: Google Directory, the example website used in this guide is no longer available as it [has been shut down by Google](#). The concepts in this guide are still valid though. If you want to update this guide to use a new (working) site, your contribution will be more than welcome!. See [Contributing to Scrapy](#) for information on how to do so.

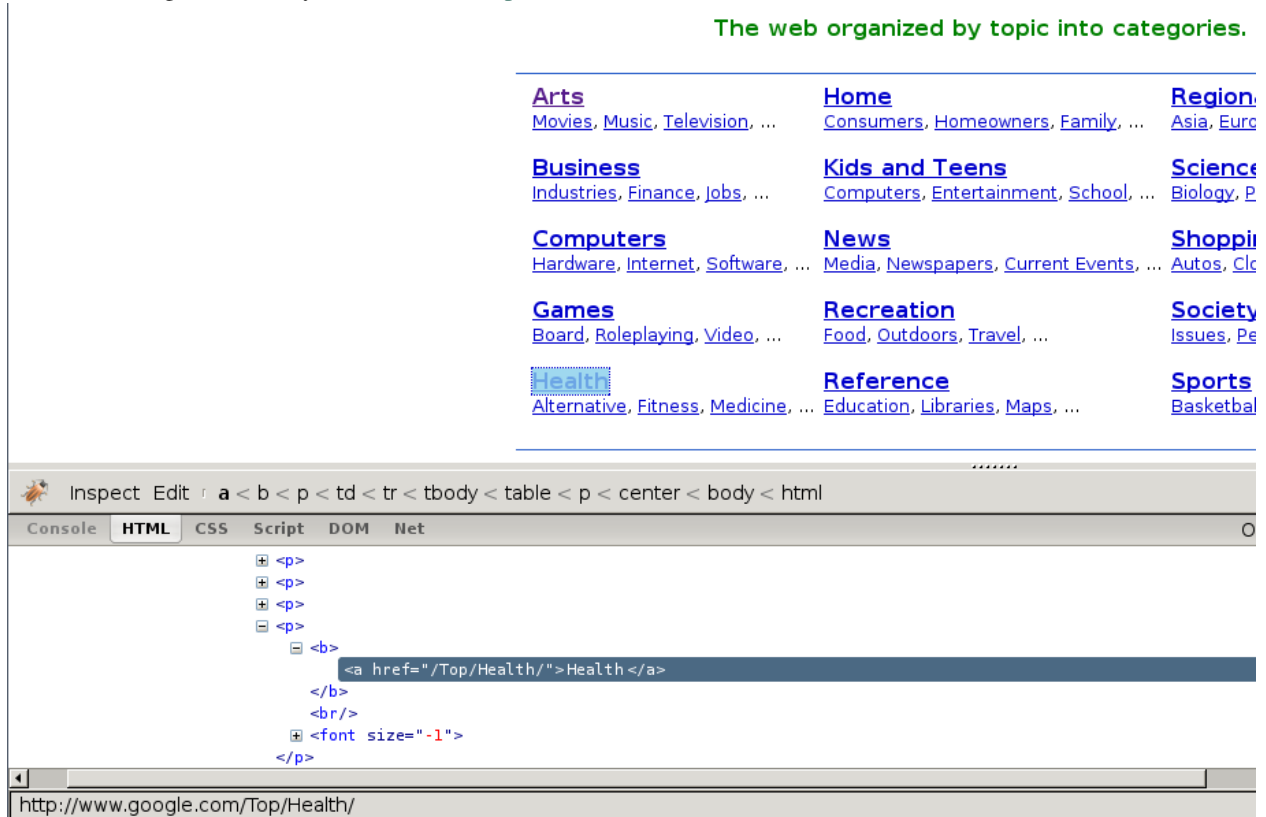
5.7.1 Introduction

This document explains how to use [Firebug](#) (a Firefox add-on) to make the scraping process easier and more fun. For other useful Firefox add-ons see [Useful Firefox add-ons for scraping](#). There are some caveats with using Firefox add-ons to inspect pages, see [Caveats with inspecting the live browser DOM](#).

In this example, we'll show how to use [Firebug](#) to scrape data from the [Google Directory](#), which contains the same data as the [Open Directory Project](#) used in the [tutorial](#) but with a different face.

Firebug comes with a very useful feature called [Inspect Element](#) which allows you to inspect the HTML code of the different page elements just by hovering your mouse over them. Otherwise you would have to search for the tags manually through the HTML body which can be a very tedious task.

In the following screenshot you can see the [Inspect Element](#) tool in action.



At first sight, we can see that the directory is divided in categories, which are also divided in subcategories.

However, it seems that there are more subcategories than the ones being shown in this page, so we'll keep looking:

[Environmental Health](#) (359) [Products and Shopping](#) (61) [Weight Loss](#) (357)
[Fitness](#) (961) [Professions](#) (1692) [Women's Health](#) (764)
[Healthcare Industry](#) (6380)

Related Categories:

- [Business > Business Services > Consulting > Medical and Life Sciences](#) (321)
- [Kids and Teens > Health](#) (1160)
- [Recreation > Humor > Medical](#) (26)
- [Science > Social Sciences > Communication > Health Communication](#) (3)
- [Shopping > Health](#) (7391)
- [Society > Issues > Health](#) (2592)

| Web Pages | Viewing in Google PageRank order | Vi |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------|----|
| WebMD - http://www.webmd.com/ A health resources for consumers, physicians, nurses, and educators. Includes news, chat forums, health quizzes and consumer product updates. | | |
| Health On the Net Foundation - http://www.hon.ch/ Guides lay persons and non-medical users and medical practitioners to useful and reliable online medical and health information. | | |
| BBC Health - http://www.bbc.co.uk/health/ Features current news plus archives, guides by subject, "Ask a Doctor" inquiry feature, a searchable conditions database, medical advice, and more. | | |
| AOL Health - http://www.aolhealth.com/ | | |

Inspect Edit **a** < font < td < tr < tbody < table < form < body < html

Console HTML CSS Script DOM Net

```

<table width="100%" cellspacing="0" cellpadding="1" border="0">
  <tbody>
    <tr valign="top">
      <td width="6%">
        <td>
          <font face="arial,sans-serif">
            <a href="http://www.webmd.com/">WebMD</a>
            <font size="-1" color="#6f6f6f">
              <br/>
              <font size="-1"> A health resources for consumers, physicians, nurses, and educators. Includes news, chat forums, health quizzes and consumer product updates. </font>
            </font>
          </td>
        </td>
      </tr>
    </tbody>
  </table>
  
```

Done

As expected, the subcategories contain links to other subcategories, and also links to actual websites, which is the purpose of the directory.

5.7.2 Getting links to follow

By looking at the category URLs we can see they share a pattern:

```
http://directory.google.com/Category/Subcategory/Another_Subcategory
```

Once we know that, we are able to construct a regular expression to follow those links. For example, the following one:

```
directory\.google\.com/[A-Z][a-zA-Z_/>+]$
```

So, based on that regular expression we can create the first crawling rule:

```
Rule(SgmlLinkExtractor(allow='directory.google.com/[A-Z][a-zA-Z_/>+',
    'parse_category',
    follow=True,
),
```

The `Rule` object instructs `CrawlSpider` based spiders how to follow the category links. `parse_category` will be a method of the spider which will process and extract data from those pages.

This is how the spider would look so far:

```
from scrapy.contrib.linkextractors.sgml import SgmlLinkExtractor
from scrapy.contrib.spiders import CrawlSpider, Rule
```

```

class GoogleDirectorySpider(CrawlSpider):
    name = 'directory.google.com'
    allowed_domains = ['directory.google.com']
    start_urls = ['http://directory.google.com/']

    rules = (
        Rule(SgmlLinkExtractor(allow='directory\.google\.com/[A-Z][a-zA-Z_/_]+$',
                                'parse_category', follow=True,
                                ),
        ),

    def parse_category(self, response):
        # write the category page data extraction code here
        pass

```

5.7.3 Extracting the data

Now we're going to write the code to extract data from those pages.

With the help of Firebug, we'll take a look at some page containing links to websites (say <http://directory.google.com/Top/Arts/Awards/>) and find out how we can extract those links using *XPath selectors*. We'll also use the *Scrapy shell* to test those XPath's and make sure they work as we expect.

Shopping > Health (7391)
Society > Issues > Health (2592)

| Web Pages | Viewing in Google PageRank order | View in alph |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------|--------------|
| WebMD - http://www.webmd.com/ A health resources for consumers, physicians, nurses, and educators. Includes news, chat forums, health quizzes and consumer product | | |
| Health On the Net Foundation - http://www.hon.ch/ Guides lay persons and non-medical users and medical practitioners to useful and reliable online medical and health information. Provide | | |
| BBC Health - http://www.bbc.co.uk/health/ Features current news plus archives, guides by subject, "Ask a Doctor" inquiry feature, a searchable conditions database, message board | | |
| AOL Health - http://www.aolhealth.com/ Find advice, information about diseases and drugs, fitness tips, and news items. | | |

Inspect Edit td < tr < tbody < table < form < body < html

Console HTML CSS Script DOM Net Options

```

<table width="100%" cellspacing="0" cellpadding="0" border="0">
  <table width="100%" cellspacing="0" cellpadding="1" border="0">
    <tbody>
      <tr valign="top">
        <td width="6%">
          <nobr>
            <a href="http://www.google.com/intl/en/dirhelp.html#pagerank">
          </nobr>
        </td>
      </tr>
    </tbody>
  </table>
</table>

```

Done

```

[<XPathSelector (a) xpath=//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td//a>,
<XPathSelector (a) xpath=//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td//a>,
<XPathSelector (a) xpath=//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td//a>,
<XPathSelector (a) xpath=//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td//a>,
<XPathSelector (a) xpath=//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td//a>,
<XPathSelector (a) xpath=//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td//a>]

In [5]: hxs.x('//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td//a').extract()
Out[5]:
[u'<a href="http://www.webmd.com/">WebMD</a>',
u'<a href="http://www.hon.ch/">Health On the Net Foundation</a>',
u'<a href="http://www.bbc.co.uk/health/">BBC Health</a>',
u'<a href="http://www.aolhealth.com/">AOL Health</a>',
u'<a href="http://www.intelihealth.com/">InteliHealth</a>',
u'<a href="http://www.judgehealth.org.uk/">Judge: Web Sites for Health</a>']

In [6]:

```

As you can see, the page markup is not very descriptive: the elements don't contain `id`, `class` or any attribute that clearly identifies them, so we'll use the ranking bars as a reference point to select the data to extract when we construct our XPath.

After using FireBug, we can see that each link is inside a `td` tag, which is itself inside a `tr` tag that also contains the link's ranking bar (in another `td`).

So we can select the ranking bar, then find its parent (the `tr`), and then finally, the link's `td` (which contains the data we want to scrape).

This results in the following XPath:

```
//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td//a
```

It's important to use the *Scrapy shell* to test these complex XPath expressions and make sure they work as expected.

Basically, that expression will look for the ranking bar's `td` element, and then select any `td` element who has a descendant `a` element whose `href` attribute contains the string `#pagerank`

Of course, this is not the only XPath, and maybe not the simpler one to select that data. Another approach could be, for example, to find any `font` tags that have that grey colour of the links,

Finally, we can write our `parse_category()` method:

```
def parse_category(self, response):
    hxs = HtmlXPathSelector(response)

    # The path to website links in directory page
    links = hxs.select('//td[descendant::a[contains(@href, "#pagerank")]]/following-sibling::td/font')

    for link in links:
        item = DirectoryItem()
        item['name'] = link.select('a/text()').extract()
        item['url'] = link.select('a/@href').extract()
        item['description'] = link.select('font[2]/text()').extract()
        yield item
```

Be aware that you may find some elements which appear in Firebug but not in the original HTML, such as the typical case of `<tbody>` elements.

or tags which Therefer in page HTML sources may on Firebug inspects the live DOM

5.8 Debugging memory leaks

In Scrapy, objects such as Requests, Responses and Items have a finite lifetime: they are created, used for a while, and finally destroyed.

From all those objects, the Request is probably the one with the longest lifetime, as it stays waiting in the Scheduler queue until it's time to process it. For more info see *Architecture overview*.

As these Scrapy objects have a (rather long) lifetime, there is always the risk of accumulating them in memory without releasing them properly and thus causing what is known as a “memory leak”.

To help debugging memory leaks, Scrapy provides a built-in mechanism for tracking objects references called *trackref*, and you can also use a third-party library called *Guppy* for more advanced memory debugging (see below for more info). Both mechanisms must be used from the *Telnet Console*.

5.8.1 Common causes of memory leaks

It happens quite often (sometimes by accident, sometimes on purpose) that the Scrapy developer passes objects referenced in Requests (for example, using the `meta` attribute or the request callback function) and that effectively bounds the lifetime of those referenced objects to the lifetime of the Request. This is, by far, the most common cause of memory leaks in Scrapy projects, and a quite difficult one to debug for newcomers.

In big projects, the spiders are typically written by different people and some of those spiders could be “leaking” and thus affecting the rest of the other (well-written) spiders when they get to run concurrently, which, in turn, affects the whole crawling process.

At the same time, it’s hard to avoid the reasons that cause these leaks without restricting the power of the framework, so we have decided not to restrict the functionality but provide useful tools for debugging these leaks, which quite often consist in an answer to the question: *which spider is leaking?*.

The leak could also come from a custom middleware, pipeline or extension that you have written, if you are not releasing the (previously allocated) resources properly. For example, if you’re allocating resources on `spider_opened` but not releasing them on `spider_closed`.

5.8.2 Debugging memory leaks with `trackref`

`trackref` is a module provided by Scrapy to debug the most common cases of memory leaks. It basically tracks the references to all live Requests, Responses, Item and Selector objects.

You can enter the telnet console and inspect how many objects (of the classes mentioned above) are currently alive using the `prefs()` function which is an alias to the `print_live_refs()` function:

```
telnet localhost 6023
```

```
>>> prefs()
Live References
```

| | | |
|---------------|-----|-----------------|
| ExampleSpider | 1 | oldest: 15s ago |
| HtmlResponse | 10 | oldest: 1s ago |
| XPathSelector | 2 | oldest: 0s ago |
| FormRequest | 878 | oldest: 7s ago |

As you can see, that report also shows the “age” of the oldest object in each class.

If you do have leaks, chances are you can figure out which spider is leaking by looking at the oldest request or response. You can get the oldest object of each class using the `get_oldest()` function like this (from the telnet console).

Which objects are tracked?

The objects tracked by `trackrefs` are all from these classes (and all its subclasses):

- `scrapy.http.Request`
- `scrapy.http.Response`
- `scrapy.item.Item`
- `scrapy.selector.XPathSelector`
- `scrapy.spider.BaseSpider`
- `scrapy.selector.document.Libxml2Document`

A real example

Let's see a concrete example of an hypothetical case of memory leaks.

Suppose we have some spider with a line similar to this one:

```
return Request("http://www.somenastyspider.com/product.php?pid=%d" % product_id,
               callback=self.parse, meta={'referer': response})
```

That line is passing a response reference inside a request which effectively ties the response lifetime to the requests' one, and that would definitely cause memory leaks.

Let's see how we can discover which one is the nasty spider (without knowing it a-priori, of course) by using the `trackref` tool.

After the crawler is running for a few minutes and we notice its memory usage has grown a lot, we can enter its telnet console and check the live references:

```
>>> prefs()
Live References

SomenastySpider          1   oldest: 15s ago
HtmlResponse             3890  oldest: 265s ago
XPathSelector            2    oldest: 0s ago
Request                  3878  oldest: 250s ago
```

The fact that there are so many live responses (and that they're so old) is definitely suspicious, as responses should have a relatively short lifetime compared to Requests. So let's check the oldest response:

```
>>> from scrapy.utils.trackref import get_oldest
>>> r = get_oldest('HtmlResponse')
>>> r.url
'http://www.somenastyspider.com/product.php?pid=123'
```

There it is. By looking at the URL of the oldest response we can see it belongs to the `somenastyspider.com` spider. We can now go and check the code of that spider to discover the nasty line that is generating the leaks (passing response references inside requests).

If you want to iterate over all objects, instead of getting the oldest one, you can use the `iter_all()` function:

```
>>> from scrapy.utils.trackref import iter_all
>>> [r.url for r in iter_all('HtmlResponse')]
['http://www.somenastyspider.com/product.php?pid=123',
 'http://www.somenastyspider.com/product.php?pid=584',
 ...]
```

Too many spiders?

If your project has too many spiders, the output of `prefs()` can be difficult to read. For this reason, that function has a `ignore` argument which can be used to ignore a particular class (and all its subclasses). For example, using:

```
>>> from scrapy.spider import BaseSpider
>>> prefs(ignore=BaseSpider)
```

Won't show any live references to spiders.

scrapy.utils.trackref module

Here are the functions available in the `trackref` module.

`class scrapy.utils.trackref.object_ref`

Inherit from this class (instead of `object`) if you want to track live instances with the `trackref` module.

`scrapy.utils.trackref.print_live_refs(class_name, ignore=NoneType)`

Print a report of live references, grouped by class name.

Parameters `ignore` (*class or classes tuple*) – if given, all objects from the specified class (or tuple of classes) will be ignored.

`scrapy.utils.trackref.get_oldest(class_name)`

Return the oldest object alive with the given class name, or `None` if none is found. Use `print_live_refs()` first to get a list of all tracked live objects per class name.

`scrapy.utils.trackref.iter_all(class_name)`

Return an iterator over all objects alive with the given class name, or `None` if none is found. Use `print_live_refs()` first to get a list of all tracked live objects per class name.

5.8.3 Debugging memory leaks with Guppy

`trackref` provides a very convenient mechanism for tracking down memory leaks, but it only keeps track of the objects that are more likely to cause memory leaks (Requests, Responses, Items, and Selectors). However, there are other cases where the memory leaks could come from other (more or less obscure) objects. If this is your case, and you can't find your leaks using `trackref`, you still have another resource: the [Guppy library](#).

If you use `setuptools`, you can install Guppy with the following command:

```
easy_install guppy
```

The telnet console also comes with a built-in shortcut (`hpy`) for accessing Guppy heap objects. Here's an example to view all Python objects available in the heap using Guppy:

```
>>> x = hpy.heap()
>>> x.bytype
Partition of a set of 297033 objects. Total size = 52587824 bytes.
  Index  Count    %      Size    % Cumulative  % Type
    0    22307    8 16423880   31 16423880   31 dict
    1   122285   41 12441544   24 28865424   55 str
    2    68346   23  5966696   11 34832120   66 tuple
    3      227    0  5836528   11 40668648   77 unicode
    4     2461    1  2222272    4 42890920   82 type
    5   16870    6  2024400    4 44915320   85 function
    6   13949    5  1673880    3 46589200   89 types.CodeType
    7   13422    5  1653104    3 48242304   92 list
    8    3735    1  1173680    2 49415984   94 _sre.SRE_Pattern
    9    1209    0   456936    1 49872920   95 scrapy.http.headers.Headers
<1676 more rows. Type e.g. '_.more' to view.>
```

You can see that most space is used by dicts. Then, if you want to see from which attribute those dicts are referenced, you could do:

```
>>> x.bytype[0].byvia
Partition of a set of 22307 objects. Total size = 16423880 bytes.
  Index  Count    %      Size    % Cumulative  % Referred Via:
    0   10982   49  9416336   57  9416336   57 '.__dict__'
    1    1820    8  2681504   16 12097840   74 '.__dict__', '.func_globals'
    2    3097   14  1122904    7 13220744   80
    3     990    4   277200    2 13497944   82 "['cookies']"
    4     987    4   276360    2 13774304   84 "['cache']"
    5     985    4   275800    2 14050104   86 "['meta']"
```

```

6      897      4      251160      2      14301264      87      '[2]'
7         1      0      196888      1      14498152      88      "['moduleDict']", "['modules']"
8       672      3      188160      1      14686312      89      "['cb_kwargs']"
9        27      0      155016      1      14841328      90      '[1]'
<333 more rows. Type e.g. '_.more' to view.>

```

As you can see, the Guppy module is very powerful but also requires some deep knowledge about Python internals. For more info about Guppy, refer to the [Guppy documentation](#).

5.8.4 Leaks without leaks

Sometimes, you may notice that the memory usage of your Scrapy process will only increase, but never decrease. Unfortunately, this could happen even though neither Scrapy nor your project are leaking memory. This is due to a (not so well) known problem of Python, which may not return released memory to the operating system in some cases. For more information on this issue see:

- [Python Memory Management](#)
- [Python Memory Management Part 2](#)
- [Python Memory Management Part 3](#)

The improvements proposed by Evan Jones, which are detailed in [this paper](#), got merged in Python 2.5, but this only reduces the problem, it doesn't fix it completely. To quote the paper:

Unfortunately, this patch can only free an arena if there are no more objects allocated in it anymore. This means that fragmentation is a large issue. An application could have many megabytes of free memory, scattered throughout all the arenas, but it will be unable to free any of it. This is a problem experienced by all memory allocators. The only way to solve it is to move to a compacting garbage collector, which is able to move objects in memory. This would require significant changes to the Python interpreter.

This problem will be fixed in future Scrapy releases, where we plan to adopt a new process model and run spiders in a pool of recyclable sub-processes.

5.9 Downloading Item Images

Scrapy provides an *item pipeline* for downloading images attached to a particular item, for example, when you scrape products and also want to download their images locally.

This pipeline, called the Images Pipeline and implemented in the `ImagesPipeline` class, provides a convenient way for downloading and storing images locally with some additional features:

- Convert all downloaded images to a common format (JPG) and mode (RGB)
- Avoid re-downloading images which were downloaded recently
- Thumbnail generation
- Check images width/height to make sure they meet a minimum constraint

This pipeline also keeps an internal queue of those images which are currently being scheduled for download, and connects those items that arrive containing the same image, to that queue. This avoids downloading the same image more than once when it's shared by several items.

The [Python Imaging Library](#) is used for thumbnailing and normalizing images to JPEG/RGB format, so you need to install that library in order to use the images pipeline.

5.9.1 Using the Images Pipeline

The typical workflow, when using the `ImagesPipeline` goes like this:

1. In a Spider, you scrape an item and put the URLs of its images into a `image_urls` field.
2. The item is returned from the spider and goes to the item pipeline.
3. When the item reaches the `ImagesPipeline`, the URLs in the `image_urls` field are scheduled for download using the standard Scrapy scheduler and downloader (which means the scheduler and downloader middlewares are reused), but with a higher priority, processing them before other pages are scraped. The item remains “locked” at that particular pipeline stage until the images have finish downloading (or fail for some reason).
4. When the images are downloaded another field (`images`) will be populated with the results. This field will contain a list of dicts with information about the images downloaded, such as the downloaded path, the original scraped url (taken from the `image_urls` field) , and the image checksum. The images in the list of the `images` field will retain the same order of the original `image_urls` field. If some image failed downloading, an error will be logged and the image won’t be present in the `images` field.

5.9.2 Usage example

In order to use the image pipeline you just need to *enable it* and define an item with the `image_urls` and `images` fields:

```
from scrapy.item import Item

class MyItem(Item):

    # ... other item fields ...
    image_urls = Field()
    images = Field()
```

If you need something more complex and want to override the custom images pipeline behaviour, see *Implementing your custom Images Pipeline*.

5.9.3 Enabling your Images Pipeline

To enable your images pipeline you must first add it to your project `ITEM_PIPELINES` setting:

```
ITEM_PIPELINES = ['scrapy.contrib.pipeline.images.ImagesPipeline']
```

And set the `IMAGES_STORE` setting to a valid directory that will be used for storing the downloaded images. Otherwise the pipeline will remain disabled, even if you include it in the `ITEM_PIPELINES` setting.

For example:

```
IMAGES_STORE = '/path/to/valid/dir'
```

5.9.4 Images Storage

File system is currently the only officially supported storage, but there is also (undocumented) support for [Amazon S3](#).

File system storage

The images are stored in files (one per image), using a [SHA1 hash](#) of their URLs for the file names.

For example, the following image URL:

```
http://www.example.com/image.jpg
```

Whose *SHA1 hash* is:

```
3afec3b4765f8f0a07b78f98c07b83f013567a0a
```

Will be downloaded and stored in the following file:

```
<IMAGES_STORE>/full/3afec3b4765f8f0a07b78f98c07b83f013567a0a.jpg
```

Where:

- `<IMAGES_STORE>` is the directory defined in `IMAGES_STORE` setting
- `full` is a sub-directory to separate full images from thumbnails (if used). For more info see *Thumbnail generation*.

5.9.5 Additional features

Image expiration

The Image Pipeline avoids downloading images that were downloaded recently. To adjust this retention delay use the `IMAGES_EXPIRES` setting, which specifies the delay in number of days:

```
# 90 days of delay for image expiration
IMAGES_EXPIRES = 90
```

Thumbnail generation

The Images Pipeline can automatically create thumbnails of the downloaded images. In order use this feature, you must set `IMAGES_THUMBS` to a dictionary where the keys are the thumbnail names and the values are their dimensions.

For example:

```
IMAGES_THUMBS = {
    'small': (50, 50),
    'big': (270, 270),
}
```

When you use this feature, the Images Pipeline will create thumbnails of the each specified size with this format:

```
<IMAGES_STORE>/thumbs/<size_name>/<image_id>.jpg
```

Where:

- `<size_name>` is the one specified in the `IMAGES_THUMBS` dictionary keys (`small`, `big`, etc)
- `<image_id>` is the [SHA1 hash](#) of the image url

Example of image files stored using `small` and `big` thumbnail names:

```
<IMAGES_STORE>/full/63bbfea82b8880ed33cdb762aa11fab722a90a24.jpg
<IMAGES_STORE>/thumbs/small/63bbfea82b8880ed33cdb762aa11fab722a90a24.jpg
<IMAGES_STORE>/thumbs/big/63bbfea82b8880ed33cdb762aa11fab722a90a24.jpg
```

The first one is the full image, as downloaded from the site.

Filtering out small images

You can drop images which are too small, by specifying the minimum allowed size in the `IMAGES_MIN_HEIGHT` and `IMAGES_MIN_WIDTH` settings.

For example:

```
IMAGES_MIN_HEIGHT = 110
IMAGES_MIN_WIDTH = 110
```

Note: these size constraints don't affect thumbnail generation at all.

By default, there are no size constraints, so all images are processed.

5.9.6 Implementing your custom Images Pipeline

Here are the methods that you should override in your custom Images Pipeline:

```
class scrapy.contrib.pipeline.images.ImagesPipeline
```

```
    get_media_requests(item, info)
```

As seen on the workflow, the pipeline will get the URLs of the images to download from the item. In order to do this, you must override the `get_media_requests()` method and return a Request for each image URL:

```
def get_media_requests(self, item, info):
    for image_url in item['image_urls']:
        yield Request(image_url)
```

Those requests will be processed by the pipeline and, when they have finished downloading, the results will be sent to the `item_completed()` method, as a list of 2-element tuples. Each tuple will contain (success, image_info_or_failure) where:

- success is a boolean which is True if the image was downloaded successfully or False if it failed for some reason
- image_info_or_error is a dict containing the following keys (if success is True) or a `TwistedFailure` if there was a problem.
 - url - the url where the image was downloaded from. This is the url of the request returned from the `get_media_requests()` method.
 - path - the path (relative to `IMAGES_STORE`) where the image was stored
 - checksum - a `MD5` hash of the image contents

The list of tuples received by `item_completed()` is guaranteed to retain the same order of the requests returned from the `get_media_requests()` method.

Here's a typical value of the results argument:

```
[ (True,
  {'checksum': '2b00042f7481c7b056c4b410d28f33cf',
   'path': 'full/7d97e98f8af710c7e7fe703abc8f639e0ee507c4.jpg',
   'url': 'http://www.example.com/images/product1.jpg'}),
  (True,
   {'checksum': 'b9628c4ab9b595f72f280b90c4fd093d',
    'path': 'full/1ca5879492b8fd606df1964ea3c1e2f4520f076f.jpg',
    'url': 'http://www.example.com/images/product2.jpg'}),
  (False,
   Failure(...))]
```

By default the `get_media_requests()` method returns `None` which means there are no images to download for the item.

item_completed (results, items, info)

The `ImagesPipeline.item_completed()` method called when all image requests for a single item have completed (either finished downloading, or failed for some reason).

The `item_completed()` method must return the output that will be sent to subsequent item pipeline stages, so you must return (or drop) the item, as you would in any pipeline.

Here is an example of the `item_completed()` method where we store the downloaded image paths (passed in results) in the `image_paths` item field, and we drop the item if it doesn't contain any images:

```
from scrapy.exceptions import DropItem

def item_completed(self, results, item, info):
    image_paths = [x['path'] for ok, x in results if ok]
    if not image_paths:
        raise DropItem("Item contains no images")
    item['image_paths'] = image_paths
    return item
```

By default, the `item_completed()` method returns the item.

5.9.7 Custom Images pipeline example

Here is a full example of the Images Pipeline whose methods are exemplified above:

```
from scrapy.contrib.pipeline.images import ImagesPipeline
from scrapy.exceptions import DropItem
from scrapy.http import Request

class MyImagesPipeline(ImagesPipeline):

    def get_media_requests(self, item, info):
        for image_url in item['image_urls']:
            yield Request(image_url)

    def item_completed(self, results, item, info):
        image_paths = [x['path'] for ok, x in results if ok]
        if not image_paths:
            raise DropItem("Item contains no images")
        item['image_paths'] = image_paths
        return item
```

5.10 Ubuntu packages

New in version 0.10. [Scrapinghub](#) publishes apt-gettable packages which are generally fresher than those in Ubuntu, and more stable too since they're continuously built from [Github repo](#) (master & stable branches) and so they contain the latest bug fixes.

To use the packages, just add the following line to your `/etc/apt/sources.list`, and then run `aptitude update` and `aptitude install scrapy-0.13`:

```
deb http://archive.scrapy.org/ubuntu DISTRO main
```

Replacing `DISTRO` with the name of your Ubuntu release, which you can get with command:

```
lsb_release -cs
```

Supported Ubuntu releases are: karmic, lucid, maverick, natty, oneiric, precise.

For Ubuntu Precise (12.04):

```
deb http://archive.scrapy.org/ubuntu precise main
```

For Ubuntu Oneiric (11.10):

```
deb http://archive.scrapy.org/ubuntu oneiric main
```

For Ubuntu Natty (11.04):

```
deb http://archive.scrapy.org/ubuntu natty main
```

For Ubuntu Maverick (10.10):

```
deb http://archive.scrapy.org/ubuntu maverick main
```

For Ubuntu Lucid (10.04):

```
deb http://archive.scrapy.org/ubuntu lucid main
```

For Ubuntu Karmic (9.10):

```
deb http://archive.scrapy.org/ubuntu karmic main
```

Warning: Please note that these packages are updated frequently, and so if you find you can't download the packages, try updating your apt package lists first, e.g., with `apt-get update` or `aptitude update`.

The public GPG key used to sign these packages can be imported into you APT keyring as follows:

```
curl -s http://archive.scrapy.org/ubuntu/archive.key | sudo apt-key add -
```

5.11 Scrapy Service (scrapyd)

New in version 0.10. Scrapy comes with a built-in service, called “Scrapyd”, which allows you to deploy (aka. upload) your projects and control their spiders using a JSON web service.

5.11.1 Projects and versions

Scrapyd can manage multiple projects and each project can have multiple versions uploaded, but only the latest one will be used for launching new spiders.

A common (and useful) convention to use for the version name is the revision number of the version control tool you're using to track your Scrapy project code. For example: `r23`. The versions are not compared alphabetically but using a smarter algorithm (the same `distutils` uses) so `r10` compares greater to `r9`, for example.

5.11.2 How Scrapyd works

Scrapyd is an application (typically run as a daemon) that continually polls for spiders that need to run.

When a spider needs to run, a process is started to crawl the spider:

```
scrapy crawl myspider
```

Scrapyd also runs multiple processes in parallel, allocating them in a fixed number of slots given by the `max_proc` and `max_proc_per_cpu` options, starting as many processes as possible to handle the load.

In addition to dispatching and managing processes, Scrapyd provides a *JSON web service* to upload new project versions (as eggs) and schedule spiders. This feature is optional and can be disabled if you want to implement your own custom Scrapyd. The components are pluggable and can be changed, if you're familiar with the *Twisted Application Framework* which Scrapyd is implemented in.

Starting from 0.11, Scrapyd also provides a minimal *web interface*.

5.11.3 Starting Scrapyd

Scrapyd is implemented using the standard *Twisted Application Framework*. To start the service, use the `extras/scrapyd.tac` file provided in the Scrapy distribution, like this:

```
twistd -ny extras/scrapyd.tac
```

That should get your Scrapyd started.

Or, if you want to start Scrapyd from inside a Scrapy project you can use the `server` command, like this:

```
scrapy server
```

5.11.4 Installing Scrapyd

How to deploy Scrapyd on your servers depends on the platform you're using. Scrapy comes with Ubuntu packages for Scrapyd ready for deploying it as a system service, to ease the installation and administration, but you can create packages for other distribution or operating systems (including Windows). If you do so, and want to contribute them, send a message to scrapy-developers@googlegroups.com and say hi. The community will appreciate it.

Installing Scrapyd in Ubuntu

When deploying Scrapyd, it's very useful to have a version already packaged for your system. For this reason, Scrapyd comes with Ubuntu packages ready to use in your Ubuntu servers.

So, if you plan to deploy Scrapyd on a Ubuntu server, just add the Ubuntu repositories as described in *Ubuntu packages* and then run:

```
aptitude install scrapyd-X.YY
```

Where `X.YY` is the Scrapy version, for example: `0.14`.

This will install Scrapyd in your Ubuntu server creating a `scrapy` user which Scrapyd will run as. It will also create some directories and files that are listed below:

`/etc/scrapyd`

Scrapyd configuration files. See *Scrapy Configuration file*.

`/var/log/scrapyd/scrapyd.log`

Scrapyd main log file.

`/var/log/scrapyd/scrapyd.out`

The standard output captured from Scrapyd process and any sub-process spawned from it.

`/var/log/scrapyd/scrapyd.err`

The standard error captured from Scrapyd and any sub-process spawned from it. Remember to check this file if you're having problems, as the errors may not get logged to the `scrapyd.log` file.

`/var/log/scrapyd/project`

Besides the main service log file, Scrapyd stores one log file per crawling process in:

```
/var/log/scrapyd/PROJECT/SPIDER/ID.log
```

Where `ID` is a unique id for the run.

`/var/lib/scrapyd/`

Directory used to store data files (uploaded eggs and spider queues).

5.11.5 Scrapy Configuration file

Scrapyd searches for configuration files in the following locations, and parses them in order with the latest ones taking more priority:

- `/etc/scrapyd/scrapyd.conf` (Unix)
- `c:\scrapyd\scrapyd.conf` (Windows)
- `/etc/scrapyd/conf.d/*` (in alphabetical order, Unix)
- `scrapyd.conf`

The configuration file supports the following options (see default values in the *example*).

http_port

The TCP port where the HTTP JSON API will listen. Defaults to 6800.

bind_address

The IP address where the HTTP JSON API will listen. Defaults to 0.0.0.0 (all)

max_proc

The maximum number of concurrent Scrapy process that will be started. If unset or 0 it will use the number of cpus available in the system multiplied by the value in `max_proc_per_cpu` option. Defaults to 0.

max_proc_per_cpu

The maximum number of concurrent Scrapy process that will be started per cpu. Defaults to 4.

debug

Whether debug mode is enabled. Defaults to `off`. When debug mode is enabled the full Python traceback will be returned (as plain text responses) when there is an error processing a JSON API call.

eggs_dir

The directory where the project eggs will be stored.

dbs_dir

The directory where the project databases will be stored (this includes the spider queues).

logs_dir

The directory where the Scrapy logs will be stored. If you want to disable storing logs set this option empty, like this:

```
logs_dir =
```

items_dir

New in version 0.15. The directory where the Scrapy items will be stored. If you want to disable storing feeds of scraped items (perhaps, because you use a database or other storage) set this option empty, like this:

```
items_dir =
```

jobs_to_keep

New in version 0.15. The number of finished jobs to keep per spider. Defaults to 5. This includes logs and items.

This setting was named `logs_to_keep` in previous versions.

runner

The module that will be used for launching sub-processes. You can customize the Scrapy processes launched from Scrapyd by using your own module.

application

A function that returns the (Twisted) Application object to use. This can be used if you want to extend Scrapyd by adding and removing your own components and services.

For more info see [Twisted Application Framework](#)

Example configuration file

Here is an example configuration file with all the defaults:

```
[scrapyd]
eggs_dir      = eggs
logs_dir      = logs
items_dir     = items
jobs_to_keep  = 5
dbs_dir       = dbs
max_proc      = 0
max_proc_per_cpu = 4
finished_to_keep = 100
http_port     = 6800
debug        = off
runner        = scrapyd.runner
application   = scrapyd.app.application
launcher      = scrapyd.launcher.Launcher

[services]
schedule.json    = scrapyd.webservice.Schedule
cancel.json      = scrapyd.webservice.Cancel
addversion.json  = scrapyd.webservice.AddVersion
listprojects.json = scrapyd.webservice.ListProjects
listversions.json = scrapyd.webservice.ListVersions
listspiders.json = scrapyd.webservice.ListSpiders
delproject.json  = scrapyd.webservice.DeleteProject
delversion.json  = scrapyd.webservice.DeleteVersion
listjobs.json    = scrapyd.webservice.ListJobs
```

5.11.6 Deploying your project

Deploying your project into a Scrapyd server typically involves two steps:

1. building a [Python egg](#) of your project. This is called “eggifying” your project. You’ll need to install [setuptools](#) for this. See [Egg caveats](#) below.
2. uploading the egg to the Scrapyd server

The simplest way to deploy your project is by using the [deploy](#) command, which automates the process of building the egg uploading it using the Scrapyd HTTP JSON API.

The [deploy](#) command supports multiple targets (Scrapyd servers that can host your project) and each target supports multiple projects.

Each time you deploy a new version of a project, you can name it for later reference.

Show and define targets

To see all available targets type:

```
scrapy deploy -l
```

This will return a list of available targets and their URLs. For example:

```
scrapyd                http://localhost:6800/
```

You can define targets by adding them to your project's `scrapy.cfg` file, or any other supported location like `~/.scrapy.cfg`, `/etc/scrapy.cfg`, or `c:\scrapy\scrapy.cfg` (in Windows).

Here's an example of defining a new target `scrapyd2` with restricted access through HTTP basic authentication:

```
[deploy:scrapyd2]
url = http://scrapyd.mydomain.com/api/scrapyd/
username = john
password = secret
```

Note: The `deploy` command also supports `netrc` for getting the credentials.

Now, if you type `scrapy deploy -l` you'll see:

```
scrapyd                http://localhost:6800/
scrapyd2               http://scrapyd.mydomain.com/api/scrapyd/
```

See available projects

To see all available projects in a specific target use:

```
scrapy deploy -L scrapyd
```

It would return something like this:

```
project1
project2
```

Deploying a project

Finally, to deploy your project use:

```
scrapy deploy scrapyd -p project1
```

This will eggify your project and upload it to the target, printing the JSON response returned from the Scrapy server. If you have a `setup.py` file in your project, that one will be used. Otherwise a `setup.py` file will be created automatically (based on a simple template) that you can edit later.

After running that command you will see something like this, meaning your project was uploaded successfully:

```
Deploying myproject-1287453519 to http://localhost:6800/addversion.json
Server response (200):
{"status": "ok", "spiders": ["spider1", "spider2"]}
```

By default `scrapy deploy` uses the current timestamp for generating the project version, as you can see in the output above. However, you can pass a custom version with the `--version` option:

```
scrapy deploy scrapyd -p project1 --version 54
```

Also, if you use Mercurial for tracking your project source code, you can use `HG` for the version which will be replaced by the current Mercurial revision, for example `r382`:

```
scrapy deploy scrapyd -p project1 --version HG
```

And, if you use Git for tracking your project source code, you can use `GIT` for the version which will be replaced by the SHA1 of current Git revision, for example `b0582849179d1de7bd86eaa7201ea3cda4b5651f`:

```
scrapy deploy scrapyd -p project1 --version GIT
```

Support for other version discovery sources may be added in the future.

Finally, if you don't want to specify the target, project and version every time you run `scrapy deploy` you can define the defaults in the `scrapy.cfg` file. For example:

```
[deploy]
url = http://scrapyd.mydomain.com/api/scrapyd/
username = john
password = secret
project = project1
version = HG
```

This way, you can deploy your project just by using:

```
scrapy deploy
```

Local settings

Sometimes, while your working on your projects, you may want to override your certain settings with certain local settings that shouldn't be deployed to Scrapyd, but only used locally to develop and debug your spiders.

One way to deal with this is to have a `local_settings.py` at the root of your project (where the `scrapy.cfg` file resides) and add these lines to the end of your project settings:

```
try:
    from local_settings import *
except ImportError:
    pass
```

`scrapy deploy` won't deploy anything outside the project module so the `local_settings.py` file won't be deployed.

Here's the directory structure, to illustrate:

```
scrapy.cfg
local_settings.py
myproject/
  __init__.py
  settings.py
  spiders/
  ...
```

Egg caveats

There are some things to keep in mind when building eggs of your Scrapy project:

- make sure no local development settings are included in the egg when you build it. The `find_packages` function may be picking up your custom settings. In most cases you want to upload the egg with the default project settings.
- you shouldn't use `__file__` in your project code as it doesn't play well with eggs. Consider using `pkgutil.get_data()` instead.
- be careful when writing to disk in your project (in any spider, extension or middleware) as Scrapy will probably run with a different user which may not have write access to certain directories. If you can, avoid writing to disk and always use `tempfile` for temporary files.

5.11.7 Scheduling a spider run

To schedule a spider run:

```
$ curl http://localhost:6800/schedule.json -d project=myproject -d spider=spider2
{"status": "ok", "jobid": "26d1b1a6d6f111e0be5c001e648c57f8"}
```

For more resources see: [JSON API reference](#) for more available resources.

5.11.8 Web Interface

New in version 0.11. Scrapy comes with a minimal web interface (for monitoring running processes and accessing logs) which can be accessed at <http://localhost:6800/>

5.11.9 JSON API reference

The following section describes the available resources in Scrapy JSON API.

addversion.json

Add a version to a project, creating the project if it doesn't exist.

- Supported Request Methods: POST
- Parameters:
 - `project` (string, required) - the project name
 - `version` (string, required) - the project version
 - `egg` (file, required) - a Python egg containing the project's code

Example request:

```
$ curl http://localhost:6800/addversion.json -F project=myproject -F version=r23 -F egg=@myproject.egg
```

Example response:

```
{"status": "ok", "spiders": 3}
```

schedule.json

Schedule a spider run (also known as a job), returning the job id.

- Supported Request Methods: POST
- Parameters:
 - `project` (string, required) - the project name
 - `spider` (string, required) - the spider name
 - `setting` (string, optional) - a scrapy setting to use when running the spider
 - any other parameter is passed as spider argument

Example request:

```
$ curl http://localhost:6800/schedule.json -d project=myproject -d spider=somespider
```

Example response:

```
{"status": "ok", "jobid": "6487ec79947edab326d6db28a2d86511e8247444"}
```

Example request passing a spider argument (`arg1`) and a setting (`DOWNLOAD_DELAY`):

```
$ curl http://localhost:6800/schedule.json -d project=myproject -d spider=somespider -d setting=DOWNN
```

cancel.json

New in version 0.15. Cancel a spider run (aka. job). If the job is pending, it will be removed. If the job is running, it will be terminated.

- Supported Request Methods: POST
- Parameters:
 - `project` (string, required) - the project name
 - `job` (string, required) - the job id

Example request:

```
$ curl http://localhost:6800/cancel.json -d project=myproject -d job=6487ec79947edab326d6db28a2d8651
```

Example response:

```
{"status": "ok", "prevstate": "running"}
```

listprojects.json

Get the list of projects uploaded to this Scrapy server.

- Supported Request Methods: GET
- Parameters: none

Example request:

```
$ curl http://localhost:6800/listprojects.json
```

Example response:

```
{"status": "ok", "projects": ["myproject", "otherproject"]}
```

listversions.json

Get the list of versions available for some project. The versions are returned in order, the last one is the currently used version.

- Supported Request Methods: GET
- Parameters:
 - project (string, required) - the project name

Example request:

```
$ curl http://localhost:6800/listversions.json?project=myproject
```

Example response:

```
{"status": "ok", "versions": ["r99", "r156"]}
```

listspiders.json

Get the list of spiders available in the last version of some project.

- Supported Request Methods: GET
- Parameters:
 - project (string, required) - the project name

Example request:

```
$ curl http://localhost:6800/listspiders.json?project=myproject
```

Example response:

```
{"status": "ok", "spiders": ["spider1", "spider2", "spider3"]}
```

listjobs.json

New in version 0.15. Get the list of pending, running and finished jobs of some project.

- Supported Request Methods: GET
- Parameters:
 - project (string, required) - the project name

Example request:

```
$ curl http://localhost:6800/listjobs.json?project=myproject
```

Example response:

```
{"status": "ok",  
  "pending": [{"id": "78391cc0fcaf11e1b0090800272a6d06", "spider": "spider1"}],  
  "running": [{"id": "422e608f9f28cef127b3d5ef93fe9399", "spider": "spider2"}],  
  "finished": [{"id": "2f16646cfcaf11e1b0090800272a6d06", "spider": "spider3", "start_time": "2012-09-"}]
```

Note: All job data is kept in memory and will be reset when the Scrapy service is restarted. See [issue 173](#).

delversion.json

Delete a project version. If there are no more versions available for a given project, that project will be deleted too.

- Supported Request Methods: POST
- Parameters:
 - `project` (string, required) - the project name
 - `version` (string, required) - the project version

Example request:

```
$ curl http://localhost:6800/delversion.json -d project=myproject -d version=r99
```

Example response:

```
{"status": "ok"}
```

delproject.json

Delete a project and all its uploaded versions.

- Supported Request Methods: POST
- Parameters:
 - `project` (string, required) - the project name

Example request:

```
$ curl http://localhost:6800/delproject.json -d project=myproject
```

Example response:

```
{"status": "ok"}
```

5.12 AutoThrottle extension

This is an extension for automatically throttling crawling speed based on load of both the Scrapy server and the website you are crawling.

5.12.1 Design goals

1. be nicer to sites instead of using default download delay of zero
2. automatically adjust scrapy to the optimum crawling speed, so the user doesn't have to tune the download delays and concurrent requests to find the optimum one. the user only needs to specify the maximum concurrent requests it allows, and the extension does the rest.

5.12.2 How it works

In Scrapy, the download latency is measured as the time elapsed between establishing the TCP connection and receiving the HTTP headers.

Note that these latencies are very hard to measure accurately in a cooperative multitasking environment because Scrapy may be busy processing a spider callback, for example, and unable to attend downloads. However, these latencies should still give a reasonable estimate of how busy Scrapy (and ultimately, the server) is, and this extension builds on that premise.

5.12.3 Throttling algorithm

This adjusts download delays and concurrency based on the following rules:

1. spiders always start with one concurrent request and a download delay of `AUTOTHROTTLE_START_DELAY`
2. when a response is received, the download delay is adjusted to the average of previous download delay and the latency of the response.
3. after `AUTOTHROTTLE_CONCURRENCY_CHECK_PERIOD` responses have passed, the average latency of this period is checked against the previous one and:
 - if the latency remained constant (within standard deviation limits), it is increased
 - if the latency has increased (beyond standard deviation limits) and the concurrency is higher than 1, the concurrency is decreased

Note: The AutoThrottle extension honours the standard Scrapy settings for concurrency and delay. This means that it will never set a download delay lower than `DOWNLOAD_DELAY` or a concurrency higher than `CONCURRENT_REQUESTS_PER_DOMAIN` (or `CONCURRENT_REQUESTS_PER_IP`, depending on which one you use).

5.12.4 Settings

The settings used to control the AutoThrottle extension are:

- `AUTOTHROTTLE_ENABLED`
- `AUTOTHROTTLE_START_DELAY`
- `AUTOTHROTTLE_CONCURRENCY_CHECK_PERIOD`
- `AUTOTHROTTLE_DEBUG`
- `DOWNLOAD_DELAY`
- `CONCURRENT_REQUESTS_PER_DOMAIN`
- `CONCURRENT_REQUESTS_PER_IP`

For more information see *Throttling algorithm*.

AUTOTHROTTLE_ENABLED

Default: `False`

Enables the AutoThrottle extension.

AUTOTHROTTLE_START_DELAY

Default: 5.0

The initial download delay (in seconds).

AUTOTHROTTLE_CONCURRENCY_CHECK_PERIOD

Default: 10

How many responses should pass to perform concurrency adjustments.

AUTOTHROTTLE_DEBUG

Default: False

Enable AutoThrottle debug mode which will display stats on every response received, so you can see how the throttling parameters are being adjusted in real time.

5.13 Jobs: pausing and resuming crawls

Sometimes, for big sites, it's desirable to pause crawls and be able to resume them later.

Scrapy supports this functionality out of the box by providing the following facilities:

- a scheduler that persists scheduled requests on disk
- a duplicates filter that persists visited requests on disk
- an extension that keeps some spider state (key/value pairs) persistent between batches

5.13.1 Job directory

To enable persistence support you just need to define a *job directory* through the `JOBDIR` setting. This directory will be for storing all required data to keep the state of a single job (ie. a spider run). It's important to note that this directory must not be shared by different spiders, or even different jobs/runs of the same spider, as it's meant to be used for storing the state of a *single* job.

5.13.2 How to use it

To start a spider with persistence supported enabled, run it like this:

```
scrapy crawl somespider -s JOBDIR=crawls/somespider-1
```

Then, you can stop the spider safely at any time (by pressing Ctrl-C or sending a signal), and resume it later by issuing the same command:

```
scrapy crawl somespider -s JOBDIR=crawls/somespider-1
```


5.13.3 Keeping persistent state between batches

Sometimes you'll want to keep some persistent spider state between pause/resume batches. You can use the `spider.state` attribute for that, which should be a dict. There's a built-in extension that takes care of serializing, storing and loading that attribute from the job directory, when the spider starts and stops.

Here's an example of a callback that uses the spider state (other spider code is omitted for brevity):

```
def parse_item(self, response):
    # parse item here
    self.state['items_count'] = self.state.get('items_count', 0) + 1
```

5.13.4 Persistence gotchas

There are a few things to keep in mind if you want to be able to use the Scrapy persistence support:

Cookies expiration

Cookies may expire. So, if you don't resume your spider quickly the requests scheduled may no longer work. This won't be an issue if you spider doesn't rely on cookies.

Request serialization

Requests must be serializable by the *pickle* module, in order for persistence to work, so you should make sure that your requests are serializable.

The most common issue here is to use lambda functions on request callbacks that can't be persisted.

So, for example, this won't work:

```
def some_callback(self, response):
    somearg = 'test'
    return Request('http://www.example.com', callback=lambda r: self.other_callback(r, somearg))

def other_callback(self, response, somearg):
    print "the argument passed is:", somearg
```

But this will:

```
def some_callback(self, response):
    somearg = 'test'
    return Request('http://www.example.com', meta={'somearg': somearg})

def other_callback(self, response):
    somearg = response.meta['somearg']
    print "the argument passed is:", somearg
```

5.14 DjangoItem

DjangoItem is a class of item that gets its fields definition from a Django model, you simply create a DjangoItem and specify what Django model it relates to.

Besides of getting the model fields defined on your item, DjangoItem provides a method to create and populate a Django model instance with the item data.

5.14.1 Using DjangoItem

DjangoItem works much like ModelForms in Django, you create a subclass and define its `django_model` attribute to be a valid Django model. With this you will get an item with a field for each Django model field.

In addition, you can define fields that aren't present in the model and even override fields that are present in the model defining them in the item.

Let's see some examples:

Django model for the examples:

```
class Person(models.Model):
    name = models.CharField(max_length=255)
    age = models.IntegerField()
```

Defining a basic DjangoItem:

```
class PersonItem(DjangoItem):
    django_model = Person
```

DjangoItem work just like Item:

```
p = PersonItem()
p['name'] = 'John'
p['age'] = '22'
```

To obtain the Django model from the item, we call the extra method `save()` of the DjangoItem:

```
>>> person = p.save()
>>> person.name
'John'
>>> person.age
'22'
>>> person.id
1
```

As you see the model is already saved when we call `save()`, we can prevent this by calling it with `commit=False`. We can use `commit=False` in `save()` method to obtain an unsaved model:

```
>>> person = p.save(commit=False)
>>> person.name
'John'
>>> person.age
'22'
>>> person.id
None
```

As said before, we can add other fields to the item:

```
class PersonItem(DjangoItem):
    django_model = Person
    sex = Field()

p = PersonItem()
p['name'] = 'John'
p['age'] = '22'
p['sex'] = 'M'
```

Note: fields added to the item won't be taken into account when doing a `save()`

And we can override the fields of the model with your own:

```
class PersonItem(DjangoItem):
    django_model = Person
    name = Field(default='No Name')
```

This is useful to provide properties to the field, like a default or any other property that your project uses.

5.14.2 DjangoItem caveats

DjangoItem is a rather convenient way to integrate Scrapy projects with Django models, but bear in mind that Django ORM may not scale well if you scrape a lot of items (ie. millions) with Scrapy. This is because a relational backend is often not a good choice for a write intensive application (such as a web crawler), specially if the database is highly normalized and with many indices.

Frequently Asked Questions Get answers to most frequently asked questions.

Debugging Spiders Learn how to debug common problems of your scrapy spider.

Spiders Contracts Learn how to use contracts for testing your spiders.

Common Practices Get familiar with some Scrapy common practices.

Broad Crawls Tune Scrapy for crawling a lot domains in parallel.

Using Firefox for scraping Learn how to scrape with Firefox and some useful add-ons.

Using Firebug for scraping Learn how to scrape efficiently using Firebug.

Debugging memory leaks Learn how to find and get rid of memory leaks in your crawler.

Downloading Item Images Download static images associated with your scraped items.

Ubuntu packages Install latest Scrapy packages easily on Ubuntu

Scrapy Service (scrapyd) Deploying your Scrapy project in production.

AutoThrottle extension Adjust crawl rate dynamically based on load.

Jobs: pausing and resuming crawls Learn how to pause and resume crawls for large spiders.

DjangoItem Write scraped items using Django models.

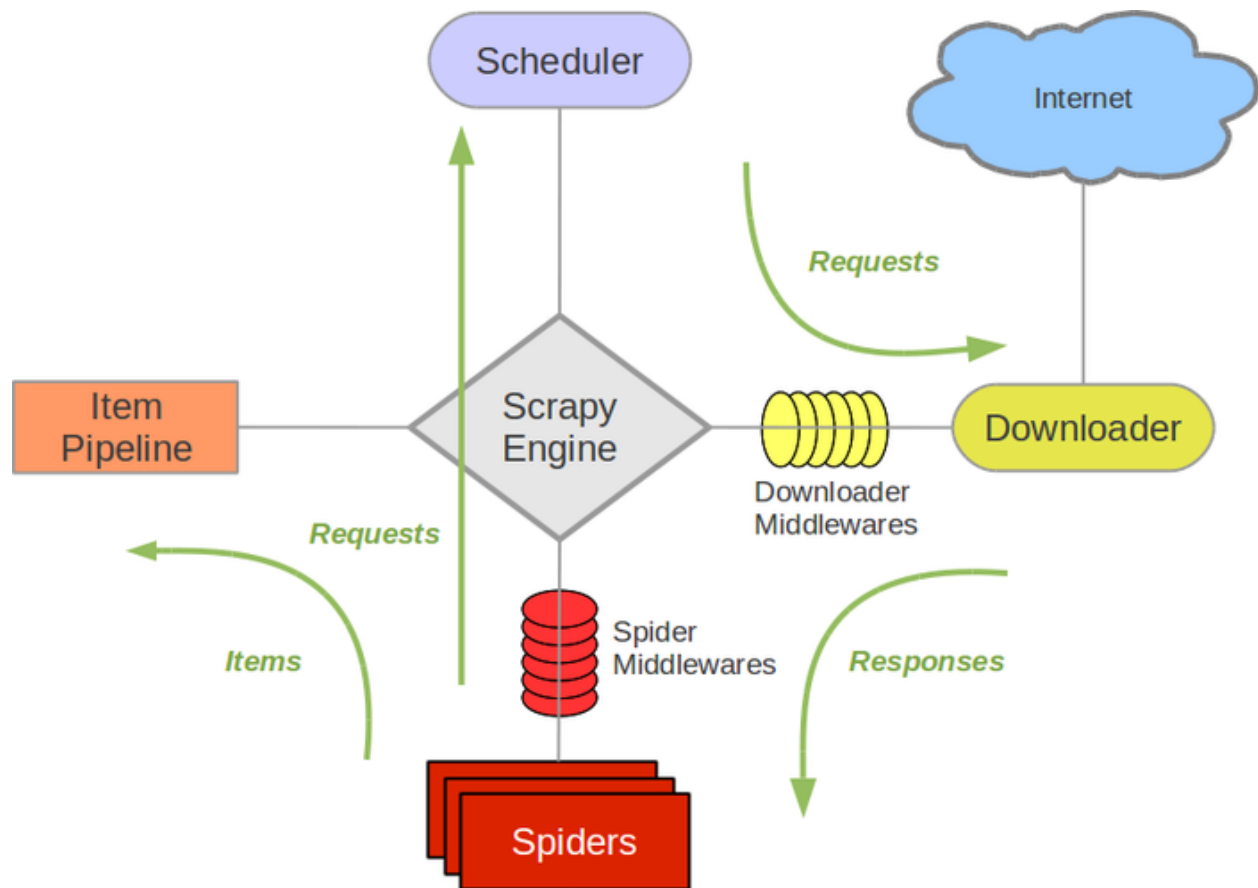
Extending Scrapy

6.1 Architecture overview

This document describes the architecture of Scrapy and how its components interact.

6.1.1 Overview

The following diagram shows an overview of the Scrapy architecture with its components and an outline of the data flow that takes place inside the system (shown by the green arrows). A brief description of the components is included below with links for more detailed information about them. The data flow is also described below.



6.1.2 Components

Scrapy Engine

The engine is responsible for controlling the data flow between all components of the system, and triggering events when certain actions occur. See the Data Flow section below for more details.

Scheduler

The Scheduler receives requests from the engine and enqueues them for feeding them later (also to the engine) when the engine requests them.

Downloader

The Downloader is responsible for fetching web pages and feeding them to the engine which, in turn, feeds them to the spiders.

Spiders

Spiders are custom classes written by Scrapy users to parse responses and extract items (aka scraped items) from them or additional URLs (requests) to follow. Each spider is able to handle a specific domain (or group of domains). For more information see [Spiders](#).

Item Pipeline

The Item Pipeline is responsible for processing the items once they have been extracted (or scraped) by the spiders. Typical tasks include cleansing, validation and persistence (like storing the item in a database). For more information see [Item Pipeline](#).

Downloader middlewares

Downloader middlewares are specific hooks that sit between the Engine and the Downloader and process requests when they pass from the Engine to the Downloader, and responses that pass from Downloader to the Engine. They provide a convenient mechanism for extending Scrapy functionality by plugging custom code. For more information see [Downloader Middleware](#).

Spider middlewares

Spider middlewares are specific hooks that sit between the Engine and the Spiders and are able to process spider input (responses) and output (items and requests). They provide a convenient mechanism for extending Scrapy functionality by plugging custom code. For more information see [Spider Middleware](#).

6.1.3 Data flow

The data flow in Scrapy is controlled by the execution engine, and goes like this:

1. The Engine opens a domain, locates the Spider that handles that domain, and asks the spider for the first URLs to crawl.
2. The Engine gets the first URLs to crawl from the Spider and schedules them in the Scheduler, as Requests.
3. The Engine asks the Scheduler for the next URLs to crawl.
4. The Scheduler returns the next URLs to crawl to the Engine and the Engine sends them to the Downloader, passing through the Downloader Middleware (request direction).
5. Once the page finishes downloading the Downloader generates a Response (with that page) and sends it to the Engine, passing through the Downloader Middleware (response direction).
6. The Engine receives the Response from the Downloader and sends it to the Spider for processing, passing through the Spider Middleware (input direction).
7. The Spider processes the Response and returns scraped Items and new Requests (to follow) to the Engine.
8. The Engine sends scraped Items (returned by the Spider) to the Item Pipeline and Requests (returned by spider) to the Scheduler
9. The process repeats (from step 2) until there are no more requests from the Scheduler, and the Engine closes the domain.

6.1.4 Event-driven networking

Scrapy is written with [Twisted](#), a popular event-driven networking framework for Python. Thus, it's implemented using a non-blocking (aka asynchronous) code for concurrency.

For more information about asynchronous programming and Twisted see these links:

- [Asynchronous Programming with Twisted](#)
- [Twisted - hello, asynchronous programming](#)

6.2 Downloader Middleware

The downloader middleware is a framework of hooks into Scrapy's request/response processing. It's a light, low-level system for globally altering Scrapy's requests and responses.

6.2.1 Activating a downloader middleware

To activate a downloader middleware component, add it to the `DOWNLOADER_MIDDLEWARES` setting, which is a dict whose keys are the middleware class paths and their values are the middleware orders.

Here's an example:

```
DOWNLOADER_MIDDLEWARES = {
    'myproject.middlewares.CustomDownloaderMiddleware': 543,
}
```

The `DOWNLOADER_MIDDLEWARES` setting is merged with the `DOWNLOADER_MIDDLEWARES_BASE` setting defined in Scrapy (and not meant to be overridden) and then sorted by order to get the final sorted list of enabled middlewares: the first middleware is the one closer to the engine and the last is the one closer to the downloader.

To decide which order to assign to your middleware see the `DOWNLOADER_MIDDLEWARES_BASE` setting and pick a value according to where you want to insert the middleware. The order does matter because each middleware performs a different action and your middleware could depend on some previous (or subsequent) middleware being applied.

If you want to disable a built-in middleware (the ones defined in `DOWNLOADER_MIDDLEWARES_BASE` and enabled by default) you must define it in your project's `DOWNLOADER_MIDDLEWARES` setting and assign `None` as its value. For example, if you want to disable the off-site middleware:

```
DOWNLOADER_MIDDLEWARES = {
    'myproject.middlewares.CustomDownloaderMiddleware': 543,
    'scrapy.contrib.downloadermiddleware.useragent.UserAgentMiddleware': None,
}
```

Finally, keep in mind that some middlewares may need to be enabled through a particular setting. See each middleware documentation for more info.

6.2.2 Writing your own downloader middleware

Writing your own downloader middleware is easy. Each middleware component is a single Python class that defines one or more of the following methods:

```
class scrapy.contrib.downloadermiddleware.DownloaderMiddleware
```

```
    process_request(request, spider)
```

This method is called for each request that goes through the download middleware.

`process_request()` should return either `None`, a `Response` object, or a `Request` object.

If it returns `None`, Scrapy will continue processing this request, executing all other middlewares until, finally, the appropriate downloader handler is called the request performed (and its response downloaded).

If it returns a `Response` object, Scrapy won't bother calling ANY other request or exception middleware, or the appropriate download function; it'll return that `Response`. Response middleware is always called on every `Response`.

If it returns a `Request` object, the returned request will be rescheduled (in the Scheduler) to be downloaded in the future. The callback of the original request will always be called. If the new request has a

callback it will be called with the response downloaded, and the output of that callback will then be passed to the original callback. If the new request doesn't have a callback, the response downloaded will be just passed to the original request callback.

If it returns an `IgnoreRequest` exception, the entire request will be dropped completely and its callback never called.

Parameters

- **request** (`Request` object) – the request being processed
- **spider** (`BaseSpider` object) – the spider for which this request is intended

process_response (*request, response, spider*)

`process_response()` should return a `Response` object or raise a `IgnoreRequest` exception.

If it returns a `Response` (it could be the same given response, or a brand-new one), that response will continue to be processed with the `process_response()` of the next middleware in the pipeline.

If it returns an `IgnoreRequest` exception, the response will be dropped completely and its callback never called.

Parameters

- **request** (is a `Request` object) – the request that originated the response
- **response** (`Response` object) – the response being processed
- **spider** (`BaseSpider` object) – the spider for which this response is intended

process_exception (*request, exception, spider*)

Scrapy calls `process_exception()` when a download handler or a `process_request()` (from a downloader middleware) raises an exception.

`process_exception()` should return either `None`, `Response` or `Request` object.

If it returns `None`, Scrapy will continue processing this exception, executing any other exception middleware, until no middleware is left and the default exception handling kicks in.

If it returns a `Response` object, the response middleware kicks in, and won't bother calling any other exception middleware.

If it returns a `Request` object, the returned request is used to instruct an immediate redirection. The original request won't finish until the redirected request is completed. This stops the `process_exception()` middleware the same as returning `Response` would do.

Parameters

- **request** (is a `Request` object) – the request that generated the exception
- **exception** (an `Exception` object) – the raised exception
- **spider** (`BaseSpider` object) – the spider for which this request is intended

6.2.3 Built-in downloader middleware reference

This page describes all downloader middleware components that come with Scrapy. For information on how to use them and how to write your own downloader middleware, see the [downloader middleware usage guide](#).

For a list of the components enabled by default (and their orders) see the `DOWNLOADER_MIDDLEWARES_BASE` setting.

CookiesMiddleware

class scrapy.contrib.downloadermiddleware.cookies.CookiesMiddleware

This middleware enables working with sites that require cookies, such as those that use sessions. It keeps track of cookies sent by web servers, and send them back on subsequent requests (from that spider), just like web browsers do.

The following settings can be used to configure the cookie middleware:

- `COOKIES_ENABLED`
- `COOKIES_DEBUG`

Multiple cookie sessions per spider

New in version 0.15. There is support for keeping multiple cookie sessions per spider by using the `cookiejar` Request meta key. By default it uses a single cookie jar (session), but you can pass an identifier to use different ones.

For example:

```
for i, url in enumerate(urls):
    yield Request("http://www.example.com", meta={'cookiejar': i},
                  callback=self.parse_page)
```

Keep in mind that the `cookiejar` meta key is not “sticky”. You need to keep passing it along on subsequent requests. For example:

```
def parse_page(self, response):
    # do some processing
    return Request("http://www.example.com/otherpage",
                   meta={'cookiejar': response.meta['cookiejar']},
                   callback=self.parse_other_page)
```

COOKIES_ENABLED

Default: True

Whether to enable the cookies middleware. If disabled, no cookies will be sent to web servers.

COOKIES_DEBUG

Default: False

If enabled, Scrapy will log all cookies sent in requests (ie. Cookie header) and all cookies received in responses (ie. Set-Cookie header).

Here’s an example of a log with `COOKIES_DEBUG` enabled:

```
2011-04-06 14:35:10-0300 [diningcity] INFO: Spider opened
2011-04-06 14:35:10-0300 [diningcity] DEBUG: Sending cookies to: <GET http://www.diningcity.com/netherlands/>
Cookie: clientlanguage_nl=en_EN
2011-04-06 14:35:14-0300 [diningcity] DEBUG: Received cookies from: <200 http://www.diningcity.com/netherlands/>
Set-Cookie: JSESSIONID=B~FA4DC0C496C8762AE4F1A620EAB34F38; Path=/
Set-Cookie: ip_isocode=US
Set-Cookie: clientlanguage_nl=en_EN; Expires=Thu, 07-Apr-2011 21:21:34 GMT; Path=/
2011-04-06 14:49:50-0300 [diningcity] DEBUG: Crawled (200) <GET http://www.diningcity.com/netherlands/>
[...]
```

DefaultHeadersMiddleware

class scrapy.contrib.downloadermiddleware.defaultheaders.**DefaultHeadersMiddleware**
This middleware sets all default requests headers specified in the `DEFAULT_REQUEST_HEADERS` setting.

DownloadTimeoutMiddleware

class scrapy.contrib.downloadermiddleware.downloadtimeout.**DownloadTimeoutMiddleware**
This middleware sets the download timeout for requests specified in the `DOWNLOAD_TIMEOUT` setting.

HttpAuthMiddleware

class scrapy.contrib.downloadermiddleware.httpauth.**HttpAuthMiddleware**
This middleware authenticates all requests generated from certain spiders using [Basic access authentication](#) (aka. HTTP auth).

To enable HTTP authentication from certain spiders, set the `http_user` and `http_pass` attributes of those spiders.

Example:

```
class SomeIntranetSiteSpider(CrawlSpider):  
  
    http_user = 'someuser'  
    http_pass = 'somepass'  
    name = 'intranet.example.com'  
  
    # .. rest of the spider code omitted ...
```

HttpCacheMiddleware

class scrapy.contrib.downloadermiddleware.httpcache.**HttpCacheMiddleware**
This middleware provides low-level cache to all HTTP requests and responses. Every request and its corresponding response are cached. When the same request is seen again, the response is returned without transferring anything from the Internet.

The HTTP cache is useful for testing spiders faster (without having to wait for downloads every time) and for trying your spider offline, when an Internet connection is not available.

Scrapy ships with two storage backends for the HTTP cache middleware:

- *DBM storage backend (default)*
- *File system backend*

You can change the storage backend with the `HTTPCACHE_STORAGE` setting. Or you can also implement your own backend.

DBM storage backend (default)

New in version 0.13. A [DBM](#) storage backend is available for the HTTP cache middleware. To use it (note: it is the default storage backend) set `HTTPCACHE_STORAGE` to `scrapy.contrib.httpcache.DbmCacheStorage`.

By default, it uses the [anydbm](#) module, but you can change it with the `HTTPCACHE_DBM_MODULE` setting.

File system backend

A file system storage backend is also available for the HTTP cache middleware. To use it (instead of the default `DBM` storage backend) set `HTTPCACHE_STORAGE` to `scrapy.contrib.downloadermiddleware.httpcache.FilesystemCacheStorage`.

Each request/response pair is stored in a different directory containing the following files:

- `request_body` - the plain request body
- `request_headers` - the request headers (in raw HTTP format)
- `response_body` - the plain response body
- `response_headers` - the request headers (in raw HTTP format)
- `meta` - some metadata of this cache resource in Python `repr()` format (grep-friendly format)
- `pickled_meta` - the same metadata in `meta` but pickled for more efficient deserialization

The directory name is made from the request fingerprint (see `scrapy.utils.request.fingerprint`), and one level of subdirectories is used to avoid creating too many files into the same directory (which is inefficient in many file systems). An example directory could be:

```
/path/to/cache/dir/example.com/72/72811f648e718090f041317756c03adb0ada46c7
```

HTTPCache middleware settings

The `HttpCacheMiddleware` can be configured through the following settings:

HTTPCACHE_ENABLED New in version 0.11. Default: `False`

Whether the HTTP cache will be enabled. Changed in version 0.11: Before 0.11, `HTTPCACHE_DIR` was used to enable cache.

HTTPCACHE_EXPIRATION_SECS Default: `0`

Expiration time for cached requests, in seconds.

Cached requests older than this time will be re-downloaded. If zero, cached requests will never expire. Changed in version 0.11: Before 0.11, zero meant cached requests always expire.

HTTPCACHE_DIR Default: `'httpcache'`

The directory to use for storing the (low-level) HTTP cache. If empty, the HTTP cache will be disabled. If a relative path is given, is taken relative to the project data dir. For more info see: *Default structure of Scrapy projects*.

HTTPCACHE_IGNORE_HTTP_CODES New in version 0.10. Default: `[]`

Don't cache response with these HTTP codes.

HTTPCACHE_IGNORE_MISSING Default: `False`

If enabled, requests not found in the cache will be ignored instead of downloaded.

HTTPCACHE_IGNORE_SCHEMES New in version 0.10. Default: `['file']`

Don't cache responses with these URI schemes.

HTTPCACHE_STORAGE Default: `'scrapy.contrib.httppcache.DbmCacheStorage'`

The class which implements the cache storage backend.

HTTPCACHE_DBM_MODULE New in version 0.13. Default: `'anydbm'`

The database module to use in the *DBM storage backend*. This setting is specific to the DBM backend.

HttpCompressionMiddleware

class scrapy.contrib.downloadermiddleware.httpcompression.**HttpCompressionMiddleware**
This middleware allows compressed (gzip, deflate) traffic to be sent/received from web sites.

ChunkedTransferMiddleware

class scrapy.contrib.downloadermiddleware.chunked.**ChunkedTransferMiddleware**
This middleware adds support for *chunked transfer encoding*

HttpProxyMiddleware

New in version 0.8.

class scrapy.contrib.downloadermiddleware.httpproxy.**HttpProxyMiddleware**
This middleware sets the HTTP proxy to use for requests, by setting the `proxy` meta value to *Request* objects.
Like the Python standard library modules *urllib* and *urllib2*, it obeys the following environment variables:

- `http_proxy`
- `https_proxy`
- `no_proxy`

RedirectMiddleware

class scrapy.contrib.downloadermiddleware.redirect.**RedirectMiddleware**
This middleware handles redirection of requests based on response status and meta-refresh html tag.

The urls which the request goes through (while being redirected) can be found in the `redirect_urls` *Request.meta* key.

The *RedirectMiddleware* can be configured through the following settings (see the settings documentation for more info):

- `REDIRECT_ENABLED`
- `REDIRECT_MAX_TIMES`
- `REDIRECT_MAX_METAREFRESH_DELAY`

If *Request.meta* contains the `dont_redirect` key, the request will be ignored by this middleware.

RedirectMiddleware settings

REDIRECT_ENABLED New in version 0.13. Default: `True`

Whether the Redirect middleware will be enabled.

REDIRECT_MAX_TIMES Default: `20`

The maximum number of redirections that will be follow for a single request.

REDIRECT_MAX_METAREFRESH_DELAY Default: `100`

The maximum meta-refresh delay (in seconds) to follow the redirection.

RetryMiddleware

class scrapy.contrib.downloadermiddleware.retry.**RetryMiddleware**

A middleware to retry failed requests that are potentially caused by temporary problems such as a connection timeout or HTTP 500 error.

Failed pages are collected on the scraping process and rescheduled at the end, once the spider has finished crawling all regular (non failed) pages. Once there are no more failed pages to retry, this middleware sends a signal (`retry_complete`), so other extensions could connect to that signal.

The `RetryMiddleware` can be configured through the following settings (see the settings documentation for more info):

- `RETRY_ENABLED`
- `RETRY_TIMES`
- `RETRY_HTTP_CODES`

About HTTP errors to consider:

You may want to remove 400 from `RETRY_HTTP_CODES`, if you stick to the HTTP protocol. It's included by default because it's a common code used to indicate server overload, which would be something we want to retry. If `Request.meta` contains the `dont_retry` key, the request will be ignored by this middleware.

RetryMiddleware Settings

RETRY_ENABLED New in version 0.13. Default: `True`

Whether the Retry middleware will be enabled.

RETRY_TIMES Default: `2`

Maximum number of times to retry, in addition to the first download.

RETRY_HTTP_CODES Default: `[500, 503, 504, 400, 408]`

Which HTTP response codes to retry. Other errors (DNS lookup issues, connections lost, etc) are always retried.

RobotsTxtMiddleware

class scrapy.contrib.downloadermiddleware.robotstxt.**RobotsTxtMiddleware**

This middleware filters out requests forbidden by the robots.txt exclusion standard.

To make sure Scrapy respects robots.txt make sure the middleware is enabled and the `ROBOTSTXT_OBEY` setting is enabled.

Warning: Keep in mind that, if you crawl using multiple concurrent requests per domain, Scrapy could still download some forbidden pages if they were requested before the robots.txt file was downloaded. This is a known limitation of the current robots.txt middleware and will be fixed in the future.

DownloaderStats

class scrapy.contrib.downloadermiddleware.stats.**DownloaderStats**

Middleware that stores stats of all requests, responses and exceptions that pass through it.

To use this middleware you must enable the `DOWNLOADER_STATS` setting.

UserAgentMiddleware

class scrapy.contrib.downloadermiddleware.useragent.**UserAgentMiddleware**

Middleware that allows spiders to override the default user agent.

In order for a spider to override the default user agent, its `user_agent` attribute must be set.

6.3 Spider Middleware

The spider middleware is a framework of hooks into Scrapy's spider processing mechanism where you can plug custom functionality to process the requests that are sent to *Spiders* for processing and to process the responses and items that are generated from spiders.

6.3.1 Activating a spider middleware

To activate a spider middleware component, add it to the `SPIDER_MIDDLEWARES` setting, which is a dict whose keys are the middleware class path and their values are the middleware orders.

Here's an example:

```
SPIDER_MIDDLEWARES = {
    'myproject.middlewares.CustomSpiderMiddleware': 543,
}
```

The `SPIDER_MIDDLEWARES` setting is merged with the `SPIDER_MIDDLEWARES_BASE` setting defined in Scrapy (and not meant to be overridden) and then sorted by order to get the final sorted list of enabled middlewares: the first middleware is the one closer to the engine and the last is the one closer to the spider.

To decide which order to assign to your middleware see the `SPIDER_MIDDLEWARES_BASE` setting and pick a value according to where you want to insert the middleware. The order does matter because each middleware performs a different action and your middleware could depend on some previous (or subsequent) middleware being applied.

If you want to disable a builtin middleware (the ones defined in `SPIDER_MIDDLEWARES_BASE`, and enabled by default) you must define it in your project `SPIDER_MIDDLEWARES` setting and assign `None` as its value. For example, if you want to disable the off-site middleware:

```
SPIDER_MIDDLEWARES = {
    'myproject.middlewares.CustomSpiderMiddleware': 543,
    'scrapy.contrib.spidermiddleware.offsite.OffsiteMiddleware': None,
}
```

Finally, keep in mind that some middlewares may need to be enabled through a particular setting. See each middleware documentation for more info.

6.3.2 Writing your own spider middleware

Writing your own spider middleware is easy. Each middleware component is a single Python class that defines one or more of the following methods:

class `scrapy.contrib.spidermiddleware.SpiderMiddleware`

`process_spider_input` (*response*, *spider*)

This method is called for each response that goes through the spider middleware and into the spider, for processing.

`process_spider_input()` should return `None` or raise an exception.

If it returns `None`, Scrapy will continue processing this response, executing all other middlewares until, finally, the response is handled to the spider for processing.

If it raises an exception, Scrapy won't bother calling any other spider middleware `process_spider_input()` and will call the request errback. The output of the errback is chained back in the other direction for `process_spider_output()` to process it, or `process_spider_exception()` if it raised an exception.

Parameters

- **response** (`Response` object) – the response being processed
- **spider** (`BaseSpider` object) – the spider for which this response is intended

`process_spider_output` (*response*, *result*, *spider*)

This method is called with the results returned from the Spider, after it has processed the response.

`process_spider_output()` must return an iterable of `Request` or `Item` objects.

Parameters

- **response** (class:~`scrapy.http.Response` object) – the response which generated this output from the spider
- **result** (an iterable of `Request` or `Item` objects) – the result returned by the spider
- **spider** (`BaseSpider` object) – the spider whose result is being processed

`process_spider_exception` (*response*, *exception*, *spider*)

This method is called when a spider or `process_spider_input()` method (from other spider middleware) raises an exception.

`process_spider_exception()` should return either `None` or an iterable of `Response` or `Item` objects.

If it returns `None`, Scrapy will continue processing this exception, executing any other `process_spider_exception()` in the following middleware components, until no middleware components are left and the exception reaches the engine (where it's logged and discarded).

If it returns an iterable the `process_spider_output()` pipeline kicks in, and no other `process_spider_exception()` will be called.

Parameters

- **response** (`Response` object) – the response being processed when the exception was raised
- **exception** (`Exception` object) – the exception raised
- **spider** (`scrapy.spider.BaseSpider` object) – the spider which raised the exception

`process_start_requests` (*start_requests, spider*)

New in version 0.15. This method is called with the start requests of the spider, and works similarly to the `process_spider_output()` method, except that it doesn't have a response associated and must return only requests (not items).

It receives an iterable (in the `start_requests` parameter) and must return another iterable of `Request` objects.

Note: When implementing this method in your spider middleware, you should always return an iterable (that follows the input one) and not consume all `start_requests` iterator because it can be very large (or even unbounded) and cause a memory overflow. The Scrapy engine is designed to pull start requests while it has capacity to process them, so the start requests iterator can be effectively endless where there is some other condition for stopping the spider (like a time limit or item/page count).

Parameters

- **start_requests** (an iterable of `Request`) – the start requests
- **spider** (`BaseSpider` object) – the spider to whom the start requests belong

6.3.3 Built-in spider middleware reference

This page describes all spider middleware components that come with Scrapy. For information on how to use them and how to write your own spider middleware, see the [spider middleware usage guide](#).

For a list of the components enabled by default (and their orders) see the `SPIDER_MIDDLEWARES_BASE` setting.

DepthMiddleware

class `scrapy.contrib.spidermiddleware.depth.DepthMiddleware`

DepthMiddleware is a scrape middleware used for tracking the depth of each `Request` inside the site being scraped. It can be used to limit the maximum depth to scrape or things like that.

The `DepthMiddleware` can be configured through the following settings (see the settings documentation for more info):

- `DEPTH_LIMIT` - The maximum depth that will be allowed to crawl for any site. If zero, no limit will be imposed.
- `DEPTH_STATS` - Whether to collect depth stats.

- `DEPTH_PRIORITY` - Whether to prioritize the requests based on their depth.

HttpErrorMiddleware

class scrapy.contrib.spidermiddleware.httperror.HttpErrorMiddleware

Filter out unsuccessful (erroneous) HTTP responses so that spiders don't have to deal with them, which (most of the time) imposes an overhead, consumes more resources, and makes the spider logic more complex.

According to the [HTTP standard](#), successful responses are those whose status codes are in the 200-300 range.

If you still want to process response codes outside that range, you can specify which response codes the spider is able to handle using the `handle_httpstatus_list` spider attribute.

For example, if you want your spider to handle 404 responses you can do this:

```
class MySpider(CrawlSpider):
    handle_httpstatus_list = [404]
```

The `handle_httpstatus_list` key of `Request.meta` can also be used to specify which response codes to allow on a per-request basis.

Keep in mind, however, that it's usually a bad idea to handle non-200 responses, unless you really know what you're doing.

For more information see: [HTTP Status Code Definitions](#).

OffsiteMiddleware

class scrapy.contrib.spidermiddleware.offsite.OffsiteMiddleware

Filters out Requests for URLs outside the domains covered by the spider.

This middleware filters out every request whose host names aren't in the spider's `allowed_domains` attribute.

When your spider returns a request for a domain not belonging to those covered by the spider, this middleware will log a debug message similar to this one:

```
DEBUG: Filtered offsite request to 'www.othersite.com': <GET http://www.othersite.com/some/page.
```

To avoid filling the log with too much noise, it will only print one of these messages for each new domain filtered. So, for example, if another request for `www.othersite.com` is filtered, no log message will be printed. But if a request for `someothersite.com` is filtered, a message will be printed (but only for the first request filtered).

If the spider doesn't define an `allowed_domains` attribute, or the attribute is empty, the offsite middleware will allow all requests.

If the request has the `dont_filter` attribute set, the offsite middleware will allow the request even if its domain is not listed in allowed domains.

RefererMiddleware

class scrapy.contrib.spidermiddleware.referer.RefererMiddleware

Populates Request referer field, based on the Response which originated it.

RefererMiddleware settings

REFERER_ENABLED New in version 0.15. Default: `True`

Whether to enable referer middleware.

UrlLengthMiddleware

class scrapy.contrib.spidermiddleware.urllength.UrlLengthMiddleware

Filters out requests with URLs longer than `URLLENGTH_LIMIT`

The `UrlLengthMiddleware` can be configured through the following settings (see the settings documentation for more info):

- `URLLENGTH_LIMIT` - The maximum URL length to allow for crawled URLs.

6.4 Extensions

The extensions framework provides a mechanism for inserting your own custom functionality into Scrapy.

Extensions are just regular classes that are instantiated at Scrapy startup, when extensions are initialized.

6.4.1 Extension settings

Extensions use the *Scrapy settings* to manage their settings, just like any other Scrapy code.

It is customary for extensions to prefix their settings with their own name, to avoid collision with existing (and future) extensions. For example, an hypothetical extension to handle [Google Sitemaps](#) would use settings like `GOOGLE-SITEMAP_ENABLED`, `GOOGLESITEMAP_DEPTH`, and so on.

6.4.2 Loading & activating extensions

Extensions are loaded and activated at startup by instantiating a single instance of the extension class. Therefore, all the extension initialization code must be performed in the class constructor (`__init__` method).

To make an extension available, add it to the `EXTENSIONS` setting in your Scrapy settings. In `EXTENSIONS`, each extension is represented by a string: the full Python path to the extension's class name. For example:

```
EXTENSIONS = {
    'scrapy.contrib.corestats.CoreStats': 500,
    'scrapy.webservice.WebService': 500,
    'scrapy.telnet.TelnetConsole': 500,
}
```

As you can see, the `EXTENSIONS` setting is a dict where the keys are the extension paths, and their values are the orders, which define the extension *loading* order. Extensions orders are not as important as middleware orders though, and they are typically irrelevant, ie. it doesn't matter in which order the extensions are loaded because they don't depend on each other [1].

However, this feature can be exploited if you need to add an extension which depends on other extensions already loaded.

[1] This is why the `EXTENSIONS_BASE` setting in Scrapy (which contains all built-in extensions enabled by default) defines all the extensions with the same order (500).

6.4.3 Available, enabled and disabled extensions

Not all available extensions will be enabled. Some of them usually depend on a particular setting. For example, the HTTP Cache extension is available by default but disabled unless the `HTTPCACHE_ENABLED` setting is set.

6.4.4 Disabling an extension

In order to disable an extension that comes enabled by default (ie. those included in the `EXTENSIONS_BASE` setting) you must set its order to `None`. For example:

```
EXTENSIONS = {
    'scrapy.contrib.corestats.CoreStats': None,
}
```

6.4.5 Writing your own extension

Writing your own extension is easy. Each extension is a single Python class which doesn't need to implement any particular method.

The main entry point for a Scrapy extension (this also includes middlewares and pipelines) is the `from_crawler` class method which receives a `Crawler` instance which is the main object controlling the Scrapy crawler. Through that object you can access settings, signals, stats, and also control the crawler behaviour, if your extension needs to such thing.

Typically, extensions connect to *signals* and perform tasks triggered by them.

Finally, if the `from_crawler` method raises the `NotConfigured` exception, the extension will be disabled. Otherwise, the extension will be enabled.

Sample extension

Here we will implement a simple extension to illustrate the concepts described in the previous section. This extension will log a message every time:

- a spider is opened
- a spider is closed
- a specific number of items are scraped

The extension will be enabled through the `MYEXT_ENABLED` setting and the number of items will be specified through the `MYEXT_ITEMCOUNT` setting.

Here is the code of such extension:

```
from scrapy import signals
from scrapy.exceptions import NotConfigured

class SpiderOpenCloseLogging(object):

    def __init__(self, item_count):
        self.item_count = item_count
        self.items_scraped = 0

    @classmethod
    def from_crawler(cls, crawler):
        # first check if the extension should be enabled and raise
```

```
# NotConfigured otherwise
if not crawler.settings.getbool('MYEXT_ENABLED'):
    raise NotConfigured

# get the number of items from settings
item_count = crawler.settings.getint('MYEXT_ITEMCOUNT', 1000)

# instantiate the extension object
ext = cls(item_count)

# connect the extension object to signals
crawler.signals.connect(ext.spider_opened, signal=signals.spider_opened)
crawler.signals.connect(ext.spider_closed, signal=signals.spider_closed)
crawler.signals.connect(ext.item_scraped, signal=signals.item_scraped)

# return the extension object
return ext

def spider_opened(self, spider):
    spider.log("opened spider %s" % spider.name)

def spider_closed(self, spider):
    spider.log("closed spider %s" % spider.name)

def item_scraped(self, item, spider):
    self.items_scraped += 1
    if self.items_scraped == self.item_count:
        spider.log("scraped %d items, resetting counter" % self.items_scraped)
        self.item_count = 0
```

6.4.6 Built-in extensions reference

General purpose extensions

Log Stats extension

class scrapy.contrib.logstats.**LogStats**

Log basic stats like crawled pages and scraped items.

Core Stats extension

class scrapy.contrib.corestats.**CoreStats**

Enable the collection of core statistics, provided the stats collection is enabled (see [Stats Collection](#)).

Web service extension

class scrapy.webservice.**WebService**

See *topics-webservice*.

Telnet console extension

class scrapy.telnet.**TelnetConsole**

Provides a telnet console for getting into a Python interpreter inside the currently running Scrapy process, which can be very useful for debugging.

The telnet console must be enabled by the `TELNETCONSOLE_ENABLED` setting, and the server will listen in the port specified in `TELNETCONSOLE_PORT`.

Memory usage extension

class scrapy.contrib.memusage.**MemoryUsage**

Note: This extension does not work in Windows.

Monitors the memory used by the Scrapy process that runs the spider and:

1. sends a notification e-mail when it exceeds a certain value 2. closes the spider when it exceeds a certain value

The notification e-mails can be triggered when a certain warning value is reached (`MEMUSAGE_WARNING_MB`) and when the maximum value is reached (`MEMUSAGE_LIMIT_MB`) which will also cause the spider to be closed and the Scrapy process to be terminated.

This extension is enabled by the `MEMUSAGE_ENABLED` setting and can be configured with the following settings:

- `MEMUSAGE_LIMIT_MB`
- `MEMUSAGE_WARNING_MB`
- `MEMUSAGE_NOTIFY_MAIL`
- `MEMUSAGE_REPORT`

Memory debugger extension

class scrapy.contrib.memdebug.**MemoryDebugger**

An extension for debugging memory usage. It collects information about:

- objects uncollected by the Python garbage collector
- libxml2 memory leaks
- objects left alive that shouldn't. For more info, see *Debugging memory leaks with trackref*

To enable this extension, turn on the `MEMDEBUG_ENABLED` setting. The info will be stored in the stats.

Close spider extension

class scrapy.contrib.closespider.**CloseSpider**

Closes a spider automatically when some conditions are met, using a specific closing reason for each condition.

The conditions for closing a spider can be configured through the following settings:

- `CLOSESPIDER_TIMEOUT`
- `CLOSESPIDER_ITEMCOUNT`

- `CLOSESPIDER_PAGECOUNT`
- `CLOSESPIDER_ERRORCOUNT`

CLOSESPIDER_TIMEOUT Default: 0

An integer which specifies a number of seconds. If the spider remains open for more than that number of second, it will be automatically closed with the reason `closespider_timeout`. If zero (or non set), spiders won't be closed by timeout.

CLOSESPIDER_ITEMCOUNT Default: 0

An integer which specifies a number of items. If the spider scrapes more than that amount if items and those items are passed by the item pipeline, the spider will be closed with the reason `closespider_itemcount`. If zero (or non set), spiders won't be closed by number of passed items.

CLOSESPIDER_PAGECOUNT New in version 0.11. Default: 0

An integer which specifies the maximum number of responses to crawl. If the spider crawls more than that, the spider will be closed with the reason `closespider_pagecount`. If zero (or non set), spiders won't be closed by number of crawled responses.

CLOSESPIDER_ERRORCOUNT New in version 0.11. Default: 0

An integer which specifies the maximum number of errors to receive before closing the spider. If the spider generates more than that number of errors, it will be closed with the reason `closespider_errorcount`. If zero (or non set), spiders won't be closed by number of errors.

StatsMailer extension

```
class scrapy.contrib.statsmailer.StatsMailer
```

This simple extension can be used to send a notification e-mail every time a domain has finished scraping, including the Scrapy stats collected. The email will be sent to all recipients specified in the `STATSMAILER_RCPTS` setting.

Debugging extensions

Stack trace dump extension

```
class scrapy.contrib.debug.StackTraceDump
```

Dumps information about the running process when a `SIGQUIT` or `SIGUSR2` signal is received. The information dumped is the following:

1. engine status (using `scrapy.utils.engine.get_engine_status()`)
2. live references (see *Debugging memory leaks with trackref*)
3. stack trace of all threads

After the stack trace and engine status is dumped, the Scrapy process continues running normally.

This extension only works on POSIX-compliant platforms (ie. not Windows), because the `SIGQUIT` and `SIGUSR2` signals are not available on Windows.

There are at least two ways to send Scrapy the `SIGQUIT` signal:

1. By pressing Ctrl-C while a Scrapy process is running (Linux only?)
2. By running this command (assuming <pid> is the process id of the Scrapy process):

```
kill -QUIT <pid>
```

Debugger extension

class scrapy.contrib.debug.**Debugger**

Invokes a [Python debugger](#) inside a running Scrapy process when a `SIGUSR2` signal is received. After the debugger is exited, the Scrapy process continues running normally.

For more info see *Debugging in Python*.

This extension only works on POSIX-compliant platforms (ie. not Windows).

6.5 Core API

New in version 0.15. This section documents the Scrapy core API, and it's intended for developers of extensions and middlewares.

6.5.1 Crawler API

The main entry point to Scrapy API is the `Crawler` object, passed to extensions through the `from_crawler` class method. This object provides access to all Scrapy core components, and it's the only way for extensions to access them and hook their functionality into Scrapy. The Extension Manager is responsible for loading and keeping track of installed extensions and it's configured through the `EXTENSIONS` setting which contains a dictionary of all available extensions and their order similar to how you [configure the downloader middlewares](#).

class scrapy.crawler.**Crawler**(*settings*)

The Crawler object must be instantiated with a `scrapy.settings.Settings` object.

settings

The settings manager of this crawler.

This is used by extensions & middlewares to access the Scrapy settings of this crawler.

For an introduction on Scrapy settings see [Settings](#).

For the API see `Settings` class.

signals

The signals manager of this crawler.

This is used by extensions & middlewares to hook themselves into Scrapy functionality.

For an introduction on signals see [Signals](#).

For the API see `SignalManager` class.

stats

The stats collector of this crawler.

This is used from extensions & middlewares to record stats of their behaviour, or access stats collected by other extensions.

For an introduction on stats collection see [Stats Collection](#).

For the API see `StatsCollector` class.

extensions

The extension manager that keeps track of enabled extensions.

Most extensions won't need to access this attribute.

For an introduction on extensions and a list of available extensions on Scrapy see [Extensions](#).

spiders

The spider manager which takes care of loading and instantiating spiders.

Most extensions won't need to access this attribute.

engine

The execution engine, which coordinates the core crawling logic between the scheduler, downloader and spiders.

Some extension may want to access the Scrapy engine, to modify inspect or modify the downloader and scheduler behaviour, although this is an advanced use and this API is not yet stable.

configure()

Configure the crawler.

This loads extensions, middlewares and spiders, leaving the crawler ready to be started. It also configures the execution engine.

start()

Start the crawler. This calls `configure()` if it hasn't been called yet.

6.5.2 Settings API

class scrapy.settings.Settings

This object that provides access to Scrapy settings.

overrides

Global overrides are the ones that take most precedence, and are usually populated by command-line options.

Overrides should be populated *before* configuring the Crawler object (through the `configure()` method), otherwise they won't have any effect. You don't typically need to worry about overrides unless you are implementing your own Scrapy command.

get(name, default=None)

Get a setting value without affecting its original type.

Parameters

- **name** (*string*) – the setting name
- **default** (*any*) – the value to return if no setting is found

getbool(name, default=False)

Get a setting value as a boolean. For example, both 1 and '1', and True return True, while 0, '0', False and None return False ``

For example, settings populated through environment variables set to '0' will return False when using this method.

Parameters

- **name** (*string*) – the setting name
- **default** (*any*) – the value to return if no setting is found

getint (*name*, *default=0*)

Get a setting value as an int

Parameters

- **name** (*string*) – the setting name
- **default** (*any*) – the value to return if no setting is found

getfloat (*name*, *default=0.0*)

Get a setting value as a float

Parameters

- **name** (*string*) – the setting name
- **default** (*any*) – the value to return if no setting is found

getlist (*name*, *default=None*)

Get a setting value as a list. If the setting original type is a list it will be returned verbatim. If it's a string it will be split by ",".

For example, settings populated through environment variables set to 'one,two' will return a list ['one', 'two'] when using this method.

Parameters

- **name** (*string*) – the setting name
- **default** (*any*) – the value to return if no setting is found

6.5.3 Signals API

class scrapy.signalmanager.SignalManager

connect (*receiver*, *signal*)

Connect a receiver function to a signal.

The signal can be any object, although Scrapy comes with some predefined signals that are documented in the [Signals](#) section.

Parameters

- **receiver** (*callable*) – the function to be connected
- **signal** (*object*) – the signal to connect to

send_catch_log (*signal*, ***kwargs*)

Send a signal, catch exceptions and log them.

The keyword arguments are passed to the signal handlers (connected through the [connect\(\)](#) method).

send_catch_log_deferred (*signal*, ***kwargs*)

Like [send_catch_log\(\)](#) but supports returning [deferreds](#) from signal handlers.

Returns a [deferred](#) that gets fired once all signal handlers deferreds were fired. Send a signal, catch exceptions and log them.

The keyword arguments are passed to the signal handlers (connected through the [connect\(\)](#) method).

disconnect (*receiver*, *signal*)

Disconnect a receiver function from a signal. This has the opposite effect of the [connect\(\)](#) method, and the arguments are the same.

disconnect_all (*signal*)

Disconnect all receivers from the given signal.

Parameters **signal** (*object*) – the signal to disconnect from

6.5.4 Stats Collector API

There are several Stats Collectors available under the `scrapy.statscol` module and they all implement the Stats Collector API defined by the `StatsCollector` class (which they all inherit from).

class scrapy.statscol.**StatsCollector**

get_value (*key*, *default=None*)

Return the value for the given stats key or default if it doesn't exist.

get_stats ()

Get all stats from the currently running spider as a dict.

set_value (*key*, *value*)

Set the given value for the given stats key.

set_stats (*stats*)

Override the current stats with the dict passed in *stats* argument.

inc_value (*key*, *count=1*, *start=0*)

Increment the value of the given stats key, by the given count, assuming the start value given (when it's not set).

max_value (*key*, *value*)

Set the given value for the given key only if current value for the same key is lower than value. If there is no current value for the given key, the value is always set.

min_value (*key*, *value*)

Set the given value for the given key only if current value for the same key is greater than value. If there is no current value for the given key, the value is always set.

clear_stats ()

Clear all stats.

The following methods are not part of the stats collection api but instead used when implementing custom stats collectors:

open_spider (*spider*)

Open the given spider for stats collection.

close_spider (*spider*)

Close the given spider. After this is called, no more specific stats can be accessed or collected.

Architecture overview Understand the Scrapy architecture.

Downloader Middleware Customize how pages get requested and downloaded.

Spider Middleware Customize the input and output of your spiders.

Extensions Extend Scrapy with your custom functionality

Core API Use it on extensions and middlewares to extend Scrapy functionality

Reference

7.1 Requests and Responses

Scrapy uses `Request` and `Response` objects for crawling web sites.

Typically, `Request` objects are generated in the spiders and pass across the system until they reach the Downloader, which executes the request and returns a `Response` object which travels back to the spider that issued the request.

Both `Request` and `Response` classes have subclasses which add functionality not required in the base classes. These are described below in *Request subclasses* and *Response subclasses*.

7.1.1 Request objects

```
class scrapy.http.Request (url[, method='GET', body, headers, cookies, meta, encoding='utf-8', priority=0, dont_filter=False, callback, errback ])
```

A `Request` object represents an HTTP request, which is usually generated in the Spider and executed by the Downloader, and thus generating a `Response`.

Parameters

- **url** (*string*) – the URL of this request
- **method** (*string*) – the HTTP method of this request. Defaults to 'GET'.
- **meta** (*dict*) – the initial values for the `Request.meta` attribute. If given, the dict passed in this parameter will be shallow copied.
- **body** (*str or unicode*) – the request body. If a `unicode` is passed, then it's encoded to `str` using the `encoding` passed (which defaults to `utf-8`). If `body` is not given,, an empty string is stored. Regardless of the type of this argument, the final value stored will be a `str` (never `unicode` or `None`).
- **headers** (*dict*) – the headers of this request. The dict values can be strings (for single valued headers) or lists (for multi-valued headers).
- **cookies** (*dict or list*) – the request cookies. These can be sent in two forms.
 1. Using a dict:

```
request_with_cookies = Request(url="http://www.example.com",
                               cookies={'currency': 'USD', 'country': 'UY'})
```

2. Using a list of dicts:

```
request_with_cookies = Request(url="http://www.example.com",
                               cookies=[{'name': 'currency',
                                         'value': 'USD',
                                         'domain': 'example.com',
                                         'path': '/currency'}])
```

The latter form allows for customizing the domain and path attributes of the cookie. These is only useful if the cookies are saved for later requests.

When some site returns cookies (in a response) those are stored in the cookies for that domain and will be sent again in future requests. That's the typical behaviour of any regular web browser. However, if, for some reason, you want to avoid merging with existing cookies you can instruct Scrapy to do so by setting the `dont_merge_cookies` key in the `Request.meta`.

Example of request without merging cookies:

```
request_with_cookies = Request(url="http://www.example.com",
                               cookies={'currency': 'USD', 'country': 'UY'},
                               meta={'dont_merge_cookies': True})
```

For more info see [CookiesMiddleware](#).

- **encoding** (*string*) – the encoding of this request (defaults to 'utf-8'). This encoding will be used to percent-encode the URL and to convert the body to `str` (if given as unicode).
- **priority** (*int*) – the priority of this request (defaults to 0). The priority is used by the scheduler to define the order used to process requests.
- **dont_filter** (*boolean*) – indicates that this request should not be filtered by the scheduler. This is used when you want to perform an identical request multiple times, to ignore the duplicates filter. Use it with care, or you will get into crawling loops. Default to `False`.
- **callback** (*callable*) – the function that will be called with the response of this request (once its downloaded) as its first parameter. For more information see [Passing additional data to callback functions](#) below. If a Request doesn't specify a callback, the spider's `parse()` method will be used.
- **errback** (*callable*) – a function that will be called if any exception was raised while processing the request. This includes pages that failed with 404 HTTP errors and such. It receives a [Twisted Failure](#) instance as first parameter.

url

A string containing the URL of this request. Keep in mind that this attribute contains the escaped URL, so it can differ from the URL passed in the constructor.

This attribute is read-only. To change the URL of a Request use `replace()`.

method

A string representing the HTTP method in the request. This is guaranteed to be uppercase. Example: "GET", "POST", "PUT", etc

headers

A dictionary-like object which contains the request headers.

body

A str that contains the request body.

This attribute is read-only. To change the body of a Request use `replace()`.

meta

A dict that contains arbitrary metadata for this request. This dict is empty for new Requests, and is usually populated by different Scrapy components (extensions, middlewares, etc). So the data contained in this dict depends on the extensions you have enabled.

See *Request.meta special keys* for a list of special meta keys recognized by Scrapy.

This dict is *shallow copied* when the request is cloned using the `copy()` or `replace()` methods, and can also be accessed, in your spider, from the `response.meta` attribute.

copy()

Return a new Request which is a copy of this Request. See also: *Passing additional data to callback functions*.

replace([url, method, headers, body, cookies, meta, encoding, dont_filter, callback, errback])

Return a Request object with the same members, except for those members given new values by whichever keyword arguments are specified. The attribute `Request.meta` is copied by default (unless a new value is given in the `meta` argument). See also *Passing additional data to callback functions*.

Passing additional data to callback functions

The callback of a request is a function that will be called when the response of that request is downloaded. The callback function will be called with the downloaded `Response` object as its first argument.

Example:

```
def parse_page1(self, response):
    return Request("http://www.example.com/some_page.html",
                  callback=self.parse_page2)

def parse_page2(self, response):
    # this would log http://www.example.com/some_page.html
    self.log("Visited %s" % response.url)
```

In some cases you may be interested in passing arguments to those callback functions so you can receive the arguments later, in the second callback. You can use the `Request.meta` attribute for that.

Here's an example of how to pass an item using this mechanism, to populate different fields from different pages:

```
def parse_page1(self, response):
    item = MyItem()
    item['main_url'] = response.url
    request = Request("http://www.example.com/some_page.html",
                     callback=self.parse_page2)
    request.meta['item'] = item
    return request

def parse_page2(self, response):
    item = response.meta['item']
    item['other_url'] = response.url
    return item
```

7.1.2 Request.meta special keys

The `Request.meta` attribute can contain any arbitrary data, but there are some special keys recognized by Scrapy and its built-in extensions.

Those are:

- `dont_redirect`
- `dont_retry`
- `handle_httpstatus_list`
- `dont_merge_cookies` (see `cookies` parameter of `Request` constructor)
- `cookiejar`
- `redirect_urls`

7.1.3 Request subclasses

Here is the list of built-in `Request` subclasses. You can also subclass it to implement your own custom functionality.

FormRequest objects

The `FormRequest` class extends the base `Request` with functionality for dealing with HTML forms. It uses `lxml.html.forms` to pre-populate form fields with form data from `Response` objects.

class scrapy.http.**FormRequest** (*url*[, *formdata*, ...])

The `FormRequest` class adds a new argument to the constructor. The remaining arguments are the same as for the `Request` class and are not documented here.

Parameters `formdata` (*dict or iterable of tuples*) – is a dictionary (or iterable of (key, value) tuples) containing HTML Form data which will be url-encoded and assigned to the body of the request.

The `FormRequest` objects support the following class method in addition to the standard `Request` methods:

classmethod `from_response` (*response*[, *formname*=None, *formnumber*=0, *formdata*=None, *dont_click*=False, ...])

Returns a new `FormRequest` object with its form field values pre-populated with those found in the HTML `<form>` element contained in the given response. For an example see [Using FormRequest.from_response\(\) to simulate a user login](#).

The policy is to automatically simulate a click, by default, on any form control that looks clickable, like a `<input type="submit">`. Even though this is quite convenient, and often the desired behaviour, sometimes it can cause problems which could be hard to debug. For example, when working with forms that are filled and/or submitted using javascript, the default `from_response()` behaviour may not be the most appropriate. To disable this behaviour you can set the `dont_click` argument to `True`. Also, if you want to change the control clicked (instead of disabling it) you can also use the `clickdata` argument.

Parameters

- **response** (`Response` object) – the response containing a HTML form which will be used to pre-populate the form fields
- **formname** (*string*) – if given, the form with name attribute set to this value will be used. Otherwise, `formnumber` will be used for selecting the form.

- **formnumber** (*integer*) – the number of form to use, when the response contains multiple forms. The first one (and also the default) is 0.
- **formdata** (*dict*) – fields to override in the form data. If a field was already present in the response `<form>` element, its value is overridden by the one passed in this parameter.
- **dont_click** (*boolean*) – If True, the form data will be submitted without clicking in any element.

The other parameters of this class method are passed directly to the `FormRequest` constructor. New in version 0.10.3: The `formname` parameter.

Request usage examples

Using FormRequest to send data via HTTP POST

If you want to simulate a HTML Form POST in your spider and send a couple of key-value fields, you can return a `FormRequest` object (from your spider) like this:

```
return [FormRequest(url="http://www.example.com/post/action",
                    formdata={'name': 'John Doe', 'age': '27'},
                    callback=self.after_post)]
```

Using FormRequest.from_response() to simulate a user login

It is usual for web sites to provide pre-populated form fields through `<input type="hidden">` elements, such as session related data or authentication tokens (for login pages). When scraping, you'll want these fields to be automatically pre-populated and only override a couple of them, such as the user name and password. You can use the `FormRequest.from_response()` method for this job. Here's an example spider which uses it:

```
class LoginSpider(BaseSpider):
    name = 'example.com'
    start_urls = ['http://www.example.com/users/login.php']

    def parse(self, response):
        return [FormRequest.from_response(response,
                                         formdata={'username': 'john', 'password': 'secret'},
                                         callback=self.after_login)]

    def after_login(self, response):
        # check login succeed before going on
        if "authentication failed" in response.body:
            self.log("Login failed", level=log.ERROR)
            return

        # continue scraping with authenticated session...
```

7.1.4 Response objects

`class scrapy.http.Response(url[, status=200, headers, body, flags])`

A `Response` object represents an HTTP response, which is usually downloaded (by the Downloader) and fed to the Spiders for processing.

Parameters

- **url** (*string*) – the URL of this response
- **headers** (*dict*) – the headers of this response. The dict values can be strings (for single valued headers) or lists (for multi-valued headers).
- **status** (*integer*) – the HTTP status of the response. Defaults to 200.
- **body** (*str*) – the response body. It must be str, not unicode, unless you’re using a encoding-aware *Response subclass*, such as `TextResponse`.
- **meta** (*dict*) – the initial values for the `Response.meta` attribute. If given, the dict will be shallow copied.
- **flags** (*list*) – is a list containing the initial values for the `Response.flags` attribute. If given, the list will be shallow copied.

url

A string containing the URL of the response.

This attribute is read-only. To change the URL of a Response use `replace()`.

status

An integer representing the HTTP status of the response. Example: 200, 404.

headers

A dictionary-like object which contains the response headers.

body

A str containing the body of this Response. Keep in mind that `Response.body` is always a str. If you want the unicode version use `TextResponse.body_as_unicode()` (only available in `TextResponse` and subclasses).

This attribute is read-only. To change the body of a Response use `replace()`.

request

The `Request` object that generated this response. This attribute is assigned in the Scrapy engine, after the response and the request have passed through all *Downloader Middlewares*. In particular, this means that:

- HTTP redirections will cause the original request (to the URL before redirection) to be assigned to the redirected response (with the final URL after redirection).
- `Response.request.url` doesn’t always equal `Response.url`
- This attribute is only available in the spider code, and in the *Spider Middlewares*, but not in *Downloader Middlewares* (although you have the `Request` available there by other means) and handlers of the `response_downloaded` signal.

meta

A shortcut to the `Request.meta` attribute of the `Response.request` object (ie. `self.request.meta`).

Unlike the `Response.request` attribute, the `Response.meta` attribute is propagated along redirects and retries, so you will get the original `Request.meta` sent from your spider.

See Also:

`Request.meta` attribute

flags

A list that contains flags for this response. Flags are labels used for tagging Responses. For example: ‘cached’, ‘redirected’, etc. And they’re shown on the string representation of the Response (`__str__` method) which is used by the engine for logging.

copy()

Returns a new Response which is a copy of this Response.

replace() (*url, status, headers, body, meta, flags, cls*)

Returns a Response object with the same members, except for those members given new values by whichever keyword arguments are specified. The attribute `Response.meta` is copied by default (unless a new value is given in the `meta` argument).

7.1.5 Response subclasses

Here is the list of available built-in Response subclasses. You can also subclass the Response class to implement your own functionality.

TextResponse objects

class scrapy.http.TextResponse(*url[, encoding[, ...]]*)

TextResponse objects adds encoding capabilities to the base Response class, which is meant to be used only for binary data, such as images, sounds or any media file.

TextResponse objects support a new constructor argument, in addition to the base Response objects. The remaining functionality is the same as for the Response class and is not documented here.

Parameters `encoding` (*string*) – is a string which contains the encoding to use for this response. If you create a TextResponse object with a unicode body, it will be encoded using this encoding (remember the body attribute is always a string). If `encoding` is `None` (default value), the encoding will be looked up in the response headers and body instead.

TextResponse objects support the following attributes in addition to the standard Response ones:

encoding

A string with the encoding of this response. The encoding is resolved by trying the following mechanisms, in order:

- 1.the encoding passed in the constructor *encoding* argument
- 2.the encoding declared in the Content-Type HTTP header. If this encoding is not valid (ie. unknown), it is ignored and the next resolution mechanism is tried.
- 3.the encoding declared in the response body. The TextResponse class doesn't provide any special functionality for this. However, the HtmlResponse and XmlResponse classes do.
- 4.the encoding inferred by looking at the response body. This is the more fragile method but also the last one tried.

TextResponse objects support the following methods in addition to the standard Response ones:

body_as_unicode()

Returns the body of the response as unicode. This is equivalent to:

```
response.body.decode(response.encoding)
```

But **not** equivalent to:

```
unicode(response.body)
```

Since, in the latter case, you would be using you system default encoding (typically *ascii*) to convert the body to uniode, instead of the response encoding.

HtmlResponse objects

`class scrapy.http.HtmlResponse(url[, ...])`

The `HtmlResponse` class is a subclass of `TextResponse` which adds encoding auto-discovering support by looking into the HTML `meta http-equiv` attribute. See `TextResponse.encoding`.

XmlResponse objects

`class scrapy.http.XmlResponse(url[, ...])`

The `XmlResponse` class is a subclass of `TextResponse` which adds encoding auto-discovering support by looking into the XML declaration line. See `TextResponse.encoding`.

7.2 Settings

The Scrapy settings allows you to customize the behaviour of all Scrapy components, including the core, extensions, pipelines and spiders themselves.

The infrastructure of the settings provides a global namespace of key-value mappings that the code can use to pull configuration values from. The settings can be populated through different mechanisms, which are described below.

The settings are also the mechanism for selecting the currently active Scrapy project (in case you have many).

For a list of available built-in settings see: *Built-in settings reference*.

7.2.1 Designating the settings

When you use Scrapy, you have to tell it which settings you're using. You can do this by using an environment variable, `SCRAPY_SETTINGS_MODULE`.

The value of `SCRAPY_SETTINGS_MODULE` should be in Python path syntax, e.g. `myproject.settings`. Note that the settings module should be on the Python [import search path](#).

7.2.2 Populating the settings

Settings can be populated using different mechanisms, each of which having a different precedence. Here is the list of them in decreasing order of precedence:

1. Global overrides (most precedence)
2. Project settings module
3. Default settings per-command
4. Default global settings (less precedence)

These mechanisms are described in more detail below.

1. Global overrides

Global overrides are the ones that take most precedence, and are usually populated by command-line options. You can also override one (or more) settings from command line using the `-s` (or `--set`) command line option.

For more information see the `overrides` Settings attribute.

Example:

```
scrapy crawl domain.com -s LOG_FILE=scrapy.log
```

2. Project settings module

The project settings module is the standard configuration file for your Scrapy project. It's where most of your custom settings will be populated. For example: `myproject.settings`.

3. Default settings per-command

Each *Scrapy tool* command can have its own default settings, which override the global default settings. Those custom command settings are specified in the `default_settings` attribute of the command class.

4. Default global settings

The global defaults are located in the `scrapy.settings.default_settings` module and documented in the *Built-in settings reference* section.

7.2.3 How to access settings

Settings can be accessed through the `scrapy.crawler.Crawler.settings` attribute of the Crawler that is passed to `from_crawler` method in extensions and middlewares:

```
class MyExtension(object):

    @classmethod
    def from_crawler(cls, crawler):
        settings = crawler.settings
        if settings['LOG_ENABLED']:
            print "log is enabled!"
```

In other words, settings can be accessed like a dict, but it's usually preferred to extract the setting in the format you need it to avoid type errors. In order to do that you'll have to use one of the methods provided the *Settings API*.

7.2.4 Rationale for setting names

Setting names are usually prefixed with the component that they configure. For example, proper setting names for a fictional `robots.txt` extension would be `ROBOTSTXT_ENABLED`, `ROBOTSTXT_OBEY`, `ROBOTSTXT_CACHEDIR`, etc.

7.2.5 Built-in settings reference

Here's a list of all available Scrapy settings, in alphabetical order, along with their default values and the scope where they apply.

The scope, where available, shows where the setting is being used, if it's tied to any particular component. In that case the module of that component will be shown, typically an extension, middleware or pipeline. It also means that the component must be enabled in order for the setting to have any effect.

AWS_ACCESS_KEY_ID

Default: None

The AWS access key used by code that requires access to [Amazon Web services](#), such as the *S3 feed storage backend*.

AWS_SECRET_ACCESS_KEY

Default: None

The AWS secret key used by code that requires access to [Amazon Web services](#), such as the *S3 feed storage backend*.

BOT_NAME

Default: 'scrapybot'

The name of the bot implemented by this Scrapy project (also known as the project name). This will be used to construct the User-Agent by default, and also for logging.

It's automatically populated with your project name when you create your project with the `startproject` command.

CONCURRENT_ITEMS

Default: 100

Maximum number of concurrent items (per response) to process in parallel in the Item Processor (also known as the *Item Pipeline*).

CONCURRENT_REQUESTS

Default: 16

The maximum number of concurrent (ie. simultaneous) requests that will be performed by the Scrapy downloader.

CONCURRENT_REQUESTS_PER_DOMAIN

Default: 8

The maximum number of concurrent (ie. simultaneous) requests that will be performed to any single domain.

CONCURRENT_REQUESTS_PER_IP

Default: 0

The maximum number of concurrent (ie. simultaneous) requests that will be performed to any single IP. If non-zero, the `CONCURRENT_REQUESTS_PER_DOMAIN` setting is ignored, and this one is used instead. In other words, concurrency limits will be applied per IP, not per domain.

DEFAULT_ITEM_CLASS

Default: `' scrapy.item.Item '`

The default class that will be used for instantiating items in the *the Scrapy shell*.

DEFAULT_REQUEST_HEADERS

Default:

```
{
    'Accept': 'text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8',
    'Accept-Language': 'en',
}
```

The default headers used for Scrapy HTTP Requests. They're populated in the `DefaultHeadersMiddleware`.

DEPTH_LIMIT

Default: 0

The maximum depth that will be allowed to crawl for any site. If zero, no limit will be imposed.

DEPTH_PRIORITY

Default: 0

An integer that is used to adjust the request priority based on its depth.

If zero, no priority adjustment is made from depth.

DEPTH_STATS

Default: `True`

Whether to collect maximum depth stats.

DEPTH_STATS_VERBOSE

Default: `False`

Whether to collect verbose depth stats. If this is enabled, the number of requests for each depth is collected in the stats.

DNSCACHE_ENABLED

Default: `True`

Whether to enable DNS in-memory cache.

DOWNLOADER_DEBUG

Default: `False`

Whether to enable the Downloader debugging mode.

DOWNLOADER_MIDDLEWARES

Default:: `{}`

A dict containing the downloader middlewares enabled in your project, and their orders. For more info see [Activating a downloader middleware](#).

DOWNLOADER_MIDDLEWARES_BASE

Default:

```
{
    'scrapy.contrib.downloadermiddleware.robotstxt.RobotsTxtMiddleware': 100,
    'scrapy.contrib.downloadermiddleware.httpauth.HttpAuthMiddleware': 300,
    'scrapy.contrib.downloadermiddleware.downloadtimeout.DownloadTimeoutMiddleware': 350,
    'scrapy.contrib.downloadermiddleware.useragent.UserAgentMiddleware': 400,
    'scrapy.contrib.downloadermiddleware.retry.RetryMiddleware': 500,
    'scrapy.contrib.downloadermiddleware.defaultheaders.DefaultHeadersMiddleware': 550,
    'scrapy.contrib.downloadermiddleware.redirect.RedirectMiddleware': 600,
    'scrapy.contrib.downloadermiddleware.cookies.CookiesMiddleware': 700,
    'scrapy.contrib.downloadermiddleware.httpproxy.HttpProxyMiddleware': 750,
    'scrapy.contrib.downloadermiddleware.httpcompression.HttpCompressionMiddleware': 800,
    'scrapy.contrib.downloadermiddleware.chunked.ChunkedTransferMiddleware': 830,
    'scrapy.contrib.downloadermiddleware.stats.DownloaderStats': 850,
    'scrapy.contrib.downloadermiddleware.httpcache.HttpCacheMiddleware': 900,
}
```

A dict containing the downloader middlewares enabled by default in Scrapy. You should never modify this setting in your project, modify `DOWNLOADER_MIDDLEWARES` instead. For more info see [Activating a downloader middleware](#).

DOWNLOADER_STATS

Default: `True`

Whether to enable downloader stats collection.

DOWNLOAD_DELAY

Default: `0`

The amount of time (in secs) that the downloader should wait before downloading consecutive pages from the same spider. This can be used to throttle the crawling speed to avoid hitting servers too hard. Decimal numbers are supported. Example:

```
DOWNLOAD_DELAY = 0.25      # 250 ms of delay
```


This setting is also affected by the `RANDOMIZE_DOWNLOAD_DELAY` setting (which is enabled by default). By default, Scrapy doesn't wait a fixed amount of time between requests, but uses a random interval between 0.5 and 1.5 * `DOWNLOAD_DELAY`.

You can also change this setting per spider.

DOWNLOAD_HANDLERS

Default: `{}`

A dict containing the request downloader handlers enabled in your project. See `DOWNLOAD_HANDLERS_BASE` for example format.

DOWNLOAD_HANDLERS_BASE

Default:

```
{
    'file': 'scrapy.core.downloader.handlers.file.FileDownloadHandler',
    'http': 'scrapy.core.downloader.handlers.http.HttpDownloadHandler',
    'https': 'scrapy.core.downloader.handlers.http.HttpDownloadHandler',
    's3': 'scrapy.core.downloader.handlers.s3.S3DownloadHandler',
}
```

A dict containing the request download handlers enabled by default in Scrapy. You should never modify this setting in your project, modify `DOWNLOAD_HANDLERS` instead.

DOWNLOAD_TIMEOUT

Default: 180

The amount of time (in secs) that the downloader will wait before timing out.

DUPEFILTER_CLASS

Default: `'scrapy.dupefilter.RFPDupeFilter'`

The class used to detect and filter duplicate requests.

The default (`RFPDupeFilter`) filters based on request fingerprint using the `scrapy.utils.request.request_fingerprint` function.

EDITOR

Default: *depends on the environment*

The editor to use for editing spiders with the `edit` command. It defaults to the `EDITOR` environment variable, if set. Otherwise, it defaults to `vi` (on Unix systems) or the `IDLE` editor (on Windows).

EXTENSIONS

Default:: `{}`

A dict containing the extensions enabled in your project, and their orders.

EXTENSIONS_BASE

Default:

```
{
    'scrapy.contrib.corestats.CoreStats': 0,
    'scrapy.webservice.WebService': 0,
    'scrapy.telnet.TelnetConsole': 0,
    'scrapy.contrib.memusage.MemoryUsage': 0,
    'scrapy.contrib.memdebug.MemoryDebugger': 0,
    'scrapy.contrib.closespider.CloseSpider': 0,
    'scrapy.contrib.feedexport.FeedExporter': 0,
    'scrapy.contrib.logstats.LogStats': 0,
    'scrapy.contrib.spiderstate.SpiderState': 0,
    'scrapy.contrib.throttle.AutoThrottle': 0,
}
```

The list of available extensions. Keep in mind that some of them need to be enabled through a setting. By default, this setting contains all stable built-in extensions.

For more information See the *extensions user guide* and the *list of available extensions*.

ITEM_PIPELINES

Default: []

The item pipelines to use (a list of classes).

Example:

```
ITEM_PIPELINES = [
    'mybot.pipeline.validate.ValidateMyItem',
    'mybot.pipeline.validate.StoreMyItem'
]
```

LOG_ENABLED

Default: True

Whether to enable logging.

LOG_ENCODING

Default: 'utf-8'

The encoding to use for logging.

LOG_FILE

Default: None

File name to use for logging output. If None, standard error will be used.

LOG_LEVEL

Default: 'DEBUG'

Minimum level to log. Available levels are: CRITICAL, ERROR, WARNING, INFO, DEBUG. For more info see [Logging](#).

LOG_STDOUT

Default: False

If True, all standard output (and error) of your process will be redirected to the log. For example if you print 'hello' it will appear in the Scrapy log.

MEMDEBUG_ENABLED

Default: False

Whether to enable memory debugging.

MEMDEBUG_NOTIFY

Default: []

When memory debugging is enabled a memory report will be sent to the specified addresses if this setting is not empty, otherwise the report will be written to the log.

Example:

```
MEMDEBUG_NOTIFY = ['user@example.com']
```

MEMUSAGE_ENABLED

Default: False

Scope: `scrapy.contrib.memusage`

Whether to enable the memory usage extension that will shutdown the Scrapy process when it exceeds a memory limit, and also notify by email when that happened.

See [Memory usage extension](#).

MEMUSAGE_LIMIT_MB

Default: 0

Scope: `scrapy.contrib.memusage`

The maximum amount of memory to allow (in megabytes) before shutting down Scrapy (if MEMUSAGE_ENABLED is True). If zero, no check will be performed.

See [Memory usage extension](#).

MEMUSAGE_NOTIFY_MAIL

Default: `False`

Scope: `scrapy.contrib.memusage`

A list of emails to notify if the memory limit has been reached.

Example:

```
MEMUSAGE_NOTIFY_MAIL = ['user@example.com']
```

See *Memory usage extension*.

MEMUSAGE_REPORT

Default: `False`

Scope: `scrapy.contrib.memusage`

Whether to send a memory usage report after each spider has been closed.

See *Memory usage extension*.

MEMUSAGE_WARNING_MB

Default: `0`

Scope: `scrapy.contrib.memusage`

The maximum amount of memory to allow (in megabytes) before sending a warning email notifying about it. If zero, no warning will be produced.

NEWSPIDER_MODULE

Default: `"`

Module where to create new spiders using the `genspider` command.

Example:

```
NEWSPIDER_MODULE = 'mybot.spiders_dev'
```

RANDOMIZE_DOWNLOAD_DELAY

Default: `True`

If enabled, Scrapy will wait a random amount of time (between 0.5 and $1.5 * \text{DOWNLOAD_DELAY}$) while fetching requests from the same spider.

This randomization decreases the chance of the crawler being detected (and subsequently blocked) by sites which analyze requests looking for statistically significant similarities in the time between their requests.

The randomization policy is the same used by `wget --random-wait` option.

If `DOWNLOAD_DELAY` is zero (default) this option has no effect.

REDIRECT_MAX_TIMES

Default: 20

Defines the maximum times a request can be redirected. After this maximum the request's response is returned as is. We used Firefox default value for the same task.

REDIRECT_MAX_METAREFRESH_DELAY

Default: 100

Some sites use meta-refresh for redirecting to a session expired page, so we restrict automatic redirection to a maximum delay (in seconds)

REDIRECT_PRIORITY_ADJUST

Default: +2

Adjust redirect request priority relative to original request. A negative priority adjust means more priority.

ROBOTSTXT_OBEY

Default: False

Scope: `scrapy.contrib.downloadermiddleware.robotstxt`

If enabled, Scrapy will respect robots.txt policies. For more information see [RobotsTxtMiddleware](#)

SCHEDULER

Default: `'scrapy.core.scheduler.Scheduler'`

The scheduler to use for crawling.

SPIDER_CONTRACTS

Default:: `{}`

A dict containing the scrapy contracts enabled in your project, used for testing spiders. For more info see [Spiders Contracts](#).

SPIDER_CONTRACTS_BASE

Default:

```
{
    'scrapy.contracts.default.UrlContract' : 1,
    'scrapy.contracts.default.ReturnsContract' : 2,
    'scrapy.contracts.default.ScrapesContract' : 3,
}
```

A dict containing the scrapy contracts enabled by default in Scrapy. You should never modify this setting in your project, modify `SPIDER_CONTRACTS` instead. For more info see [Spiders Contracts](#).

SPIDER_MIDDLEWARES

Default: {}

A dict containing the spider middlewares enabled in your project, and their orders. For more info see *Activating a spider middleware*.

SPIDER_MIDDLEWARES_BASE

Default:

```
{
    'scrapy.contrib.spidermiddleware.httperror.HttpErrorMiddleware': 50,
    'scrapy.contrib.spidermiddleware.offsite.OffsiteMiddleware': 500,
    'scrapy.contrib.spidermiddleware.referer.RefererMiddleware': 700,
    'scrapy.contrib.spidermiddleware.urllength.UrlLengthMiddleware': 800,
    'scrapy.contrib.spidermiddleware.depth.DepthMiddleware': 900,
}
```

A dict containing the spider middlewares enabled by default in Scrapy. You should never modify this setting in your project, modify `SPIDER_MIDDLEWARES` instead. For more info see *Activating a spider middleware*.

SPIDER_MODULES

Default: []

A list of modules where Scrapy will look for spiders.

Example:

```
SPIDER_MODULES = ['mybot.spiders_prod', 'mybot.spiders_dev']
```

STATS_CLASS

Default: 'scrapy.statscol.MemoryStatsCollector'

The class to use for collecting stats, who must implement the *Stats Collector API*.

STATS_DUMP

Default: True

Dump the *Scrapy stats* (to the Scrapy log) once the spider finishes.

For more info see: *Stats Collection*.

STATSMAILER_RCPTS

Default: [] (empty list)

Send Scrapy stats after spiders finish scraping. See `StatsMailer` for more info.

TELNETCONSOLE_ENABLED

Default: `True`

A boolean which specifies if the *telnet console* will be enabled (provided its extension is also enabled).

TELNETCONSOLE_PORT

Default: `[6023, 6073]`

The port range to use for the telnet console. If set to `None` or `0`, a dynamically assigned port is used. For more info see *Telnet Console*.

TEMPLATES_DIR

Default: `templates` dir inside scrapy module

The directory where to look for templates when creating new projects with `startproject` command.

URLLENGTH_LIMIT

Default: `2083`

Scope: `contrib.spidermiddleware.urllength`

The maximum URL length to allow for crawled URLs. For more information about the default value for this setting see: <http://www.boutell.com/newfaq/misc/urllength.html>

USER_AGENT

Default: `"Scrapy/VERSION (+http://scrapy.org)"`

The default User-Agent to use when crawling, unless overridden.

7.3 Signals

Scrapy uses signals extensively to notify when certain events occur. You can catch some of those signals in your Scrapy project (using an *extension*, for example) to perform additional tasks or extend Scrapy to add functionality not provided out of the box.

Even though signals provide several arguments, the handlers that catch them don't need to accept all of them - the signal dispatching mechanism will only deliver the arguments that the handler receives.

You can connect to signals (or send your own) through the *Signals API*.

7.3.1 Deferred signal handlers

Some signals support returning *Twisted deferreds* from their handlers, see the *Built-in signals reference* below to know which ones.

7.3.2 Built-in signals reference

Here's the list of Scrapy built-in signals and their meaning.

engine_started

`scrapy.signals.engine_started()`

Sent when the Scrapy engine has started crawling.

This signal supports returning deferreds from their handlers.

Note: This signal may be fired *after* the `spider_opened` signal, depending on how the spider was started. So **don't** rely on this signal getting fired before `spider_opened`.

engine_stopped

`scrapy.signals.engine_stopped()`

Sent when the Scrapy engine is stopped (for example, when a crawling process has finished).

This signal supports returning deferreds from their handlers.

item_scraped

`scrapy.signals.item_scraped(item, response, spider)`

Sent when an item has been scraped, after it has passed all the *Item Pipeline* stages (without being dropped).

This signal supports returning deferreds from their handlers.

Parameters

- **item** (`Item` object) – the item scraped
- **response** (`Response` object) – the response from where the item was scraped
- **spider** (`BaseSpider` object) – the spider which scraped the item

item_dropped

`scrapy.signals.item_dropped(item, spider, exception)`

Sent after an item has been dropped from the *Item Pipeline* when some stage raised a `DropItem` exception.

This signal supports returning deferreds from their handlers.

Parameters

- **item** (`Item` object) – the item dropped from the *Item Pipeline*
- **spider** (`BaseSpider` object) – the spider which scraped the item
- **exception** (`DropItem` exception) – the exception (which must be a `DropItem` subclass) which caused the item to be dropped

spider_closed

`scrapy.signals.spider_closed(spider, reason)`

Sent after a spider has been closed. This can be used to release per-spider resources reserved on `spider_opened`.

This signal supports returning deferreds from their handlers.

Parameters

- **spider** (`BaseSpider` object) – the spider which has been closed
- **reason** (*str*) – a string which describes the reason why the spider was closed. If it was closed because the spider has completed scraping, the reason is `'finished'`. Otherwise, if the spider was manually closed by calling the `close_spider` engine method, then the reason is the one passed in the `reason` argument of that method (which defaults to `'cancelled'`). If the engine was shutdown (for example, by hitting Ctrl-C to stop it) the reason will be `'shutdown'`.

spider_opened

`scrapy.signals.spider_opened(spider)`

Sent after a spider has been opened for crawling. This is typically used to reserve per-spider resources, but can be used for any task that needs to be performed when a spider is opened.

This signal supports returning deferreds from their handlers.

Parameters **spider** (`BaseSpider` object) – the spider which has been opened

spider_idle

`scrapy.signals.spider_idle(spider)`

Sent when a spider has gone idle, which means the spider has no further:

- requests waiting to be downloaded
- requests scheduled
- items being processed in the item pipeline

If the idle state persists after all handlers of this signal have finished, the engine starts closing the spider. After the spider has finished closing, the `spider_closed` signal is sent.

You can, for example, schedule some requests in your `spider_idle` handler to prevent the spider from being closed.

This signal does not support returning deferreds from their handlers.

Parameters **spider** (`BaseSpider` object) – the spider which has gone idle

spider_error

`scrapy.signals.spider_error(failure, response, spider)`

Sent when a spider callback generates an error (ie. raises an exception).

Parameters

- **failure** (`Failure` object) – the exception raised as a Twisted `Failure` object
- **response** (`Response` object) – the response being processed when the exception was raised

- **spider** (`BaseSpider` object) – the spider which raised the exception

request_received

`scrapy.signals.request_received(request, spider)`

Sent when the engine receives a `Request` from a spider.

This signal does not support returning deferreds from their handlers.

Parameters

- **request** (`Request` object) – the request received
- **spider** (`BaseSpider` object) – the spider which generated the request

response_received

`scrapy.signals.response_received(response, request, spider)`

Sent when the engine receives a new `Response` from the downloader.

This signal does not support returning deferreds from their handlers.

Parameters

- **response** (`Response` object) – the response received
- **request** (`Request` object) – the request that generated the response
- **spider** (`BaseSpider` object) – the spider for which the response is intended

response_downloaded

`scrapy.signals.response_downloaded(response, request, spider)`

Sent by the downloader right after a `HTTPResponse` is downloaded.

This signal does not support returning deferreds from their handlers.

Parameters

- **response** (`Response` object) – the response downloaded
- **request** (`Request` object) – the request that generated the response
- **spider** (`BaseSpider` object) – the spider for which the response is intended

7.4 Exceptions

7.4.1 Built-in Exceptions reference

Here's a list of all exceptions included in Scrapy and their usage.

DropItem

exception `scrapy.exceptions.DropItem`

The exception that must be raised by item pipeline stages to stop processing an Item. For more information see [Item Pipeline](#).

CloseSpider

exception scrapy.exceptions.**CloseSpider** (*reason='cancelled'*)

This exception can be raised from a spider callback to request the spider to be closed/stopped. Supported arguments:

Parameters **reason** (*str*) – the reason for closing

For example:

```
def parse_page(self, response):
    if 'Bandwidth exceeded' in response.body:
        raise CloseSpider('bandwidth_exceeded')
```

IgnoreRequest

exception scrapy.exceptions.**IgnoreRequest**

This exception can be raised by the Scheduler or any downloader middleware to indicate that the request should be ignored.

NotConfigured

exception scrapy.exceptions.**NotConfigured**

This exception can be raised by some components to indicate that they will remain disabled. Those components include:

- Extensions
- Item pipelines
- Downloader middlewares
- Spider middlewares

The exception must be raised in the component constructor.

NotSupported

exception scrapy.exceptions.**NotSupported**

This exception is raised to indicate an unsupported feature.

7.5 Item Exporters

Once you have scraped your Items, you often want to persist or export those items, to use the data in some other application. That is, after all, the whole purpose of the scraping process.

For this purpose Scrapy provides a collection of Item Exporters for different output formats, such as XML, CSV or JSON.

7.5.1 Using Item Exporters

If you are in a hurry, and just want to use an Item Exporter to output scraped data see the [Feed exports](#). Otherwise, if you want to know how Item Exporters work or need more custom functionality (not covered by the default exports), continue reading below.

In order to use an Item Exporter, you must instantiate it with its required args. Each Item Exporter requires different arguments, so check each exporter documentation to be sure, in [Built-in Item Exporters reference](#). After you have instantiated you exporter, you have to:

1. call the method `start_exporting()` in order to signal the beginning of the exporting process
2. call the `export_item()` method for each item you want to export
3. and finally call the `finish_exporting()` to signal the end of the exporting process

Here you can see an [Item Pipeline](#) which uses an Item Exporter to export scraped items to different files, one per spider:

```
from scrapy import signals
from scrapy.contrib.exporter import XmlItemExporter

class XmlExportPipeline(object):

    def __init__(self):
        self.files = {}

    @classmethod
    def from_crawler(cls, crawler):
        pipeline = cls()
        crawler.signals.connect(pipeline.spider_opened, signals.spider_opened)
        crawler.signals.connect(pipeline.spider_closed, signals.spider_closed)
        return pipeline

    def spider_opened(self, spider):
        file = open('%s_products.xml' % spider.name, 'w+b')
        self.files[spider] = file
        self.exporter = XmlItemExporter(file)
        self.exporter.start_exporting()

    def spider_closed(self, spider):
        self.exporter.finish_exporting()
        file = self.files.pop(spider)
        file.close()

    def process_item(self, item, spider):
        self.exporter.export_item(item)
        return item
```

7.5.2 Serialization of item fields

By default, the field values are passed unmodified to the underlying serialization library, and the decision of how to serialize them is delegated to each particular serialization library.

However, you can customize how each field value is serialized *before it is passed to the serialization library*.

There are two ways to customize how a field will be serialized, which are described next.

1. Declaring a serializer in the field

You can declare a serializer in the *field metadata*. The serializer must be a callable which receives a value and returns its serialized form.

Example:

```
from scrapy.item import Item, Field

def serialize_price(value):
    return '$ %s' % str(value)

class Product(Item):
    name = Field()
    price = Field(serializer=serialize_price)
```

2. Overriding the `serialize_field()` method

You can also override the `serialize()` method to customize how your field value will be exported.

Make sure you call the base class `serialize()` method after your custom code.

Example:

```
from scrapy.contrib.exporter import XmlItemExporter

class ProductXmlExporter(XmlItemExporter):

    def serialize_field(self, field, name, value):
        if field == 'price':
            return '$ %s' % str(value)
        return super(Product, self).serialize_field(field, name, value)
```

7.5.3 Built-in Item Exporters reference

Here is a list of the Item Exporters bundled with Scrapy. Some of them contain output examples, which assume you're exporting these two items:

```
Item(name='Color TV', price='1200')
Item(name='DVD player', price='200')
```

BaseItemExporter

```
class scrapy.contrib.exporter.BaseItemExporter(fields_to_export=None,
                                              port_empty_fields=False, encoding='utf-8')
```

This is the (abstract) base class for all Item Exporters. It provides support for common features used by all (concrete) Item Exporters, such as defining what fields to export, whether to export empty fields, or which encoding to use.

These features can be configured through the constructor arguments which populate their respective instance attributes: `fields_to_export`, `export_empty_fields`, `encoding`.

export_item(item)

Exports the given item. This method must be implemented in subclasses.

serialize_field (*field, name, value*)

Return the serialized value for the given field. You can override this method (in your custom Item Exporters) if you want to control how a particular field or value will be serialized/exported.

By default, this method looks for a serializer *declared in the item field* and returns the result of applying that serializer to the value. If no serializer is found, it returns the value unchanged except for unicode values which are encoded to `str` using the encoding declared in the `encoding` attribute.

Parameters

- **field** (`Field` object) – the field being serialized
- **name** (*str*) – the name of the field being serialized
- **value** – the value being serialized

start_exporting ()

Signal the beginning of the exporting process. Some exporters may use this to generate some required header (for example, the `XmlItemExporter`). You must call this method before exporting any items.

finish_exporting ()

Signal the end of the exporting process. Some exporters may use this to generate some required footer (for example, the `XmlItemExporter`). You must always call this method after you have no more items to export.

fields_to_export

A list with the name of the fields that will be exported, or `None` if you want to export all fields. Defaults to `None`.

Some exporters (like `CsvItemExporter`) respect the order of the fields defined in this attribute.

export_empty_fields

Whether to include empty/unpopulated item fields in the exported data. Defaults to `False`. Some exporters (like `CsvItemExporter`) ignore this attribute and always export all empty fields.

encoding

The encoding that will be used to encode unicode values. This only affects unicode values (which are always serialized to `str` using this encoding). Other value types are passed unchanged to the specific serialization library.

XmlItemExporter

```
class scrapy.contrib.exporter.XmlItemExporter (file, item_element='item',
                                              root_element='items', **kwargs)
```

Exports Items in XML format to the specified file object.

Parameters

- **file** – the file-like object to use for exporting the data.
- **root_element** (*str*) – The name of root element in the exported XML.
- **item_element** (*str*) – The name of each item element in the exported XML.

The additional keyword arguments of this constructor are passed to the `BaseItemExporter` constructor.

A typical output of this exporter would be:

```
<?xml version="1.0" encoding="utf-8"?>
<items>
  <item>
    <name>Color TV</name>
```

```

    <price>1200</price>
  </item>
  <item>
    <name>DVD player</name>
    <price>200</price>
  </item>
</items>

```

Unless overridden in the `serialize_field()` method, multi-valued fields are exported by serializing each value inside a `<value>` element. This is for convenience, as multi-valued fields are very common.

For example, the item:

```
Item(name=['John', 'Doe'], age='23')
```

Would be serialized as:

```

<?xml version="1.0" encoding="utf-8"?>
<items>
  <item>
    <name>
      <value>John</value>
      <value>Doe</value>
    </name>
    <age>23</age>
  </item>
</items>

```

CsvItemExporter

class scrapy.contrib.exporter.CsvItemExporter(*file*, *include_headers_line=True*, *join_multivalued=',', **kwargs*)

Exports Items in CSV format to the given file-like object. If the `fields_to_export` attribute is set, it will be used to define the CSV columns and their order. The `export_empty_fields` attribute has no effect on this exporter.

Parameters

- **file** – the file-like object to use for exporting the data.
- **include_headers_line** (*str*) – If enabled, makes the exporter output a header line with the field names taken from `BaseItemExporter.fields_to_export` or the first exported item fields.
- **join_multivalued** – The char (or chars) that will be used for joining multi-valued fields, if found.

The additional keyword arguments of this constructor are passed to the `BaseItemExporter` constructor, and the leftover arguments to the `csv.writer` constructor, so you can use any `csv.writer` constructor argument to customize this exporter.

A typical output of this exporter would be:

```

product,price
Color TV,1200
DVD player,200

```

PickleItemExporter

class scrapy.contrib.exporter.**PickleItemExporter** (*file*, *protocol=0*, ***kwargs*)
Exports Items in pickle format to the given file-like object.

Parameters

- **file** – the file-like object to use for exporting the data.
- **protocol** (*int*) – The pickle protocol to use.

For more information, refer to the [pickle module documentation](#).

The additional keyword arguments of this constructor are passed to the [BaseItemExporter](#) constructor.

Pickle isn't a human readable format, so no output examples are provided.

PprintItemExporter

class scrapy.contrib.exporter.**PprintItemExporter** (*file*, ***kwargs*)
Exports Items in pretty print format to the specified file object.

Parameters **file** – the file-like object to use for exporting the data.

The additional keyword arguments of this constructor are passed to the [BaseItemExporter](#) constructor.

A typical output of this exporter would be:

```
{'name': 'Color TV', 'price': '1200'}  
{'name': 'DVD player', 'price': '200'}
```

Longer lines (when present) are pretty-formatted.

JsonItemExporter

class scrapy.contrib.exporter.**JsonItemExporter** (*file*, ***kwargs*)
Exports Items in JSON format to the specified file-like object, writing all objects as a list of objects. The additional constructor arguments are passed to the [BaseItemExporter](#) constructor, and the leftover arguments to the [JSONEncoder](#) constructor, so you can use any [JSONEncoder](#) constructor argument to customize this exporter.

Parameters **file** – the file-like object to use for exporting the data.

A typical output of this exporter would be:

```
[{"name": "Color TV", "price": "1200"},  
{"name": "DVD player", "price": "200"}]
```

Warning: JSON is very simple and flexible serialization format, but it doesn't scale well for large amounts of data since incremental (aka. stream-mode) parsing is not well supported (if at all) among JSON parsers (on any language), and most of them just parse the entire object in memory. If you want the power and simplicity of JSON with a more stream-friendly format, consider using [JsonLinesItemExporter](#) instead, or splitting the output in multiple chunks.

JsonLinesItemExporter

class scrapy.contrib.exporter.JsonLinesItemExporter (file, **kwargs)

Exports Items in JSON format to the specified file-like object, writing one JSON-encoded item per line. The additional constructor arguments are passed to the [BaseItemExporter](#) constructor, and the leftover arguments to the [JSONEncoder](#) constructor, so you can use any [JSONEncoder](#) constructor argument to customize this exporter.

Parameters file – the file-like object to use for exporting the data.

A typical output of this exporter would be:

```
{"name": "Color TV", "price": "1200"}
{"name": "DVD player", "price": "200"}
```

Unlike the one produced by [JsonItemExporter](#), the format produced by this exporter is well suited for serializing large amounts of data.

Command line tool Learn about the command-line tool and see all [available commands](#).

Requests and Responses Understand the classes used to represent HTTP requests and responses.

Settings Learn how to configure Scrapy and see all [available settings](#).

Signals See all available signals and how to work with them.

Exceptions See all available exceptions and their meaning.

Item Exporters Quickly export your scraped items to a file (XML, CSV, etc).

All the rest

8.1 Release notes

8.1.1 0.16.4 (released 2013-01-23)

- fixes spelling errors in documentation ([commit 6d2b3aa](#))
- add doc about disabling an extension. refs #132 ([commit c90de33](#))
- Fixed error message formatting. `log.err()` doesn't support cool formatting and when error occurred, the message was: "ERROR: Error processing %(item)s" ([commit c16150c](#))
- lint and improve images pipeline error logging ([commit 56b45fc](#))
- fixed doc typos ([commit 243be84](#))
- add documentation topics: Broad Crawls & Common Practices ([commit 1fbb715](#))
- fix bug in scrapy parse command when spider is not specified explicitly. closes #209 ([commit c72e682](#))
- Update docs/topics/commands.rst ([commit 28eac7a](#))

8.1.2 0.16.3 (released 2012-12-07)

- Remove concurrency limitation when using download delays and still ensure inter-request delays are enforced ([commit 487b9b5](#))
- add error details when image pipeline fails ([commit 8232569](#))
- improve mac os compatibility ([commit 8dcf8aa](#))
- setup.py: use README.rst to populate long_description ([commit 7b5310d](#))
- doc: removed obsolete references to ClientForm ([commit 80f9bb6](#))
- correct docs for default storage backend ([commit 2aa491b](#))
- doc: removed broken proxyhub link from FAQ ([commit bdf61c4](#))
- Merge branch '0.16' of github.com:scrapy/scrapy into 0.16 ([commit d5087b0](#))
- Fixed docs typo in SpiderOpenCloseLogging example ([commit 7184094](#))

8.1.3 0.16.2 (released 2012-11-09)

- scrapy contracts: python2.6 compat ([commit a4a9199](#))
- scrapy contracts verbose option ([commit ec41673](#))
- proper unittest-like output for scrapy contracts ([commit 86635e4](#))
- added open_in_browser to debugging doc ([commit c9b690d](#))
- removed reference to global scrapy stats from settings doc ([commit dd55067](#))
- Fix SpiderState bug in Windows platforms ([commit 58998f4](#))

8.1.4 0.16.1 (released 2012-10-26)

- fixed LogStats extension, which got broken after a wrong merge before the 0.16 release ([commit 8c780fd](#))
- better backwards compatibility for scrapy.conf.settings ([commit 3403089](#))
- extended documentation on how to access crawler stats from extensions ([commit c4da0b5](#))
- removed .hgtags (no longer needed now that scrapy uses git) ([commit d52c188](#))
- fix dashes under rst headers ([commit fa4f7f9](#))
- set release date for 0.16.0 in news ([commit e292246](#))

8.1.5 0.16.0 (released 2012-10-18)

Scrapy changes:

- added *Spiders Contracts*, a mechanism for testing spiders in a formal/reproducible way
- added options `-o` and `-t` to the `runspider` command
- documented *AutoThrottle extension* and added to extensions installed by default. You still need to enable it with `AUTOTHROTTLE_ENABLED`
- major Stats Collection refactoring: removed separation of global/per-spider stats, removed stats-related signals (`stats_spider_opened`, etc). Stats are much simpler now, backwards compatibility is kept on the Stats Collector API and signals.
- added `process_start_requests()` method to spider middlewares
- dropped Signals singleton. Signals should now be accessed through the `Crawler.signals` attribute. See the signals documentation for more info.
- dropped Signals singleton. Signals should now be accessed through the `Crawler.signals` attribute. See the signals documentation for more info.
- dropped Stats Collector singleton. Stats can now be accessed through the `Crawler.stats` attribute. See the stats collection documentation for more info.
- documented *Core API*
- `lxml` is now the default selectors backend instead of `libxml2`
- ported `FormRequest.from_response()` to use `lxml` instead of `ClientForm`
- removed modules: `scrapy.xlib.BeautifulSoup` and `scrapy.xlib.ClientForm`
- SitemapSpider: added support for sitemap urls ending in `.xml` and `.xml.gz`, even if they advertise a wrong content type ([commit 10ed28b](#))

- StackTraceDump extension: also dump trackref live references ([commit fe2ce93](#))
- nested items now fully supported in JSON and JSONLines exporters
- added `cookiejar` Request meta key to support multiple cookie sessions per spider
- decoupled encoding detection code to `w3lib.encoding`, and ported Scrapy code to use that module
- dropped support for Python 2.5. See <http://blog.scrapy.org/scrapy-dropping-support-for-python-25>
- dropped support for Twisted 2.5
- added `REFERER_ENABLED` setting, to control referer middleware
- changed default user agent to: `Scrapy/VERSION (+http://scrapy.org)`
- removed (undocumented) `HTMLImageLinkExtractor` class from `scrapy.contrib.linkextractors.image`
- removed per-spider settings (to be replaced by instantiating multiple crawler objects)
- `USER_AGENT` spider attribute will no longer work, use `user_agent` attribute instead
- `DOWNLOAD_TIMEOUT` spider attribute will no longer work, use `download_timeout` attribute instead
- removed `ENCODING_ALIASES` setting, as encoding auto-detection has been moved to the `w3lib` library
- promoted *DjangoItem* to main contrib
- `LogFormatter` method now return dicts (instead of strings) to support lazy formatting ([issue 164](#), [commit dcef7b0](#))
- downloader handlers (`DOWNLOAD_HANDLERS` setting) now receive settings as the first argument of the constructor
- replaced memory usage accounting with (more portable) `resource` module, removed `scrapy.utils.memory` module
- removed signal: `scrapy.mail.mail_sent`
- removed `TRACK_REFS` setting, now *trackrefs* is always enabled
- DBM is now the default storage backend for HTTP cache middleware
- number of log messages (per level) are now tracked through Scrapy stats (stat name: `log_count/LEVEL`)
- number received responses are now tracked through Scrapy stats (stat name: `response_received_count`)
- removed `scrapy.log.started` attribute

Scrapyd changes:

- New Scrapyd API methods: *listjobs.json* and *cancel.json*
- New Scrapyd settings: *items_dir* and *jobs_to_keep*
- Items are now stored on disk using feed exports, and accessible through the Scrapyd web interface
- Support making Scrapyd listen into a specific IP address (see `bind_address` option)

8.1.6 0.14.4

- added precise to supported ubuntu distros ([commit b7e46df](#))
- fixed bug in json-rpc webservice reported in <https://groups.google.com/d/topic/scrapy-users/qgVBmFybNAQ/discussion>. also removed no longer supported 'run' command from `extras/scrapy-ws.py` ([commit 340fddb](#))
- meta tag attributes for content-type http equiv can be in any order. [#123](#) ([commit 0cb68af](#))

- replace “import Image” by more standard “from PIL import Image”. closes #88 ([commit 4d17048](#))
- return trial status as bin/runtests.sh exit value. #118 ([commit b7b2e7f](#))

8.1.7 0.14.3

- forgot to include pydispatch license. #118 ([commit fd85f9c](#))
- include egg files used by testsuite in source distribution. #118 ([commit c897793](#))
- update docstring in project template to avoid confusion with genspider command, which may be considered as an advanced feature. refs #107 ([commit 2548dcc](#))
- added note to docs/topics/firebug.rst about google directory being shut down ([commit 668e352](#))
- Merge branch ‘0.14’ of github.com:scrapy/scrapy into 0.14 ([commit 835d082](#))
- dont discard slot when empty, just save in another dict in order to recycle if needed again. ([commit 8e9f607](#))
- do not fail handling unicode xpath in lxml2 backed selectors ([commit b830e95](#))
- fixed minor mistake in Request objects documentation ([commit bf3c9ee](#))
- fixed minor defect in link extractors documentation ([commit ba14f38](#))
- removed some obsolete remaining code related to sqlite support in scrapy ([commit 0665175](#))

8.1.8 0.14.2

- move buffer pointing to start of file before computing checksum. refs #92 ([commit 6a5bef2](#))
- Compute image checksum before persisting images. closes #92 ([commit 9817df1](#))
- remove leaking references in cached failures ([commit 673a120](#))
- fixed bug in MemoryUsage extension: get_engine_status() takes exactly 1 argument (0 given) ([commit 11133e9](#))
- Merge branch ‘0.14’ of github.com:scrapy/scrapy into 0.14 ([commit 1627320](#))
- fixed struct.error on http compression middleware. closes #87 ([commit 1423140](#))
- ajax crawling wasn’t expanding for unicode urls ([commit 0de3fb4](#))
- Catch start_requests iterator errors. refs #83 ([commit 454a21d](#))
- Speed-up lxml2 XPathSelector ([commit 2fbd662](#))
- updated versioning doc according to recent changes ([commit 0a070f5](#))
- scrapyd: fixed documentation link ([commit 2b4e4c3](#))
- extras/makedeb.py: no longer obtaining version from git ([commit caff0e](#))

8.1.9 0.14.1

- extras/makedeb.py: no longer obtaining version from git ([commit caff0e](#))
- bumped version to 0.14.1 ([commit 6cb9e1c](#))
- fixed reference to tutorial directory ([commit 4b86bd6](#))
- doc: removed duplicated callback argument from Request.replace() ([commit 1aecdd](#))
- fixed formatting of scrapyd doc ([commit 8bf19e6](#))

- Dump stacks for all running threads and fix engine status dumped by StackTraceDump extension ([commit 14a8e6e](#))
- added comment about why we disable ssl on boto images upload ([commit 5223575](#))
- SSL handshaking hangs when doing too many parallel connections to S3 ([commit 63d583d](#))
- change tutorial to follow changes on dmoz site ([commit bcb3198](#))
- Avoid `_disconnectedDeferred` `AttributeError` exception in Twisted $\geq 11.1.0$ ([commit 98f3f87](#))
- allow spider to set autothrottle max concurrency ([commit 175a4b5](#))

8.1.10 0.14

New features and settings

- Support for [AJAX](#) [crawleable](#) urls
- New persistent scheduler that stores requests on disk, allowing to suspend and resume crawls ([r2737](#))
- added `-o` option to `scrapy crawl`, a shortcut for dumping scraped items into a file (or standard output using `-`)
- Added support for passing custom settings to Scrapy `schedule.json` api ([r2779](#), [r2783](#))
- New `ChunkedTransferMiddleware` (enabled by default) to support [chunked transfer encoding](#) ([r2769](#))
- Add boto 2.0 support for S3 downloader handler ([r2763](#))
- Added [marshal](#) to formats supported by feed exports ([r2744](#))
- In request errbacks, offending requests are now received in `failure.request` attribute ([r2738](#))
- **Big downloader refactoring to support per domain/ip concurrency limits ([r2732](#))**
 - **`CONCURRENT_REQUESTS_PER_SPIDER` setting has been deprecated and replaced by:**
 - * `CONCURRENT_REQUESTS`, `CONCURRENT_REQUESTS_PER_DOMAIN`,
`CONCURRENT_REQUESTS_PER_IP`
 - check the documentation for more details
- Added builtin caching DNS resolver ([r2728](#))
- Moved Amazon AWS-related components/extensions (SQS spider queue, SimpleDB stats collector) to a separate project: [scaws](<https://github.com/scrapinghub/scaws>) ([r2706](#), [r2714](#))
- Moved spider queues to scrapyd: `scrapy.spiderqueue` -> `scrapyd.spiderqueue` ([r2708](#))
- Moved sqlite utils to scrapyd: `scrapy.utils.sqlite` -> `scrapyd.sqlite` ([r2781](#))
- Real support for returning iterators on `start_requests()` method. The iterator is now consumed during the crawl when the spider is getting idle ([r2704](#))
- Added `REDIRECT_ENABLED` setting to quickly enable/disable the redirect middleware ([r2697](#))
- Added `RETRY_ENABLED` setting to quickly enable/disable the retry middleware ([r2694](#))
- Added `CloseSpider` exception to manually close spiders ([r2691](#))
- Improved encoding detection by adding support for HTML5 meta charset declaration ([r2690](#))
- Refactored close spider behavior to wait for all downloads to finish and be processed by spiders, before closing the spider ([r2688](#))

- Added `SitemapSpider` (see documentation in Spiders page) (r2658)
- Added `LogStats` extension for periodically logging basic stats (like crawled pages and scraped items) (r2657)
- Make handling of gzipped responses more robust (#319, r2643). Now Scrapy will try and decompress as much as possible from a gzipped response, instead of failing with an `IOError`.
- Simplified `!MemoryDebugger` extension to use stats for dumping memory debugging info (r2639)
- Added new command to edit spiders: `scrapy edit` (r2636) and `-e` flag to `genspider` command that uses it (r2653)
- Changed default representation of items to pretty-printed dicts. (r2631). This improves default logging by making log more readable in the default case, for both Scraped and Dropped lines.
- Added `spider_error` signal (r2628)
- Added `COOKIES_ENABLED` setting (r2625)
- Stats are now dumped to Scrapy log (default value of `STATS_DUMP` setting has been changed to `True`). This is to make Scrapy users more aware of Scrapy stats and the data that is collected there.
- Added support for dynamically adjusting download delay and maximum concurrent requests (r2599)
- Added new DBM HTTP cache storage backend (r2576)
- Added `listjobs.json` API to Scrapyd (r2571)
- `CsvItemExporter`: added `join_multivalued` parameter (r2578)
- Added namespace support to `xmliter_lxml` (r2552)
- Improved cookies middleware by making `COOKIES_DEBUG` nicer and documenting it (r2579)
- Several improvements to Scrapyd and Link extractors

Code rearranged and removed

- **Merged `item passed` and `item scraped` concepts, as they have often proved confusing in the past. This means: (r2630)**
 - original `item_scraped` signal was removed
 - original `item_passed` signal was renamed to `item_scraped`
 - old log lines `Scraped Item...` were removed
 - old log lines `Passed Item...` were renamed to `Scraped Item...` lines and downgraded to `DEBUG` level
- **Reduced Scrapy codebase by striping part of Scrapy code into two new libraries:**
 - `w3lib` (several functions from `scrapy.utils.{http,markup,multipart,response,url}`), done in r2584
 - `scrapely` (was `scrapy.contrib.ibl`, done in r2586)
- Removed unused function: `scrapy.utils.request.request_info()` (r2577)
- Removed `googledir` project from `examples/googledir`. There's now a new example project called `dirbot` available on github: <https://github.com/scrapy/dirbot>
- Removed support for default field values in Scrapy items (r2616)
- Removed experimental `crawlspider v2` (r2632)

- Removed scheduler middleware to simplify architecture. Duplicates filter is now done in the scheduler itself, using the same dupe filtering class as before (`DUPEFILTER_CLASS` setting) (r2640)
- Removed support for passing urls to `scrapy crawl` command (use `scrapy parse` instead) (r2704)
- Removed deprecated Execution Queue (r2704)
- Removed (undocumented) spider context extension (from `scrapy.contrib.spidercontext`) (r2780)
- removed `CONCURRENT_SPIDERS` setting (use `scrapyd maxproc` instead) (r2789)
- Renamed attributes of core components: `downloader.sites` -> `downloader.slots`, `scraper.sites` -> `scraper.slots` (r2717, r2718)
- Renamed setting `CLOSESPIDER_ITEMPASSED` to `CLOSESPIDER_ITEMCOUNT` (r2655). Backwards compatibility kept.

8.1.11 0.12

The numbers like #NNN reference tickets in the old issue tracker (Trac) which is no longer available.

New features and improvements

- Passed item is now sent in the `item` argument of the `item_passed` (#273)
- Added verbose option to `scrapy version` command, useful for bug reports (#298)
- HTTP cache now stored by default in the project data dir (#279)
- Added project data storage directory (#276, #277)
- Documented file structure of Scrapy projects (see command-line tool doc)
- New lxml backend for XPath selectors (#147)
- Per-spider settings (#245)
- Support exit codes to signal errors in Scrapy commands (#248)
- Added `-c` argument to `scrapy shell` command
- Made `libxml2` optional (#260)
- New `deploy` command (#261)
- Added `CLOSESPIDER_PAGECOUNT` setting (#253)
- Added `CLOSESPIDER_ERRORCOUNT` setting (#254)

Scrapyd changes

- Scrapyd now uses one process per spider
- It stores one log file per spider run, and rotate them keeping the latest 5 logs per spider (by default)
- A minimal web ui was added, available at <http://localhost:6800> by default
- There is now a `scrapy server` command to start a Scrapyd server of the current project

Changes to settings

- added `HTTPCACHE_ENABLED` setting (False by default) to enable HTTP cache middleware
- changed `HTTPCACHE_EXPIRATION_SECS` semantics: now zero means “never expire”.

Deprecated/obsoleted functionality

- Deprecated `runserver` command in favor of `server` command which starts a Scrapy server. See also: Scrapy changes
- Deprecated `queue` command in favor of using Scrapy `schedule.json` API. See also: Scrapy changes
- Removed the `!LxmlItemLoader` (experimental contrib which never graduated to main contrib)

8.1.12 0.10

The numbers like #NNN reference tickets in the old issue tracker (Trac) which is no longer available.

New features and improvements

- New Scrapy service called `scrapy` for deploying Scrapy crawlers in production (#218) (documentation available)
- Simplified Images pipeline usage which doesn’t require subclassing your own images pipeline now (#217)
- Scrapy shell now shows the Scrapy log by default (#206)
- Refactored execution queue in a common base code and pluggable backends called “spider queues” (#220)
- New persistent spider queue (based on SQLite) (#198), available by default, which allows to start Scrapy in server mode and then schedule spiders to run.
- Added documentation for Scrapy command-line tool and all its available sub-commands. (documentation available)
- Feed exporters with pluggable backends (#197) (documentation available)
- Deferred signals (#193)
- Added two new methods to item pipeline `open_spider()`, `close_spider()` with deferred support (#195)
- Support for overriding default request headers per spider (#181)
- Replaced default Spider Manager with one with similar functionality but not depending on Twisted Plugins (#186)
- Splitting Debian package into two packages - the library and the service (#187)
- Scrapy log refactoring (#188)
- New extension for keeping persistent spider contexts among different runs (#203)
- Added `dont_redirect` request.meta key for avoiding redirects (#233)
- Added `dont_retry` request.meta key for avoiding retries (#234)

Command-line tool changes

- New *scrapy* command which replaces the old *scrapy-ctl.py* (#199) - there is only one global *scrapy* command now, instead of one *scrapy-ctl.py* per project - Added *scrapy.bat* script for running more conveniently from Windows
- Added bash completion to command-line tool (#210)
- Renamed command *start* to *runserver* (#209)

API changes

- `url` and `body` attributes of `Request` objects are now read-only (#230)
- `Request.copy()` and `Request.replace()` now also copies their `callback` and `errback` attributes (#231)
- Removed `UrlFilterMiddleware` from `scrapy.contrib` (already disabled by default)
- Offsite middleware doesn't filter out any request coming from a spider that doesn't have a `allowed_domains` attribute (#225)
- Removed Spider Manager `load()` method. Now spiders are loaded in the constructor itself.
- **Changes to Scrapy Manager (now called "Crawler"):**
 - `scrapy.core.manager.ScrapyManager` class renamed to `scrapy.crawler.Crawler`
 - `scrapy.core.manager.scrapymanager` singleton moved to `scrapy.project.crawler`
- Moved module: `scrapy.contrib.spidermanager` to `scrapy.spidermanager`
- Spider Manager singleton moved from `scrapy.spider.spiders` to the `spiders` attribute of `scrapy.project.crawler` singleton.
- **moved Stats Collector classes: (#204)**
 - `scrapy.stats.collector.StatsCollector` to `scrapy.statscol.StatsCollector`
 - `scrapy.stats.collector.SimplesdbStatsCollector` to `scrapy.contrib.statscol.SimplesdbStatsCollector`
- default per-command settings are now specified in the `default_settings` attribute of command object class (#201)
- **changed arguments of Item pipeline `process_item()` method from `(spider, item)` to `(item, spider)`**
 - backwards compatibility kept (with deprecation warning)
- **moved `scrapy.core.signals` module to `scrapy.signals`**
 - backwards compatibility kept (with deprecation warning)
- **moved `scrapy.core.exceptions` module to `scrapy.exceptions`**
 - backwards compatibility kept (with deprecation warning)
- added `handles_request()` class method to `BaseSpider`
- dropped `scrapy.log.exc()` function (use `scrapy.log.err()` instead)
- dropped component argument of `scrapy.log.msg()` function
- dropped `scrapy.log.log_level` attribute

- Added `from_settings()` class methods to Spider Manager, and Item Pipeline Manager

Changes to settings

- Added `HTTPCACHE_IGNORE_SCHEMES` setting to ignore certain schemes on `!HttpCacheMiddleware` (#225)
- Added `SPIDER_QUEUE_CLASS` setting which defines the spider queue to use (#220)
- Added `KEEP_ALIVE` setting (#220)
- Removed `SERVICE_QUEUE` setting (#220)
- Removed `COMMANDS_SETTINGS_MODULE` setting (#201)
- Renamed `REQUEST_HANDLERS` to `DOWNLOAD_HANDLERS` and make download handlers classes (instead of functions)

8.1.13 0.9

The numbers like #NNN reference tickets in the old issue tracker (Trac) which is no longer available.

New features and improvements

- Added SMTP-AUTH support to scrapy.mail
- New settings added: `MAIL_USER`, `MAIL_PASS` (r2065 | #149)
- Added new scrapy-ctl view command - To view URL in the browser, as seen by Scrapy (r2039)
- Added web service for controlling Scrapy process (this also deprecates the web console. (r2053 | #167)
- Support for running Scrapy as a service, for production systems (r1988, r2054, r2055, r2056, r2057 | #168)
- Added wrapper induction library (documentation only available in source code for now). (r2011)
- Simplified and improved response encoding support (r1961, r1969)
- Added `LOG_ENCODING` setting (r1956, documentation available)
- Added `RANDOMIZE_DOWNLOAD_DELAY` setting (enabled by default) (r1923, doc available)
- `MailSender` is no longer IO-blocking (r1955 | #146)
- `Linkextractors` and new `Crawlspider` now handle relative base tag urls (r1960 | #148)
- Several improvements to Item Loaders and processors (r2022, r2023, r2024, r2025, r2026, r2027, r2028, r2029, r2030)
- Added support for adding variables to telnet console (r2047 | #165)
- Support for requests without callbacks (r2050 | #166)

API changes

- Change `Spider.domain_name` to `Spider.name` (SEP-012, r1975)
- `Response.encoding` is now the detected encoding (r1961)
- `HttpErrorMiddleware` now returns `None` or raises an exception (r2006 | #157)
- scrapy.command modules relocation (r2035, r2036, r2037)

- Added `ExecutionQueue` for feeding spiders to scrape (r2034)
- Removed `ExecutionEngine` singleton (r2039)
- Ported `S3ImagesStore` (images pipeline) to use boto and threads (r2033)
- Moved module: `scrapy.management.telnet` to `scrapy.telnet` (r2047)

Changes to default settings

- Changed default `SCHEDULER_ORDER` to `DFO` (r1939)

8.1.14 0.8

The numbers like #NNN reference tickets in the old issue tracker (Trac) which is no longer available.

New features

- Added `DEFAULT_RESPONSE_ENCODING` setting (r1809)
- Added `dont_click` argument to `FormRequest.from_response()` method (r1813, r1816)
- Added `clickdata` argument to `FormRequest.from_response()` method (r1802, r1803)
- Added support for HTTP proxies (`HttpProxyMiddleware`) (r1781, r1785)
- Offsite spider middleware now logs messages when filtering out requests (r1841)

Backwards-incompatible changes

- Changed `scrapy.utils.response.get_meta_refresh()` signature (r1804)
- Removed deprecated `scrapy.item.ScrapedItem` class - use `scrapy.item.Item` instead (r1838)
- Removed deprecated `scrapy.xpath` module - use `scrapy.selector` instead. (r1836)
- Removed deprecated `core.signals.domain_open` signal - use `core.signals.domain_opened` instead (r1822)
- **`log.msg()` now receives a `spider` argument (r1822)**
 - Old `domain` argument has been deprecated and will be removed in 0.9. For spiders, you should always use the `spider` argument and pass spider references. If you really want to pass a string, use the `component` argument instead.
- Changed core signals `domain_opened`, `domain_closed`, `domain_idle`
- **Changed Item pipeline to use spiders instead of domains**
 - The `domain` argument of `process_item()` item pipeline method was changed to `spider`, the new signature is: `process_item(spider, item)` (r1827 | #105)
 - To quickly port your code (to work with Scrapy 0.8) just use `spider.domain_name` where you previously used `domain`.
- **Changed Stats API to use spiders instead of domains (r1849 | #113)**
 - `StatsCollector` was changed to receive spider references (instead of domains) in its methods (`set_value`, `inc_value`, etc).

- added `StatsCollector.iter_spider_stats()` method
- removed `StatsCollector.list_domains()` method
- Also, Stats signals were renamed and now pass around spider references (instead of domains). Here's a summary of the changes:
- To quickly port your code (to work with Scrapy 0.8) just use `spider.domain_name` where you previously used `domain`. `spider_stats` contains exactly the same data as `domain_stats`.
- **CloseDomain extension moved to `scrapy.contrib.closespider.CloseSpider` (r1833)**
 - **Its settings were also renamed:**
 - * `CLOSEDOMAIN_TIMEOUT` to `CLOSESPIDER_TIMEOUT`
 - * `CLOSEDOMAIN_ITEMCOUNT` to `CLOSESPIDER_ITEMCOUNT`
- Removed deprecated `SCRAPYSETTINGS_MODULE` environment variable - use `SCRAPY_SETTINGS_MODULE` instead (r1840)
- Renamed setting: `REQUESTS_PER_DOMAIN` to `CONCURRENT_REQUESTS_PER_SPIDER` (r1830, r1844)
- Renamed setting: `CONCURRENT_DOMAINS` to `CONCURRENT_SPIDERS` (r1830)
- Refactored HTTP Cache middleware
- HTTP Cache middleware has been heavily refactored, retaining the same functionality except for the domain sectorization which was removed. (r1843)
- Renamed exception: `DontCloseDomain` to `DontCloseSpider` (r1859 | #120)
- Renamed extension: `DelayedCloseDomain` to `SpiderCloseDelay` (r1861 | #121)
- Removed obsolete `scrapy.utils.markup.remove_escape_chars` function - use `scrapy.utils.markup.replace_escape_chars` instead (r1865)

8.1.15 0.7

First release of Scrapy.

8.2 Contributing to Scrapy

There are many ways to contribute to Scrapy. Here are some of them:

- Blog about Scrapy. Tell the world how you're using Scrapy. This will help newcomers with more examples and the Scrapy project to increase its visibility.
- Report bugs and request features in the [issue tracker](#), trying to follow the guidelines detailed in [Reporting bugs](#) below.
- Submit patches for new functionality and/or bug fixes. Please read [Writing patches](#) and [Submitting patches](#) below for details on how to write and submit a patch.
- Join the [scrapy-developers](#) mailing list and share your ideas on how to improve Scrapy. We're always open to suggestions.

8.2.1 Reporting bugs

Well-written bug reports are very helpful, so keep in mind the following guidelines when reporting a new bug.

- check the [FAQ](#) first to see if your issue is addressed in a well-known question
- check the [open issues](#) to see if it has already been reported. If it has, don't dismiss the report but check the ticket history and comments, you may find additional useful information to contribute.
- search the [scrapy-users](#) list to see if it has been discussed there, or if you're not sure if what you're seeing is a bug. You can also ask in the [#scrapy](#) IRC channel.
- write complete, reproducible, specific bug reports. The smaller the test case, the better. Remember that other developers won't have your project to reproduce the bug, so please include all relevant files required to reproduce it.
- include the output of `scrapy version -v` so developers working on your bug know exactly which version and platform it occurred on, which is often very helpful for reproducing it, or knowing if it was already fixed.

8.2.2 Writing patches

The better written a patch is, the higher chance that it'll get accepted and the sooner that will be merged.

Well-written patches should:

- contain the minimum amount of code required for the specific change. Small patches are easier to review and merge. So, if you're doing more than one change (or bug fix), please consider submitting one patch per change. Do not collapse multiple changes into a single patch. For big changes consider using a patch queue.
- pass all unit-tests. See [Running tests](#) below.
- include one (or more) test cases that check the bug fixed or the new functionality added. See [Writing tests](#) below.
- if you're adding or changing a public (documented) API, please include the documentation changes in the same patch. See [Documentation policies](#) below.

8.2.3 Submitting patches

The best way to submit a patch is to issue a [pull request](#) on Github, optionally creating a new issue first.

Alternatively, we also accept the patches in the traditional way of sending them to the [scrapy-developers](#) list.

Regardless of which mechanism you use, remember to explain what was fixed or the new functionality (what it is, why it's needed, etc). The more info you include, the easier will be for core developers to understand and accept your patch.

You can also discuss the new functionality (or bug fix) in [scrapy-developers](#) first, before creating the patch, but it's always good to have a patch ready to illustrate your arguments and show that you have put some additional thought into the subject.

8.2.4 Coding style

Please follow these coding conventions when writing code for inclusion in Scrapy:

- Unless otherwise specified, follow [PEP 8](#).
- It's OK to use lines longer than 80 chars if it improves the code readability.

- Don't put your name in the code you contribute. Our policy is to keep the contributor's name in the [AUTHORS](#) file distributed with Scrapy.

8.2.5 Documentation policies

- **Don't** use docstrings for documenting classes, or methods which are already documented in the official (sphinx) documentation. For example, the `ItemLoader.add_value()` method should be documented in the sphinx documentation, not its docstring.
- **Do** use docstrings for documenting functions not present in the official (sphinx) documentation, such as functions from `scrapy.utils` package and its sub-modules.

8.2.6 Tests

Tests are implemented using the [Twisted unit-testing framework](#) called `trial`.

Running tests

To run all tests go to the root directory of Scrapy source code and run:

```
bin/runtests.sh (on unix)
bin\runtests.bat (on windows)
```

To run a specific test (say `scrapy.tests.test_contrib_loader`) use:

```
bin/runtests.sh scrapy.tests.test_contrib_loader (on unix)
bin\runtests.bat scrapy.tests.test_contrib_loader (on windows)
```

Writing tests

All functionality (including new features and bug fixes) must include a test case to check that it works as expected, so please include tests for your patches if you want them to get accepted sooner.

Scrapy uses unit-tests, which are located in the `scrapy.tests` package ([scrapy/tests](#) directory). Their module name typically resembles the full path of the module they're testing. For example, the item loaders code is in:

```
scrapy.contrib.loader
```

And their unit-tests are in:

```
scrapy.tests.test_contrib_loader
```

8.3 Versioning and API Stability

8.3.1 Versioning

Scrapy uses the [odd-numbered versions](#) for development releases.

There are 3 numbers in a Scrapy version: *A.B.C*

- *A* is the major version. This will rarely change and will signify very large changes. So far, only zero is available for *A* as Scrapy hasn't yet reached 1.0.

- *B* is the release number. This will include many changes including features and things that possibly break backwards compatibility. Even Bs will be stable branches, and odd Bs will be development.
- *C* is the bugfix release number.

For example:

- *0.14.1* is the first bugfix release of the *0.14* series (safe to use in production)

8.3.2 API Stability

API stability is one of Scrapy major goals for the *1.0* release, which doesn't have a due date scheduled yet.

Methods or functions that start with a single dash (`_`) are private and should never be relied as stable. Besides those, the plan is to stabilize and document the entire API, as we approach the 1.0 release.

Also, keep in mind that stable doesn't mean complete: stable APIs could grow new methods or functionality but the existing methods should keep working the same way.

8.4 Experimental features

This section documents experimental Scrapy features that may become stable in future releases, but whose API is not yet stable. Use them with caution, and subscribe to the [mailing lists](#) to get notified of any changes.

Since it's not revised so frequently, this section may contain documentation which is outdated, incomplete or overlapping with stable documentation (until it's properly merged) . Use at your own risk.

Warning: This documentation is a work in progress. Use at your own risk.

No experimental features at this time

Release notes See what has changed in recent Scrapy versions.

Contributing to Scrapy Learn how to contribute to the Scrapy project.

Versioning and API Stability Understand Scrapy versioning and API stability.

Experimental features Learn about bleeding-edge features.

Python Module Index

S

scrapy.contracts, 86
scrapy.contracts.default, 86
scrapy.contrib.clospider, 138
scrapy.contrib.corestats, 137
scrapy.contrib.debug, 139
scrapy.contrib.downloadermiddleware, 124
scrapy.contrib.downloadermiddleware.chunked, 129
scrapy.contrib.downloadermiddleware.cookies, 126
scrapy.contrib.downloadermiddleware.defaultheaders, 127
scrapy.contrib.downloadermiddleware.downloadtimeout, 127
scrapy.contrib.downloadermiddleware.httpauth, 127
scrapy.contrib.downloadermiddleware.httpeagle, 127
scrapy.contrib.downloadermiddleware.httpcompression, 129
scrapy.contrib.downloadermiddleware.httpproxy, 129
scrapy.contrib.downloadermiddleware.redirect, 129
scrapy.contrib.downloadermiddleware.retry, 130
scrapy.contrib.downloadermiddleware.robotstxt, 131
scrapy.contrib.downloadermiddleware.stats, 131
scrapy.contrib.downloadermiddleware.useragent, 131
scrapy.contrib.exporter, 167
scrapy.contrib.linkextractors, 61
scrapy.contrib.linkextractors.sgml, 61
scrapy.contrib.loader, 44
scrapy.contrib.loader.processor, 50
scrapy.contrib.logstats, 137
scrapy.contrib.memdebug, 138
scrapy.contrib.memusage, 138
scrapy.contrib.pipeline.images, 102
scrapy.contrib.spidermiddleware, 132
scrapy.contrib.spidermiddleware.depth, 133
scrapy.contrib.spidermiddleware.httperror, 134
scrapy.contrib.spidermiddleware.offsite, 134
scrapy.contrib.spidermiddleware.referrer, 134
scrapy.contrib.spidermiddleware.urllength, 135
scrapy.contrib.spiders, 31
scrapy.contrib.statsmailer, 139
scrapy.contrib.webservice, 73
scrapy.contrib.webservice.crawler, 73
scrapy.contrib.webservice.enginestatus, 73
scrapy.contrib.webservice.stats, 73
scrapy.crawler, 140
scrapy.exceptions, 166
scrapy.http, 145
scrapy.item, 24
scrapy.log, 66
scrapy.mail, 68
scrapy.selector, 41
scrapy.settings, 141
scrapy.signalmanager, 142
scrapy.signals, 164
scrapy.spider, 29
scrapy.statscol, 68
scrapy.telnet, 70
scrapy.utils.trackref, 97
scrapy.webservice, 137